

Elm IN ACTION

Richard Feldman



MEAP



MANNING



MEAP Edition
Manning Early Access Program
Elm in Action
Version 4

Copyright 2017 Manning Publications

For more information on this and other Manning titles go to
www.manning.com

welcome

Thanks for buying the MEAP of *Elm in Action!* I hope you find this a wonderfully practical journey into a delightful language.

As an early adopter of Elm, I felt like a wide-eyed sprinter barreling down an exciting path that had not yet been paved. Now that I've had time to catch my breath, I'm eager to help pave that road for future travelers. My goal for this book is for you to enjoy learning Elm even more than I did...and I loved it so much I'm writing a book about it!

More than anything, I hope this book gives you the confidence to use Elm professionally. As much as I enjoyed using Elm on a side project, it's brought me even more joy at work. When I think back on what motivated my team to try it, it was the practical benefits: reliability, a quicker development loop, and a lower maintenance burden.

The most authentic way I know to convey the benefits of building an application in Elm is...well, to build an application in Elm. That's exactly what we'll be doing over the course of this book: developing an application from start to finish, learning the language along the way. By the end you'll have built, refactored, and tested an Elm code base. If using it at work sounds appealing, you'll be able to convey its benefits based on firsthand experience!

Throughout the MEAP process I intend both to release new chapters as I finish them, and to revise past chapters based on feedback from readers like you. I'd love it if you could leave some comments in the Author Online Forum. All feedback is helpful! Mentioning which parts felt difficult tells me where to focus my revision time, and mentioning your favorite parts tells me what not to change.

Thanks again, and enjoy!

—Richard Feldman

brief contents

PART 1: GETTING STARTED

- 1 Welcome to Elm*
- 2 Your First Elm Application*
- 3 Compiler as Assistant*

PART 2: PRODUCTION-GRADE ELM

- 4 Talking to Servers*
- 5 Talking to JavaScript*
- 6 Testing*
- 7 Tools*

PART 3: BUILDING BIGGER

- 8 Single-Page Applications*
- 9 Scaling Elm Code*
- 10 Performance Optimization*

Appendix A: Getting Set Up

1

Welcome to Elm

This chapter covers

- How and why to introduce Elm to a project
- Using `elm-repl`
- Building expressions
- Writing and importing functions
- Working with collections

Back in 2014 I set out to rewrite a side project, and ended up with a new favorite programming language. Not only was the rewritten code faster, more reliable, and easier to refactor, writing it was the most fun I'd had in over a decade writing code. Ever since that project, I've been hooked on Elm.

The rewrite in question was a writing application I'd built even longer ago, in 2011. Having tried out several writing apps over the course of writing a novel, and being satisfied with none, I decided to scratch my own itch and build my dream writing app. I called it Dreamwriter.

For those keeping score: yes, I was indeed writing code in order to write prose better.

Things went well at first. I built the basic Web app, started using it, and iterated on the design. Months later I'd written over fifty thousand words in Dreamwriter. If I'd been satisfied with that early design, the story might have ended there. However, users always want a better experience...and when the user and the developer are the same person, further iteration is inevitable.

The more I revised Dreamwriter, the more difficult it became to maintain. I'd spend hours trying to reproduce bugs that knocked me out of my writing groove. At some point the copy and paste functions stopped working, and I found myself resorting to the browser's developer tools whenever I needed to move paragraphs around.

Right around when I'd decided to scrap my unreliable code base and do a full rewrite, a blog post crossed my radar. After reading it I knew three things:

1. The Elm programming language compiled to JavaScript, just like Babel or CoffeeScript. (I already had a compile step in my build script, so this was familiar territory.)
2. Elm used the same rendering approach as React.js—which I had recently grown to love—except Elm had rendering benchmarks that outperformed React's!
3. Elm's compiler would catch a lot of the errors I'd been seeing before they could harm me in production. I did not yet know just how many it would catch.

I'd never built anything with a functional programming language like Elm before, but I decided to take the plunge. I didn't really know what I was doing, but the compiler's error messages kept picking me up whenever I stumbled. Eventually I got the revised version up and running, and began to refactor.

The refactoring experience blew me away. I revised the Elm-powered Dreamwriter gleefully, even recklessly—and no matter how dramatic my changes, the compiler always had my back. It would point out whatever corner cases I'd missed, and I'd go through and fix them. As soon as the code compiled, lo and behold, everything worked again. I felt *invincible*.

I related my Elm experience to my coworkers at NoRedInk, and they were curious but understandably cautious. How could we find out if the team liked it without taking a big risk? A full rewrite may have been fine for Dreamwriter, but it would have been crazy to attempt that for our company's entire front-end.

So we introduced Elm gently, by rewriting just one portion of one production feature in Elm. It went well, so we did a bit more. And then more.

Today our front-end is as Elm-powered as we can make it, and our team has never been happier. Our test suites are smaller, yet our product is more reliable. Our feature set has grown more complex, yet refactoring remains delightful. We swap stories with other companies using Elm about how long our production code has run without throwing a single runtime exception.

In this book we'll explore all of these benefits.

After learning some basics, we'll build an Elm Web application the way teams typically do: ship a basic version that works, but which has missing features and some technical debt. As we advance through the chapters, we'll expand and refactor this application, adding features and paying off technical debt as we learn more about Elm. We'll debug our business logic. We'll even interoperate with JavaScript, using its vast library ecosystem to avoid reinventing the wheel. By the end of the book we will have transformed our application into a more featureful product, with a more maintainable code base, than the one we initially shipped.

With any luck, we'll have a lot of fun doing it.

Welcome to Elm!

1.1 How Elm Fits In

Elm can be used either as a replacement for in-browser JavaScript code, or as a complement to it. You write some .elm files, run them through Elm's compiler, and end up with plain old .js files that the browser runs as normal. If you have separate stylesheets that you use alongside JavaScript, they'll work the same way alongside Elm.

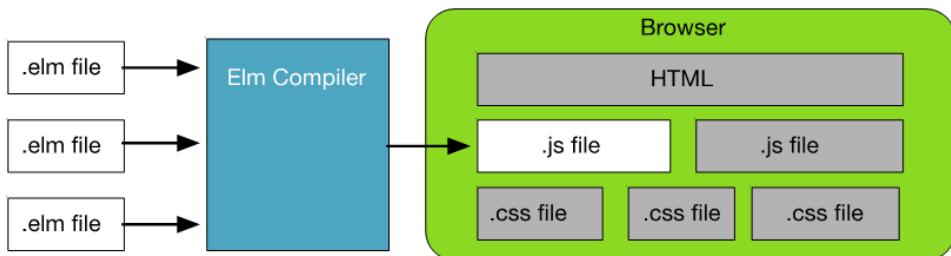


Figure 1.1 Elm files are compiled to plain old JavaScript files

The appropriate Elm-to-JavaScript ratio can vary by project. Some projects may want primarily JavaScript and only a touch of Elm for business logic or rendering. Others may want a great deal of Elm but just a pinch of JavaScript to leverage its larger ecosystem. No single answer applies to every project.

What distinguishes Elm from JavaScript is **maintainability**.

Handwritten JavaScript code is notoriously prone to runtime crashes like “undefined is not a function.” In contrast, Elm code has a reputation for never throwing runtime exceptions in practice. This is because Elm is built around a small set of simple primitives like expressions, immutable values, and managed effects. That design lets the compiler identify lurking problems just by scanning your source code. It reports these problems with such clarity that it has earned a reputation for user-friendliness even among programming legends.

That should be an inspiration for every error message.

—John Carmack, after seeing one of Elm's compiler errors

Having this level of compiler assistance makes Elm code dramatically easier to refactor and debug, especially as code bases grow larger. There is an up-front cost to learning and adopting Elm, but you reap more and more maintainability benefits the longer the project remains in active development.

TIP Most teams that use Elm in production say they used a “planting the seed” approach. Instead of waiting for a big project where they could build everything in Elm from the ground up, they rewrote a small part of their existing JavaScript code base in Elm. This was low-risk and could be rolled back if things did not go as planned,

but having that small seed planted in production meant they could grow their Elm code at a comfortable pace from then on.

Although Elm is in many ways a simpler language than JavaScript, it's also much younger. This means Elm has fewer off-the-shelf solutions available for any given problem. Elm code can interoperate with JavaScript code to piggyback the larger JavaScript library ecosystem, but Elm's design differs enough from JavaScript's that incorporating JavaScript libraries takes effort.

Balancing these tradeoffs depends on the specifics of a given project. Let's say you're on a team where people are comfortable with JavaScript but are new to Elm. Here are some projects I'd expect would benefit from learning and using Elm:

- Feature-rich Web applications whose code bases are large or will grow large
- Individual features that will be revised and maintained over an extended period of time
- Projects where most functionality comes from in-house code, not off-the-shelf libraries

In contrast, I'd stick to a more familiar language and toolset for projects like these:

- Proof-of-concept prototypes that will not be maintained long-term
- Time-crunched projects where learning a language is unrealistic given the deadline
- Projects that will consist primarily of gluing together off-the-shelf components

We'll explore these tradeoffs in more detail throughout the course of the book.

1.2 Expressions

To get our feet wet with Elm, let's tap into one of the most universal traits across the animal kingdom: the innate desire to *play*.

Researchers have developed many theories as to why we play, including to learn, to practice, to experiment, and of course for the pure fun of it. These researchers could get some high-quality data by observing a member of the *homo sapiens programmerus* species in its natural environment for play—the Read-Eval-Print Loop, or REPL.

You'll be using Elm's REPL to play as you take your first steps as an Elm programmer.

1.2.1 Using elm-repl

The Elm Platform includes a nice REPL called `elm-repl`, so if you have not installed the Elm Platform yet, head over to Appendix A to get hooked up.

Once you're ready, enter `elm-repl` at the terminal. You should see this prompt:

```
---- elm-repl 0.18.0 -----
:help for help, :exit to exit, more at https://github.com/elm-lang/elm-repl
----->
```

Alexander Graham Bell invented the telephone over a century ago. There was no customary greeting back then, so Bell suggested one: lift the receiver and bellow out a rousing “Ahoy!” Thomas Edison later proposed the alternative “Hello,” which stuck, and today programmers everywhere append “World” as the customary way to greet a new programming language.

Let’s spice things up a bit, shall we? Enter this at the prompt.

```
> "Ahoy, World!"
```

You should see this response from `elm-repl`:

```
"Ahoy, World!" : String
```

Congratulations, you are now an Elm programmer!

NOTE To focus on the basics, for the rest of this chapter we’ll omit the type annotations that `elm-repl` prints. For example, the previous code snippet would have omitted the `: String` portion of `"Ahoy, World!" : String`. We’ll get into these annotations in Chapter 3.

If you’re the curious sort, by all means feel free to play as we continue. Enter things that occur to you, and see what happens! Whenever you encounter an error you don’t understand yet, picture yourself as a tiger cub building intuition for physics through experimentation: adorable for now, but powerful in time.

1.2.2 Building Expressions

Let’s rebuild our “Ahoy, World!” greeting from two parts, and then play around from there. Try entering these into `elm-repl`.

Listing 1.1 Combining Strings

```
> "Ahoy, World!"
"Ahoy, World!"

> "Ahoy, " ++ "World!"
"Ahoy, World!"

> "Pi is " ++ toString pi ++ " (give or take)" ①
"Pi is 3.141592653589793 (give or take)"
```

① `toString` is a standalone function, not a method

In Elm, we use the `++` operator to combine strings, instead of the `+` operator JavaScript uses. At this point you may be wondering: Does Elm even have a `+` operator? What about the other arithmetic operators?

Let’s find out by experimenting in `elm-repl`!

Listing 1.2 Arithmetic Expressions

```
> 1234 + 103
```

```

1337

> 12345 - (5191 * -15) ①
90210

> 2 ^ 11
2048

> 49 / 10
4.9

> 49 // 10          ②
4

> -5 % 2           ③
1

```

- ① Nest expressions via parentheses
- ② Integer division (decimals get truncated)
- ③ Remainder after integer division

Sure enough, Elm has both a `++` and a `+` operator. They are used for different things:

- The `++` operator is for appending. Using it on a number is an error.
- The `+` operator is for addition. It can *only* be used on numbers.

You will see this preference for **being explicit** often in Elm. If two operations are sufficiently different—in this case, adding and appending—Elm implements them separately, so each implementation can **do one thing well**.

STRINGS AND CHARACTERS

Elm also distinguishes between strings and the individual UTF-8 *characters* that comprise them. Double quotes in Elm represent string literals, just like in JavaScript, but single quotes in Elm represent character literals.

Table 1.1 Strings and Characters

Elm Literal	Result
<code>"a"</code>	a string with a length of 1
<code>'a'</code>	a single character
<code>"abc"</code>	a string with a length of 3
<code>'abc'</code>	error: character literals must contain exactly 1 character
<code>""</code>	an empty string
<code>''</code>	error: character literals must contain exactly 1 character

COMMENTS

Comments also work a bit differently in Elm:

- Use `--` instead of `//` for inline comments
- Use `{- -}` instead of `/* */` for block comments

Let's see these in action!

Listing 1.3 Characters, Comments, and Constants

```
> 'a'  -- This is a comment. It will be ignored. ①
'a'

> "a" {- This is a block comment. It will also be ignored. -} ②
"a"

> milesPerHour = 88 ③
88

> milesPerHour
88
```

- ① JavaScript comment: `//`
- ② JavaScript comment: `/* ... */`
- ③ JavaScript: const `milesPerHour = 88;`

NAMING VALUES WITH CONSTANTS

In the last two lines of code above, we did something new: we assigned the *constant* `milesPerHour` to the value `88`.

DEFINITION A *constant* assigns a name to a value. Once assigned, this name cannot be later reassigned to a different value in the same scope.

There are a few things to keep in mind when naming constants.

- The name must begin with a lowercase letter. After that it can be a mix of letters, numbers, and underscores.
- By convention, all letters should be in one uninterrupted sequence. For example, `map4` is a reasonable name, but `map4ever` is not, because the sequence of letters is interrupted by the `4`.
- Because of the previous two rules, you should never use `snake_case` or `SCREAMING_SNAKE_CASE` to name constants. Use `camelCase` instead.
- If you absolutely must know whether the compiler will accept `some_raD_TH1NG__` as a valid constant name, remember: what happens in `elm-repl` stays in `elm-repl`.

ASSIGNING CONSTANTS TO EXPRESSIONS

Not only can you assign constants to literal values, you can also assign them to *expressions*.

DEFINITION An *expression* is anything that evaluates to a single value.

Here are some expressions we've seen so far.

Expression	Evaluates to
"Ahoy, " ++ "World!"	"Ahoy, World!"
2 ^ 11	2048
pi	3.141592653589793
42	42

NOTE Since an expression is anything that evaluates to a value, literal values like "Ahoy, World!" and 42 are expressions too—just expressions that have already been fully evaluated.

Expressions are the basic building block of Elm applications. This is different from JavaScript, which offers many features as *statements* instead of expressions.

Consider these two lines of JavaScript code.

```
label = (num > 0) ? "positive" : "negative" // ternary expression
label = if (num > 0) { "positive" } else { "negative" } // if-statement
```

The first line is *ternary expression*. Being an expression, it evaluates to a value, and JavaScript happily assigns that value to `label`.

The second line is an *if-statement*, and since statements do not evaluate to values, trying to assign it to `label` yields a syntax error.

This distinction does not exist in Elm, as Elm programs express logic using expressions only. As such, Elm has *if-expressions* instead of *if-statements*. As we will see in Chapter 2, every Elm application is essentially one big expression built up from lots of smaller ones!

1.2.3 Booleans and Conditionals

There aren't terribly many booleans out there—just the two, really—and working with them in Elm is similar to working with them in JavaScript. There are a few differences, though.

- You write `True` and `False` instead of `true` and `false`
- You write `/=` instead of `!=`
- To negate values, you use Elm's `not` function instead of JavaScript's `!` prefix

Let's try them out!

Listing 1.4 Boolean Expressions

```
> pi == pi      ①
True           ②

> pi /= pi     ③
False          ④

> not (pi == pi) ⑤
False

> pi <= 0 || pi >= 10
False

> 3 < pi && pi < 4 ⑥
True
```

- ① JavaScript: `pi === pi`
- ② JavaScript: `true`
- ③ JavaScript: `pi !== pi`
- ④ JavaScript: `false`
- ⑤ JavaScript: `!(pi === pi)`
- ⑥ `3 < pi < 4` would be an error

Now let's say it's a lovely afternoon at the North Pole, and we're in Santa's workshop writing a bit of UI logic to display how many elves are currently on vacation. The quick-and-dirty approach would be to add the string "`elves`" after the number of vacationing elves, but then when the count is 1 we'd display "`1 elves`", and we're better than that.

Let's polish our user experience with the *if-expression* shown in Figure 1.2.

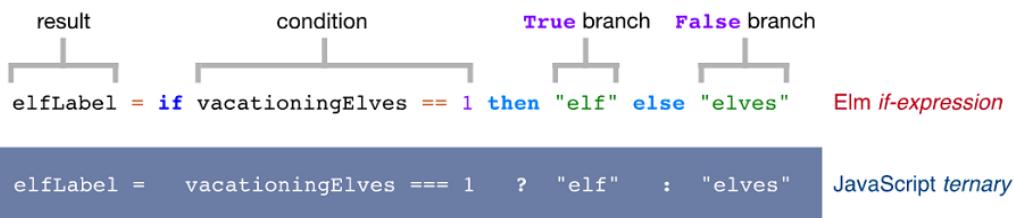


Figure 1.2 Comparing an Elm *if-expression* to a JavaScript *ternary*

Like JavaScript ternaries, Elm *if-expressions* require three ingredients:

1. A condition
2. A branch to evaluate if the condition passes
3. A branch to evaluate otherwise

Each of these ingredients must be expressions, and the *if-expression* itself evaluates to the result of whichever branch got evaluated. You'll get an error if any of these three ingredients are missing, so make sure to specify an `else` branch every time!

NOTE There is no such thing as “truthiness” in Elm. Conditions can be either `True` or `False`, and that’s it. Life is simpler this way.

Now let’s say we modified our pluralization conditional to include a third case:

- If we have one Elf, evaluate to “elf”
- Otherwise, if we have a positive number of elves, evaluate to “elves”
- Otherwise, we must have a negative number of elves, so evaluate to “anti-elves”

In JavaScript you may have used `else if` to continue branching conditionals like this. It’s common to use `else if` for the same purpose in Elm, but it’s worth noting that `else if` in Elm is nothing more than a stylish way to combine the concepts we learned a moment ago.

Check it out!

Listing 1.6 Using `else if`

```
if elfCount == 1 then          1
  "elf"
else                         1
  (if elfCount >= 0 then "elves" else "anti-elves") 1

if elfCount == 1 then          2
  "elf"
else (if elfCount >= 0 then 2
  "elves"
else                         2
  "anti-elves")               2

if elfCount == 1 then          3
  "elf"
else if elfCount >= 0 then   3
  "elves"
else                         3
  "anti-elves"                3
```

- ① Use an *if-expression* inside `else`
- ② Rearrange some whitespace
- ③ Drop the parentheses

This works because the `else` branch of an *if-expression* must be an expression, and it just so happens that *if-expressions* themselves are expressions. As shown in Figure 1.3, all it takes is putting an *if-expression* after another one’s `else`, and *voilà!* Additional branching achieved.

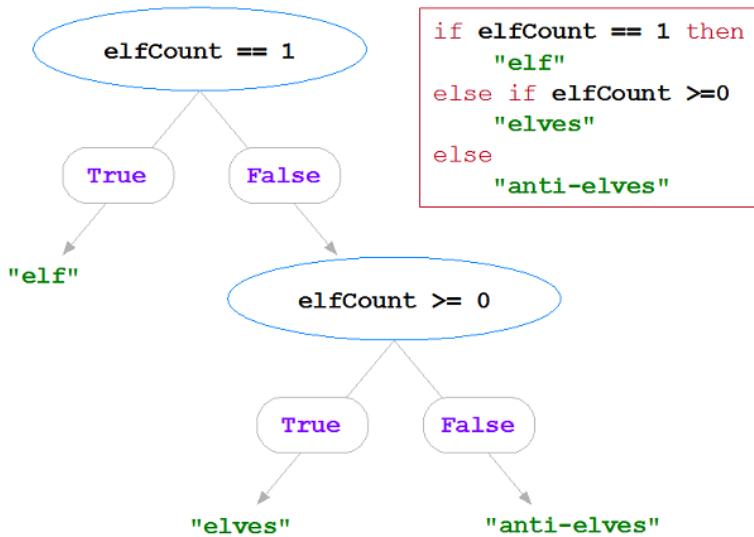


Figure 1.3 The `else if` technique: use an `if-expression` as the `else` branch of another `if-expression`

Nesting expressions for fun and profit is a recurring theme in Elm, and we'll see plenty more recipes like `else if` throughout the book.

Chapter 3 will add a powerful new conditional to our expression toolbox, one which has no analogue in JavaScript: the *case-expression*.

1.3 Functions

Earlier we wrote this expression:

```
elfLabel = if vacationingElves == 1 then "elf" else "elves"
```

Suppose it turns out a general-purpose singular/plural labeler would be really useful, and we want to reuse similar logic across the code base at Santa's Workshop. Search results might want to display "1 result" and "2 results" as appropriate, for example.

We can write a *function* to make this pluralization logic easily reusable.

DEFINITION Elm *functions* represent reusable logic. They are not objects. They have no fields, no prototypes, and no ability to store state. All they do is accept values as arguments, and then return a value.

If you thought expressions would be a recurring theme in Elm, wait 'til you see functions!

1.3.1 Defining Functions

Let's define our first function: `isOdd`. It will take a number and then:

- Return `True` if the number is odd
- Return `False` if the number is even

We can define `isOdd` in `elm-repl` and try it out right away.

Listing 1.7 Defining a function

```
> isOdd num = num % 2 == 1    ①
<function>

> isOdd 5                  ②
True

> (isOdd 5)                ③
True

> isOdd (5 + 1)            ④
False
```

① JavaScript: `function isOdd(num) { return num % 2 === 1; }`
② JavaScript: `isOdd(5)`
③ JavaScript: `(isOdd(5))`
④ JavaScript: `isOdd(5 + 1)`

As you can see, in Elm we put the function parameter name before the `=` sign. We also don't surround the function body with `{ }`. And did you notice the `return` keyword is nowhere to be seen? That's because Elm doesn't have one! In Elm, a function body is a single expression, and since an expression evaluates to a single value, Elm uses that value as the function's return value. This means all Elm functions return values!

For our `isOdd` function, the expression `num % 2 == 1` serves as the function's body, and provides its return value.

Refactoring out an early return

In JavaScript, `return` is often used to exit a function early. This is harmless when used responsibly, but can lead to unpleasant surprises when used in the middle of large functions. Elm does not support these unpleasant surprises, because it has no `return` keyword.

Let's refactor the early `return` out of this function:

```
function capitalize(str) {
  if (!str) {
    return str;    ①
  }

  return str[0].toUpperCase() + str.slice(1);
}
```

① Early return

Without making any other changes, we can refactor this early `return` into a ternary:

```
function capitalize(str) {
    return !str ? str : str[0].toUpperCase() + str.slice(1);
}
```

Poof! There it goes. Since JavaScript's ternaries are semantically analogous to Elm's *if-expressions*, this code is now much more straightforward to rewrite in Elm. More convoluted JavaScript functions may require more steps than this, but it is always possible to untangle them into plain old conditionals.

Removing an early `return` is one of many quick refactors you can do to ease the transition from legacy JavaScript to Elm, and we'll look at more of them throughout the book. When doing these, do not worry if the intermediate JavaScript code looks ugly! It's intended to be a stepping stone to nicer Elm code, not something to be maintained long-term.

Let's use what we just learned to generalize our previous `elf-labeling` expression into a reusable `pluralize` function. Our function this time will have a longer definition than last time, so let's use multiple lines to give it some breathing room. In `elm-repl`, you can enter multiple lines by adding `\` to the end of the first line and indenting the next line.

NOTE Indent with spaces only! Tab characters are syntax errors in Elm.

Listing 1.8 Using multiple REPL lines

```
> pluralize singular plural count = \
|   if count == 1 then singular else plural  ①
<function>

> pluralize "elf" "elves" 3                ②
"elves"

> pluralize "elf" "elves" (round 0.9)      ③
"elf"
```

- ① Don't forget to indent!
- ② No commas between arguments!
- ③ `(round 0.9)` returns 1

When passing multiple arguments to an Elm function, separate the arguments with whitespace and not commas. That last line of code is an example of passing the result of one function call, namely `round 0.9`, as an argument to another function. (Think about what would happen if we did not put parentheses around `(round 0.9)`...how many arguments would we then be passing to `pluralize`?)

1.3.2 Importing Functions

So far we've only used basic operators and functions we wrote ourselves. Now let's expand our repertoire of functions by using one from an external module.

DEFINITION

A **module** is a named collection of Elm functions and other values.

The `String` module is one of the core modules that ships with Elm. Additional modules can be obtained from Elm's official package repository, copy-pasting code from elsewhere, or through a back-alley rendezvous with a shadowy figure known as Dr. Deciduous. Chapter 4 will cover how to do the former, but neither the author nor Manning Publications endorses obtaining Elm modules through a shadowy back-alley rendezvous.

Let's import the `String` module and try out two of its functions, `toLower` and `toUpper`.

Listing 1.9 Importing functions

```
> String.toLower "Why don't you make TEN louder?"
"why don't you make ten louder?"

> String.toUpperCase "These go to eleven."
"THESE GO TO ELEVEN."
```

Observant readers may note a striking resemblance between Elm's `String.toUpperCase` function and the `toUpperCase()` method one finds on JavaScript strings. This is the first example of a pattern we will encounter many times!

JavaScript has several ways of organizing string-related functionality: fields on a string, methods on a string, or methods on the `String` global itself.

In contrast, Elm strings have neither fields nor methods. The `String` module houses the standard set of string-related features, and exposes them in the form of plain old functions like `toLower` and `toUpper`.

Table 1.2 String Functionality Comparison

JavaScript	Elm
<code>"storm".length</code>	<code>String.length "storm"</code>
<code>"dredge".toUpperCase()</code>	<code>String.toUpperCase "dredge"</code>
<code>String.fromCharCode(something)</code>	<code>String.fromChar something</code>

Not only is this organizational pattern consistent within the `String` module, it's consistent across Elm. Want a standard date-related feature? Look no further than the functions in the `Date` module. Regular Expression functions? Hit up the `Regex` module.

Methods are never the answer in Elm; over here it's all vanilla functions, all the time.

TIP Complete documentation for `String`, `Date`, `Regex`, and other tasty modules can be found in the *core* section of the package.elm-lang.org website.

We'll learn more about modules in the coming chapters, including how to write our own!

USING STRING.FILTER TO FILTER OUT CHARACTERS

Another useful function in the `String` module is `filter`. It lets us filter out unwanted characters from a string, such as non-numeric digits from a phone number.

To do this, we must give `filter` a function which specifies which characters to keep. The function will take a single character as an argument and return `True` if we should keep that character or `False` if we should chuck it. Figure 1.4 illustrates using `String.filter` to remove dashes from a US telephone number.

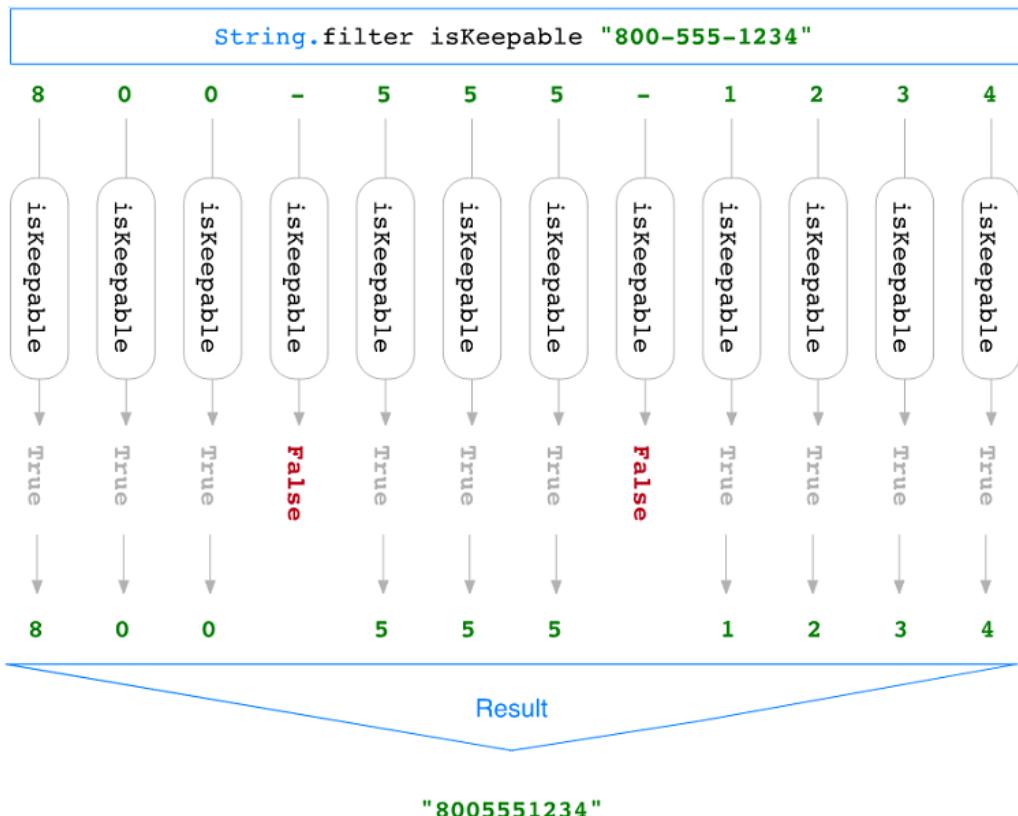


Figure 1.4 Using `String.filter` to remove dashes from a US phone number

As in JavaScript, Elm functions are first-class values that can be passed around just like any other value. This lets us provide `filter` with the function it expects by defining that function and then passing it in as a plain old argument.

Listing 1.10 Filtering with a named function

```
> isKeepable character = character /= '-' ①
<function>

> isKeepable 'z'
True

> isKeepable '-'
False

> String.filter isKeepable "800-555-1234" ②
"8005551234"
```

- ① A function describing which characters to keep
- ② Passing our function to `String.filter`

This normalizes telephone numbers splendidly. Alexander Graham Bell would be proud!

`String.filter` is one of the *higher-order functions* (that is, functions which accept other functions as arguments) that Elm uses to implement customizable logic like this.

1.3.3 Creating scope with let-expressions

Let's say we find ourselves removing dashes from phone numbers so often, we want to make a reusable function for it. We can do that with our trusty `isKeepable` function:

```
withoutDashes str = String.filter isKeepable str
```

This works, but in a larger Elm program, it might be annoying having `isKeepable` in the global scope like this. After all, its implementation is only useful to `withoutDashes`. Can we avoid globally reserving such a nicely self-documenting name?

Absolutely! We can scope `isKeepable` to the implementation of `withoutDashes` using a *let-expression*.

DEFINITION A *let-expression* adds locally scoped constants to an expression.

Figure 1.5 shows how we can implement `withoutDashes` using a single *let-expression*.

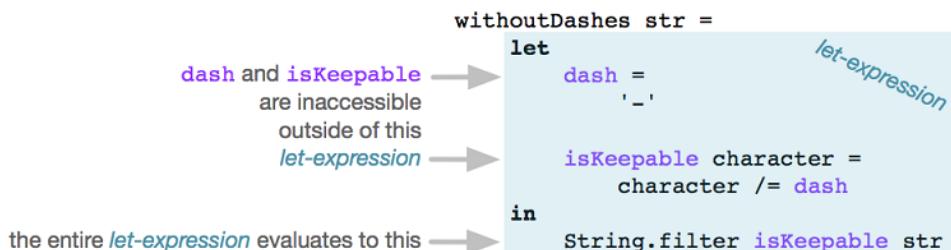


Figure 1.5 Anatomy of the wild *let-expression*

The above code does very nearly the same thing as entering the following in elm-repl:

```
> dash = '-'
> isKeepable character = character /= dash
> withoutDashes str = String.filter isKeepable str
```

In both versions, the implementation of `withoutDashes` boils down to `String.filter isKeepable str`. The only difference between the two is the scope of `dash` and `isKeepable`.

- In the elm-repl version above, `dash` and `isKeepable` are in the global scope.
- In Figure 1.5, `dash` and `isKeepable` are scoped locally to the *let-expression*.

Whenever you see a *let-expression*, you can mentally replace it with the part after its `in` keyword—in this case, `String.filter isKeepable str`. All the values between `let` and `in` are intermediate constants that are no longer in scope once the expression after `in` gets evaluated.

NOTE The indentation in Figure 1.5 is no accident! In a multiline *let-expression*, the `let` and `in` keywords must be at the same indentation level, and all other lines in the *let-expression* must be indented further than they are.

Anywhere you'd write a normal expression, you can swap in a *let-expression* instead. Because of this, you don't need to learn anything new to define locally-scoped constants inside function bodies, branches of *if-expressions*, or anywhere else.

Wherever you want some local scope, reach for a refreshing *let-expression*!

1.3.4 Anonymous Functions

Anonymous functions in Elm work the same way they do in JavaScript: like named functions but without the name.

Listing 1.11 Named and Anonymous Functions

```
function area(w, h) { return w * h; } ①
function(w, h) { return w * h; }        ②
area w h = w * h                    ③
\w h -> w * h                      ④
```

- ① JavaScript named function
- ② JavaScript anonymous function
- ③ Elm named function
- ④ Elm anonymous function

Elm's anonymous functions differ from its named functions in three ways.

1. They have no names

2. They begin with a \
3. Their parameters are followed by -> instead of =

Once defined, anonymous functions and named functions work the same way; you can always use the one in place of the other. For example, the following do exactly the same thing:

```
isKeepable char = char /= '-'
isKeepable = \char -> char /= '-'
```

Let's use an anonymous function to call `String.filter` in one line instead of two, then see if we can improve the business logic! For example, we can try using `Char.isDigit` to cast a wider net, filtering out any non-digit characters instead of just dashes.

Listing 1.12 Filtering with anonymous functions

```
> String.filter (\char -> char /= '-') "800-555-1234"
"8005551234"

> String.filter (\char -> char /= '-') "(800) 555-1234"
"(800) 5551234" ①

> import Char
> String.filter (\char -> Char.isDigit char) "(800) 555-1234"
"8005551234" ②

> String.filter Char.isDigit "(800) 555-1234" ③
"8005551234"
```

- ① Our simple filter fell short here
- ② Much better!
- ③ Refactor of previous approach

Anonymous functions are often used with higher-order functions like `String.filter`.

1.3.5 Operators

So far we've seen functions such as `String.filter`, as well as operators such as `++`, `-`, and `==`. How do operators and functions relate?

As it turns out, Elm's operators *are* functions! There are a few things that distinguish operators from normal functions:

- Operators must always accept exactly two arguments—no more, no fewer.
- Normal functions have names that begin with a letter. You typically call them by writing the name of the function followed by its arguments. This is *prefix-style* calling.
- Operators have names that contain neither letters nor numbers. You typically call them by writing the first argument, followed by the operator, followed by the second argument. This is *infix-style* calling.

Wrapping an operator in parentheses treats it as a normal function—*prefix-style* calling and all!

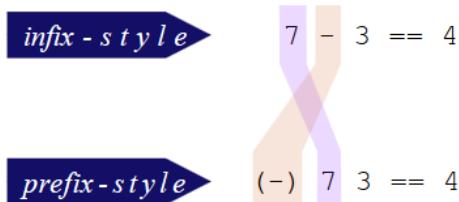


Figure 1.6 Calling the `(-)` operator in both *infix-style* and *prefix-style*

Let's play with some operators in `elm-repl`:

Listing 1.13 Operators are functions

```
> (/)
<function>

> divideBy = (/)
<function>

> 7 / 2      ①
3.5

> (/) 7 2    ②
3.5

> divideBy 7 2
3.5
```

- ① *infix-style* calling
- ② *prefix-style* calling

OPERATOR PRECEDENCE

Try entering an expression involving both arithmetic operators and `(==)` into `elm-repl`:

```
> 3 + 4 == 8 - 1
True : Bool
```

Now consider how we'd rewrite this expression in *prefix-style*:

```
> (==) ((+) 3 4) ((-) 8 1)
True : Bool
```

Notice anything about the order in which these operators appear? Reading the *infix-style* expression from left to right, you see `+` first, then `==`, and finally `-`. In the *prefix-style* expression, the order is different: first you see `==`, then `+`, and finally `-`. Why is this?

They get reordered because `(==)`, `(+)`, and `(-)` have different *precedence* values.

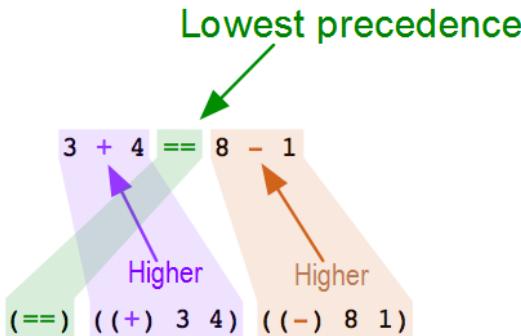


Figure 1.7 ($==$) gets evaluated after ($+$) and ($-$) because it has lower precedence

DEFINITION In any expression containing multiple operators, the operators with higher *precedence* get evaluated before those with lower precedence. This only applies to *infix-style* calls, as all *prefix-style* calls implicitly have the same precedence.

There isn't much formal documentation on operators' relative precedence values, but operators that appear in many programming languages (such as the $(==)$, $(+)$, and $(-)$ operators) tend to work similarly in Elm to how they do everywhere else.

We'll dig into operator precedence in greater depth in later chapters.

NORMAL FUNCTION CALLS HAVE TOP PRECEDENCE

Here are two ways to write the same thing:

```
> negate 1 + negate 5
-6

> (negate 1) + (negate 5)
-6
```

These two are equivalent because normal function calls have higher precedence than any operator. This means any time you want to pass the results of two normal function calls to an operator, you won't need to add any parentheses! You'll still get the result you wanted.

OPERATOR ASSOCIATIVITY

Besides precedence, the other factor that determines evaluation order for operators called in *infix-style* is whether the operators are *left-associative*, *right-associative*, or *non-associative*. Every operator is one of these.

An easy way to think about operator associativity is in terms of where the implied parentheses go. Infix expressions involving left-associative operators, such as arithmetic operators, have implied parentheses that cluster on the left:

Table 1.3 Implied parentheses for the (-) operator

Parentheses Shown	Expression	Result
none	10 - 6 - 3	1
assuming left-associative	((10 - 6) - 3)	1
assuming right-associative	(10 - (6 - 3))	7

If `(-)` were right-associative, $10 - 6 - 3$ would have parentheses clustering on the right, meaning it would evaluate to $(10 - (6 - 3))$ and the undesirable result of $10 - 6 - 3 == 7$. Good thing arithmetic operators are left-associative!

Non-associative operators cannot be chained together. For example, `foo == bar == baz` does not result in clustered parentheses, it results in an error!

In Chapter 7, we'll see how to define our own operators, including specifying their associativity and precedence values.

1.4 Collections

Elm's most basic collections are lists, records, and tuples. Each has varying degrees of similarity to JavaScript's arrays and objects, but one way in which they differ from JavaScript collections is in that Elm collections are always *immutable*.

DEFINITION An *immutable* value cannot be modified in any way once created.

This is in contrast to JavaScript, where some values (like strings and numbers) are immutable, but collections (like arrays and objects) can be mutated.

1.4.1 Lists

An Elm list has many similarities to a JavaScript array.

- You can create one with a square bracket literal, e.g. `["one fish", "two fish"]`
- You can ask for its first element
- You can ask for its length
- You can iterate over its elements in various ways

An Elm list does have some differences, though.

- It is immutable
- It has no fields or methods. You work with it using functions from the `List` module.
- Because it is a *linked list*, you can ask for its first element, but not for other individual elements. (If you need to ask for elements at various different positions, you can first convert from an Elm `List` to an Elm `Array`. We'll discuss Elm Arrays in Chapter 3.)
- You can combine it with another list using the `++` operator. In JavaScript this is done

with the `concat` method rather than an operator.

- All elements in an Elm list must have a consistent type. For example, it can be a “list of numbers” or a “list of strings,” but not a “list where strings and numbers intermingle.” (Making a list containing both strings and numbers involves first creating wrapper elements for them, using a feature called *union types* that we’ll cover in Chapter 3.)

Although Elm supports both (immutable) lists and (also immutable) arrays, lists are used far more often because they have better performance characteristics in typical Elm use cases. We’ll get more into the performance differences between linked lists and arrays in Chapter 10.

Here are some examples of how Elm lists and JavaScript arrays differ.

Table 1.4 Contrasting JavaScript Arrays and Elm Lists

JavaScript Array	Elm List
<code>[1, 2, 3].length</code>	<code>List.length [1, 2, 3]</code>
<code>["one fish", "two fish"][0]</code>	<code>List.head ["one fish", "two fish"]</code>
<code>["one fish", "two fish"][1]</code>	No arbitrary position-based element access
<code>[1, 2].concat([3, 4])</code>	<code>[1, 2] ++ [3, 4]</code>
<code>[1, 2].push(3)</code>	Cannot be modified; use e.g. <code>append</code> instead
<code>[1, "Paper", 3]</code>	All elements in a list must have a consistent type

Let’s focus on that last one. Why must all elements in an Elm list have a consistent type?

To understand how this requirement benefits us, let’s consider the `List.filter` function, which works like the `String.filter` function we used earlier.

Recall that `String.filter` takes a function which returns `True` when the given character should be kept, and `False` when it should be dropped. `List.filter` differs only in that the function you provide doesn’t necessarily receive characters—instead it receives elements from the list, whatever they may be.

Let’s see that in action. Quick! To `elm-repl`!

Listing 1.14 Filtering lists

```
> List.filter (\char -> char /= '-') [ 'z', '-', 'z' ]      ①
['Z','Z']

> List.filter (\str -> str /= "-") [ "ZZ", "-", "Top" ]    ②
["ZZ","Top"]

> import Char
> List.filter Char.isDigit [ '7', '-', '9' ]                  ③
['7','9']

> List.filter (\num -> num % 2 == 1) [ 1, 2, 3, 4, 5 ]     ④
[1,3,5]
```

- 1 Same function we passed to String.filter earlier
- 2 Strings instead of characters
- 3 Works just like with String.filter
- 4 Keep only the odd numbers

Here's how we would rewrite that last line of code in JavaScript:

```
[ 1, 2, 3, 4, 5 ].filter(function(num) { return num % 2 === 1; })
```

This looks straightforward enough, but JavaScript arrays permit inconsistent element types. Without looking it up, can you guess what happens if we change it to the following?

```
[ 1, "2", "cat", 4, "5", "" ].filter(function(num) { return num % 2 === 1; })
```

Will it crash? Will it happily return numbers? What about strings? It's a bit of a head-scratcher.

Because Elm requires consistent element types, this is a no-brainer: in Elm it would be an error. Even better, it would be an error at build time—meaning you can rest easy knowing whatever surprises would result from executing this code will not inflict pain on your users. Requiring consistent element types means all lists in Elm guarantee this level of predictability.

By the way, the above `filter()` call successfully returns `[1, "5"]`. (Like, *duh*, right?)

1.4.2 Records

We've now seen how JavaScript's mutable arrays resemble Elm's immutable lists. In a similar vein, JavaScript's mutable objects resemble Elm's immutable *records*.

DEFINITION A *record* is a collection of named fields, each with an associated value.

Whereas array and list literals between the two languages are syntactically identical, where JavaScript object literals use `:` to separate fields and values, Elm record literals use `=` instead.

Let's get a taste for some of their other differences.

JavaScript Object	Elm Record
<code>{ name: "Li", cats: 2 }</code>	<code>{ name = "Li", cats = 2 }</code>
<code>({ name: "Li", cats: 2 }).toString()</code>	<code>toString { name = "Li", cats = 2 }</code>
<code>({ name: "Li", cats: 2 }).cats</code>	<code>({ name = "Li", cats = 2 }).cats</code>
<code>({ name: "Li", cats: 2 })["cats"]</code>	Fields can only be accessed directly, using a dot
<code>({ name: "Li", cats: 2 }).cats = 3</code>	Cannot be modified. (New cat? New record!)
<code>{ NAME: "Li", CATS: 2 }</code>	Fields have the same naming rules as constants
<code>({ name: "Li", cats: 2 }).__proto__</code>	No prepackaged fields, only the ones you define
<code>Object.keys({ name: "Li", cats: 5 })</code>	No listing of field names is available on demand
<code>Object.prototype</code>	Records have no concept of inheritance

Wow—compared to objects, records sure don't do much! It's like all they do is sit around holding onto the data we gave them. (Yep.) Personally I've found Elm's records a welcome reprieve from the intricacies of JavaScript's `this` keyword.

RECORD UPDATES

Record updates let us concisely obtain a new record by copying the old one and changing only the specified values. (As we will see in Chapter 10, behind the scenes Elm does not actually copy the entire record—that would be slow!—but rather only the parts that will be different.)

Let's use this technique to represent someone obtaining an extra cat, going from `{ name = "Li", cats = 2 }` to `{ name = "Li", cats = 3 }` by way of a record update.

Listing 1.15 Record updates

```
> catLover = { name = "Li", cats = 2 }
{ name = "Li", cats = 2 }

> catLover
{ name = "Li", cats = 2 }

> withThirdCat = { catLover | cats = 3 }          ①
{ name = "Li", cats = 3 }

> withThirdCat
{ name = "Li", cats = 3 }

> catLover          ②
{ name = "Li", cats = 2 }  ②

> { catLover | cats = 88, name = "LORD OF CATS" } ③
{ name = "LORD OF CATS", cats = 88 }
```

① Record update syntax

② Original record unmodified!

③ Update multiple fields (order doesn't matter)

Record updates let us represent this incremental evolution without mutating our records or recreating them from scratch. In Chapter 2 we'll represent our application state with a record, and use record updates to make changes based on user interaction.

1.4.3 Tuples

Lists let us represent collections of **varying size**, whose elements share a **consistent type**. Records let us represent collections of **fixed fields**, but whose corresponding values may have **varied types**.

Tuples introduce no new capabilities to this mix, as there is nothing a tuple can do that a record couldn't. Compared to records, though, what tuples bring to the party is conciseness.

DEFINITION A *tuple* is a record-like value whose fields are accessed by position rather than by name.

In other words, tuples are for when you want a record, but don't want to bother naming its fields. They are often used for things like key-value pairs where writing out `{ key = "foo", value = "bar" }` would add verbosity but not much clarity.

Let's try some out!

Listing 1.17 Using Tuples

```
> ( "Tech", 9 )
("Tech",9)

> Tuple.first ( "Tech", 9 ) ①
"Tech"

> Tuple.second ( "Tech", 9 ) ②
9
```

- ① Return first element (only works on 2-element tuples)
- ② Return second element (only works on 2-element tuples)

You can only use the `Tuple.first` and `Tuple.second` functions on tuples that contain two elements. If they have more than two, you can use *tuple destructuring* to extract their values.

DEFINITION *Tuple destructuring* extracts the values inside a tuple and assigns them to constants in the current scope.

Let's use tuple destructuring to implement a function that takes a tuple of three elements.

Listing 1.18 Tuple Destructuring

```
> multiply3d ( x, y, z ) = x * y * z ①
<function>

> multiply3d ( 6, 7, 2 )
84

> multiply2d someTuple = let ( x, y ) = someTuple in x * y ②
<function>
```

- ① Destructuring a tuple into three constants: `x`, `y`, and `z`
- ② Destructuring a tuple inside a `let`-expression

As demonstrated in Listing 1.18, once you have named the values inside the tuple, you can use them just like you would any other constant.

TIP Mind the difference between a tuple and a parenthetical function call! `(foo, bar)` is a tuple, whereas `(foo bar)` is a call to the `foo` function passing `bar` as an argument. A simple mnemonic to remember the difference is “comma means tuple.”

Table 1.5 Comparing Lists, Records, and Tuples

List	Record	Tuple
Variable Length	Fixed Length	Fixed Length
Can Iterate Over	Cannot Iterate Over	Cannot Iterate Over
No Names	Named Fields	No Names
Immutable	Immutable	Immutable

Since any tuples can be represented (more verbosely) using a record instead, it's often better to refactor long tuples—say, with more than three elements—into records. Choose tuples or records based on whichever would yield more readable code; their performance characteristics are equivalent.

1.5 Summary

We're off to a fantastic start! First we discussed some of the toughest problems Web programmers face: crashing is too easy in JavaScript, and maintenance is too error-prone. Then we learned how Elm addresses these problems, with a design that prioritizes maintainability and a helpful compiler that catches would-be runtime exceptions before they can cause user pain. From there we dove in and wrote our first Elm code in `elm-repl`.

Here is a brief review of things we covered along the way.

- The `++` operator combines strings and lists, whereas the `+` operator is for addition only.
- Double quotes refer to strings. Single quotes refer to individual UTF-8 characters.
- *let-expressions* introduce scoped constants to an expression.
- There is no concept of “truthiness” in Elm, just `True` and `False`.
- `if foo /= bar then "different" else "same"` is an *if-expression*. Like JavaScript ternaries, *if-expressions* require an `else` branch and always evaluate to a value.
- Lists like `[3, 1, 4]` are immutable. Their elements must share a consistent type.
- `List.filter (\num -> num < 0) numbersList` returns a list containing all the negative numbers in the original `numbersList`
- `catLover = { name = "Li", cats = 2 }` assigns a record to the constant `catLover`. Once assigned, constants cannot be reassigned.
- `{ catLover | cats = 3 }` returns a new record that is the same as the `catLover` record, except the `cats` value is now 3.
- `(foo, bar)` destructures a tuple such as `(2, 3)`. In this example, `foo` would be 2 and `bar` would be 3.

Table 1.6 summarizes some of the differences between JavaScript and Elm.

Table 1.6 Differences between JavaScript and Elm

JavaScript	Elm
// This is an inline comment	-- This is an inline comment
/* This is a block comment */	{- This is a block comment -}
true && false	True && False
"Ahoy, " + "World!"	"Ahoy, " ++ "World!"
"A spade" === "A spade"	"A spade" == "A spade"
"Calvin" !== "Hobbes"	"Calvin" /= "Hobbes"
Math.pow(2, 11)	2 ^ 11
Math.trunc(-49 / 10)	-49 // 10
n % 2 === 1 ? "odd" : "even"	if n % 2 == 1 then "odd" else "even"
nums.filter(function(n) { ... })	List.filter (\n -> n % 2 == 1) nums
function pluralize(s, p, c) { ... }	pluralize singular plural count = ...

We also learned about several differences between normal functions and operators:

Table 1.7 Differences between normal functions and operators

Function	How to identify one	Calling style	Examples
Normal	Name begins with a letter	<i>prefix-style</i>	negate, not, toString
Operator	Name has no letters or numbers	<i>infix-style</i>	(++), (*), (==)

In Chapter 2 we'll expand on what we've learned here to create a working Elm application.

Let's go build something!

2

Your First Elm Application

This chapter covers

- Declaratively rendering a page
- Managing state with Model-View-Update
- Handling user interaction

Elm applications are built to last. They have a reputation for being scalable, easy to refactor, and difficult to crash unexpectedly. Since JavaScript applications have...erm...a different reputation, it stands to reason that Elm must be doing things differently. And so it is!

Where each line of code in a JavaScript application can potentially result in a change or effect—like “*update that text!*” or “*send this to the server!*”—the code in an Elm application builds up a **description** of what the program should do in response to various inputs. Elm’s compiler translates this description into the appropriate JavaScript commands for the browser to run at the appropriate times, and the end user may have no idea Elm was involved at all.

In this chapter we’ll build our first Elm application: *Photo Groove*, a simple photo browsing Web app where users select thumbnails to view larger versions. We’ll create a user interface using declarative rendering and manage state using the Elm Architecture. By the end, we will have a fully functioning application—and a code base we can build on for the rest of the book!

2.1 Rendering a Page

Since the very early days of the Web, browsers have been translating HTML markup into a Document Object Model (or DOM for short) which represents the structure of the current page. The DOM consists of *DOM nodes*, and it’s only by changing these nodes that Web applications can modify the current page on the fly.

In this chapter we'll work with the two most common types of DOM nodes: *elements* and *text nodes*.

- *Elements* have a `tagName` (such as "div" or "img"), and may have child DOM nodes.
- *Text nodes* have a `textContent` property instead of a `tagName`, and are childless.

As Figure 2.1 shows, elements and text nodes can freely intermingle inside the DOM.

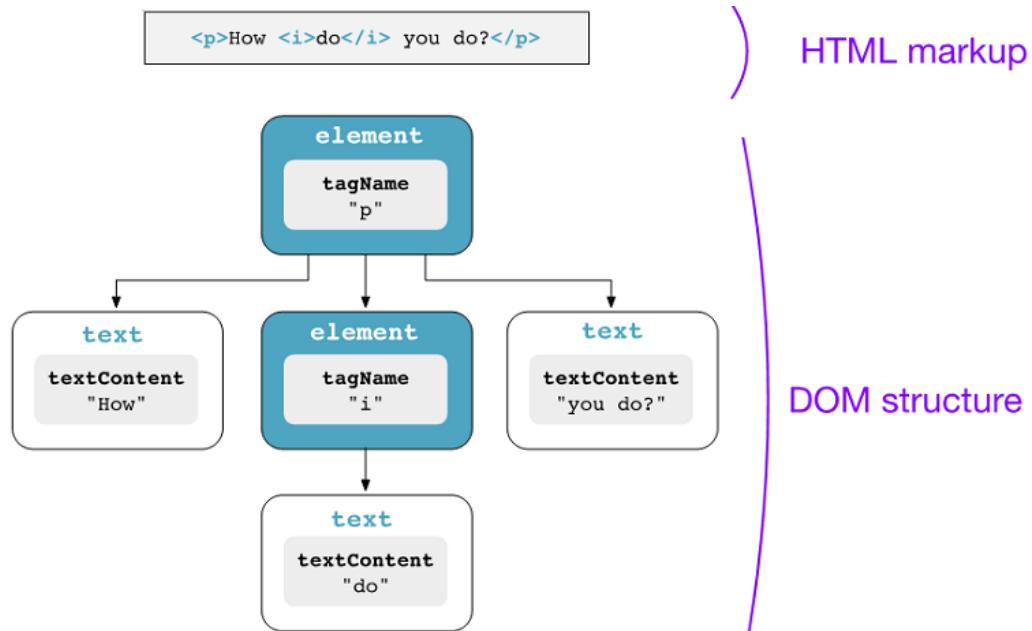


Figure 2.1 Intermingling element DOM nodes and text nodes

Here we've pulled back the curtain on the markup `<p>How <i>do</i> you do?</p>` to see that despite its two element tags—namely `<p>` and `<i>`—we are actually working with five DOM nodes here! The other three are not elements, but rather text nodes.

2.1.1 Describing a page using the Html Module

When describing how a page looks in Elm, we do not write markup. Instead, we call functions to create representations of individual DOM nodes. The most flexible of these functions is called `node`, and as Figure 2.2 shows, its arguments line up neatly with the analogous markup.

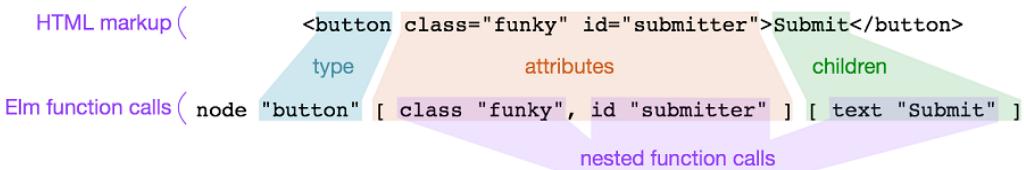


Figure 2.2 Representing a button using HTML markup (top) and Elm's node function (bottom)

There are four function calls in this line of Elm code. Can you spot them?

1. A call to the `node` function, passing three arguments: the string "button", a list of attributes, and a list of child nodes
2. A call to the `class` function, passing "funky"
3. A call to the `id` function passing "submitter"
4. A call to the `text` function passing "Submit"

These are plain old Elm functions. Each returns a representation of some portion of the DOM: a `button` element, a text node, and some `class` and `id` attributes. You can call these functions anytime you like, and pass their return values to other functions as normal.

In Elm we usually refer to a “a virtual DOM node” as “`Html`” for short. This name comes from the `Html` module, which provides a plethora of intuitively-named functions that let you avoid calling `node` directly. For example, the `Html` module’s `img` function is shorthand for calling `node` passing “`img`” as the first argument. The following two lines are equivalent:

```
node "img" [ src "logo.png" ] []
     img [ src "logo.png" ] []
```

It’s best practice to use helper functions like `img` as much as possible, and to fall back on `node` only in cases where no helper function is available. (For example, you may notice that there is no helper for the deprecated `<blink>` element. I’m not saying you should call `node "blink" [] [text "<BLINK> LIVES AGAIN"]`, but I’m not *not* saying it either.)

RENDERING A PAGE

Let’s use what we’ve learned to render our first page with Elm!

For the rest of this chapter we’ll be building an application called Photo Groove. Eventually we’ll add features like searching and viewing larger versions, but first we need to render a basic page that says “Photo Groove” across the top, with some thumbnail images below.

Since our output is visual this time, `elm-repl` won’t get us very far. Instead, let’s bring up the *Try Elm* website at <http://elm-lang.org/try> and enter the following into the left-hand pane.

Listing 2.2 Building a view in Try Elm

```
import Html exposing (div, h1, img, text)    ①
import Html.Attributes exposing (...)          ①
```

```

view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]      ②
        , div [ id "thumbnails" ]
            [ img [ src "http://elm-in-action.com/1.jpeg" ] []
            , img [ src "http://elm-in-action.com/2.jpeg" ] []
            , img [ src "http://elm-in-action.com/3.jpeg" ] [] ④
            ]
    ] ③

main = ⑤
    view "no model yet"

```

- ① We'll discuss "exposing" later
- ② h1 element with an empty attributes list
- ③ Put commas at the start of the line
- ④ img element with an empty children list
- ⑤ We'll discuss "main" later

Press **Compile**. The words "Photo Groove" should appear in the right-hand pane, followed by three images. (Depending on the size of your browser, the images may not be side-by-side.)

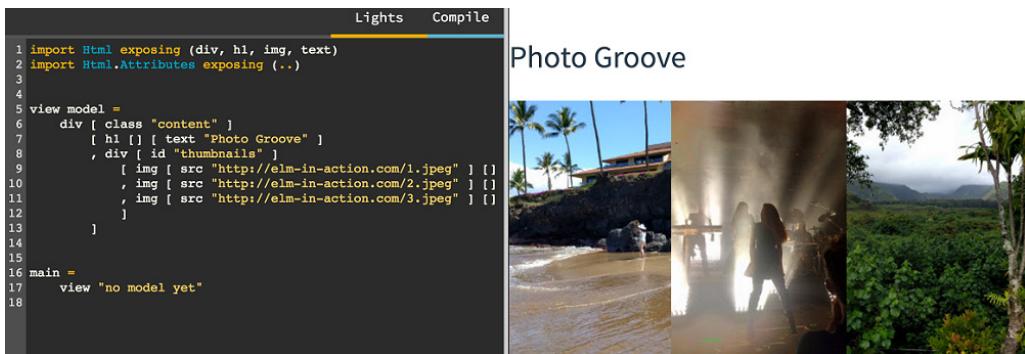


Figure 2.3 The Try Elm website

Congratulations! You've rendered your first user interface in Elm.

Figure 2.4 shows the DOM structure of the interface we just rendered.

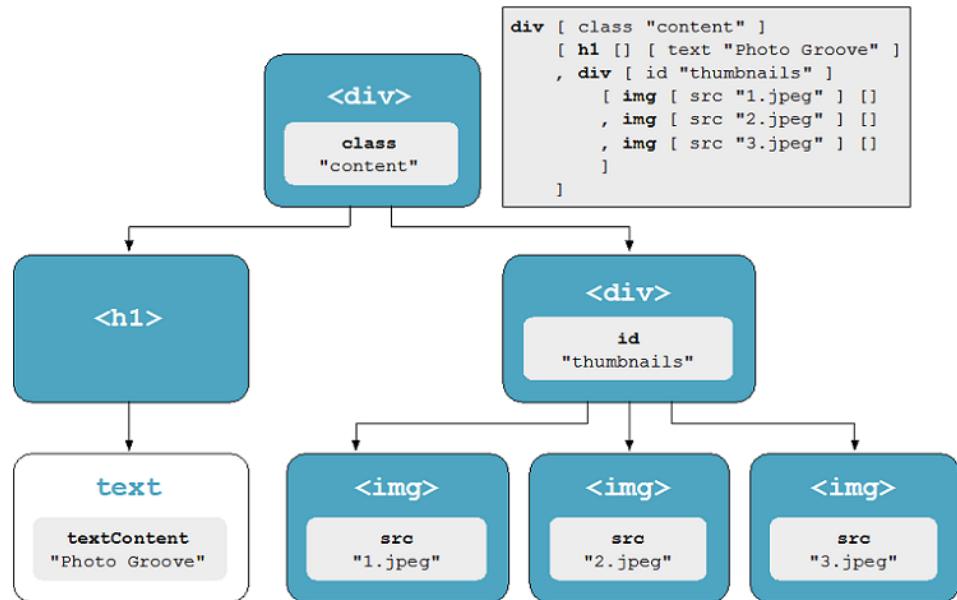


Figure 2.4 The DOM structure of your first Elm user interface

Notice how the helper functions that create elements—in this case, `div`, `h1`, and `img`—take exactly two arguments in all cases:

1. **A list of attributes.** If an element has no attributes, we pass `[]` like so:
`h1 [] [text "Photo Groove"]`
2. **A list of child DOM nodes.** If an element has no children, we pass `[]` like so:
`img [src "1.jpeg"] []`

If an element has neither attributes nor children? In that case we pass `[] []` like so:

```
br [] []
```

This pattern is consistent throughout the `Html` module, and it's worth following if you ever decide to make a custom element helper of your own using the `node` function.

IMPORTING UNQUALIFIED VALUES WITH ‘EXPOSING’

We've now used several functions from the `Html` module, but we wrote them in a different style from the way we did in Chapter 1. Back then we wrote functions like `String.filter` in a *qualified* style—that is, we included the `String` module's name right there in the function call. With the `Html` module's functions, we used an *unqualified* style—we wrote `div` instead of `Html.div`, we wrote `h1` instead of `Html.h1`, and so forth.

We could do this because we used `exposing` when we imported the `Html` module:

```
import Html exposing (div, h1, img, text)
```

This line of code both imports the `Html` module so we can use its contents, and also brings `Html.div`, `Html.h1`, `Html.img`, and `Html.text` into the global scope. That lets us refer to them as `div`, `h1`, `img`, and `text` without the prefix of `Html`.

We could have achieved a similar result by assigning them to constants:

```
import Html

div = Html.div
h1 = Html.h1
img = Html.img
text = Html.text
```

However, since this pile of code can be replaced by a single line—`import Html exposing (div, h1, img, text)`—it's normal to use `exposing` for this purpose instead.

EXPOSING EVERYTHING WITH(..)

When we imported the `Html` module, we listed exactly which values we wanted to expose: `div`, `h1`, `img`, and `text`. For the `Html.Attributes` module, we wrote this instead.

```
import Html.Attributes exposing(..)
```

Using `exposing(..)` means “expose everything,” which lets us use every value in the `Html` module in an unqualified style. Let’s change our first `import` to use `exposing(..)` instead:

```
import Html exposing(..)
```

Now we won’t need to extend the list of `div`, `h1`, `img`, and `text` whenever we want to use a new element type. Everything the `Html` module has to offer is now in our global scope!

When to use qualified over unqualified imports

In Chapter 1 we wrote out `String.toUpperCase` and `List.filter`, instead of `toUpperCase` and `filter`. Here we’re doing the opposite, writing `img` and `div` instead of `Html.img` and `Html.div`.

This begs the question: when is it a good idea to use the qualified style (with the module name prefixed) over the unqualified style? The unqualified style is more concise, so why not use `exposing(..)` every time?

There are two primary downsides to unqualified imports. One is that unqualified names can become ambiguous. Try this in `elm-repl`:

```
> import String exposing(..)
> import List exposing(..)
> filter
```

You’ll get an error saying that `filter` is ambiguous. After importing and exposing both `String.filter` and `List.filter`, it’s no longer clear which of the two you meant! (In cases like this you can still use the qualified style to resolve the ambiguity, so if you now put `String.filter` or `List.filter` into `elm-repl`, they will still work as normal.)

Unqualified imports are also less self-documenting. Suppose you come across code that says `partition foo bar`, and you’ve never seen `partition` before. Naturally you wonder: “How can I find out what `partition` does? Is it

defined in this file?" You search through the file and can't find it, so it must come from an import. You scroll up to the imports and discover a long list of exposing (...) declarations. Argh! partition could be in any of those!

This could take awhile...

Suppose instead you see the code `List.partition foo bar`. You want to know what `List.partition` does, so you bring up the documentation for the `List` module on <http://package.elm-lang.org>. You learn about `List.partition`, then get on with your day!

Scenarios like this are why it's best practice to write things in a *qualified* way by default.

Still, sometimes there's a good reason to prefer the unqualified style—like how unqualified `Html` functions are designed to resemble HTML markup. In these cases, it's best to limit yourself to one exposing (...) (or perhaps one "family" of them, such as `Html` and `Html.Attributes`) per file. This way if you encounter an unfamiliar function of mysterious origin, you'll have the fewest modules to hunt through to find its documentation!

2.1.2 Building a Project

Now that we've gotten something on the screen, let's add some styles! We could style our elements inline using the `Html.Attributes.style` attribute, but instead we'll organize things the traditional way: with our CSS declarations in a separate file.

The only way to get multiple files involved in the same page is to give a browser some HTML markup, so our first step in the process of styling our application will be to create an `index.html` file. Since *Try Elm* doesn't let you bring your own HTML file to the party, this means we'll need to start building our project locally.

1. Make a new directory called `PhotoGroove` and open it in a terminal.
2. Create a file called `PhotoGroove.elm` inside this directory.
3. Copy the *Try Elm* code from Listing 2.2 into `PhotoGroove.elm`.
4. Add the line `module PhotoGroove exposing(..)` to the top of `PhotoGroove.elm`.
5. Change `import Html exposing (div, h1, img, text)` to use `exposing(..)`.
6. When the dust settles, we should have a `PhotoGroove.elm` file that looks like this:

Listing 2.3 PhotoGroove.elm

```
module PhotoGroove exposing(..) ①

import Html exposing(..)          ②
import Html.Attributes exposing(..)

view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , div [ id "thumbnails" ]
            [ img [ src "http://elm-in-action.com/1.jpeg" ] []
            , img [ src "http://elm-in-action.com/2.jpeg" ] []
            , img [ src "http://elm-in-action.com/3.jpeg" ] []
            ]
        ]
main = ③
    view "no model yet"
```

- 1 "This file describes a module called PhotoGroove"
- 2 Remember: this uses "exposing(..)" now!
- 3 This will be the application's entry point

Why Commas in Front?

When writing a multi-line literal in JavaScript, the usual convention is to put commas at the end of each line. Consider the following code:

```
rules = [
    rule("Do not talk about Sandwich Club."),
    rule("Do NOT talk about Sandwich Club.")
    rule("No eating in the common area.")
]
```

Did you spot the mistake? There's a comma missing after the second call to `rule`, meaning this is not syntactically valid JavaScript. Running this code will result in a `SyntaxError`.

Now consider the equivalent Elm code, with the same missing comma:

```
rules = [
    rule "Do not talk about Sandwich Club.",
    rule "Do NOT talk about Sandwich Club."
    rule "No eating in the common area."
]
```

The mistake is just as easy to overlook, but harder to fix because this is syntactically valid Elm code—just not the code you intended to write!

The missing comma means the above code is essentially equivalent to the following:

```
rules = [
    (rule "Do not..."),
    (rule "Do NOT..." rule "No eating...")
]
```

Instead of calling `rule` three times, each time with one argument, here the second call to `rule` is receiving three arguments—and there is no third call! This means instead of the syntax error JavaScript gave you, you'll get a seemingly nonsensical error about arguments.

Now try to make this mistake when writing in a commas-first style:

```
rules =
  [ rule "Do not talk about Sandwich Club."
    rule "Do NOT talk about Sandwich Club."
    , rule "No eating in the common area."
  ]
```

This style makes it blindingly obvious that a comma is missing. Now we don't even need to compile our code to identify the problem!

It may feel different at first, but the commas-first style gives you one less potential error to worry about once you get used to it.

DECLARING THE PHOTOGROOVE MODULE

By writing `module PhotoGroove exposing (...)` at the top of our `PhotoGroove.elm` file, we defined a new module of our own. This means other files in our project can now import this `PhotoGroove` module just like they would the `String` or `Html` modules, for example like so:

```
import PhotoGroove exposing (main)
```

Because we wrote `exposing (...)` after `module PhotoGroove`, we are exposing all of our top-level values—namely `main` and `view`—for other modules to import.

If we wanted to expose `main` but leave `view` invisible to other modules, we could have written `exposing (main)` instead. In such a case, another module that imported `PhotoGroove` would get an error if it tried to access `PhotoGroove.view`. Only exposed values can be accessed by other modules!

INSTALLING PACKAGE DEPENDENCIES

At this point we have a `PhotoGroove.elm` file, but it's missing a dependency. All this fancy rendering logic we've been writing comes from the `Html` module, but unlike the modules we've been using so far—like `String` and `List`—the `Html` module does not come preinstalled. To access it, we need to install the *package* that contains the `Html` module.

DEFINITION An Elm *package* is an installable collection of modules.

The `elm-package` utility downloads and installs packages when you give it the name of the package you want. Package names consist of a username followed by a / and then the package name; in this case, the package we seek is named `elm-lang/html`.

Let's use `elm-package` to get some `Html` going. Run this command in the terminal:

```
elm-package install elm-lang/html
```

You should see something like this:

```
Some new packages are needed. Here is the upgrade plan.
```

(it lists the packages here)

```
Do you approve of this plan? (Y/n)
```

Answer `y`, and you should shortly see the text `Packages configured successfully!` You should now have a file called `elm-package.json` in your current directory, which `elm-package` will have created for you automatically. We'll dig into `elm-package.json` in future chapters.

COMPILING TO JAVASCRIPT AND RUNNING THE APPLICATION

Now we're ready to compile our Elm code into JavaScript. Run this in the terminal:

```
elm-make PhotoGroove.elm --output elm.js
```

This will compile our `PhotoGroove.elm` file into a JavaScript file we can give to a browser. (That generated JavaScript file will be called `elm.js`, because we passed `--output elm.js` to `elm-make`.) Now all we need is a HTML file that will load up our compiled `elm.js` file!

Make a file called `index.html` and put the following inside it.

Listing 2.4 index.html

```
<!doctype html>
<html>
    <head>
        <style>
            body { background-color: rgb(44, 44, 44); color: white; }
            img { border: 1px solid white; margin: 5px; }
            .large { width: 500px; float: right; }
            .selected { margin: 0; border: 6px solid #60b5cc; }
            .content { margin: 40px auto; width: 960px; }
            #thumbnails { width: 440px; float: left; }
            h1 { font-family: Verdana; color: #60b5cc; }
        </style>
    </head>

    <body>
        <div id="elm-area"></div> ①

        <script src="elm.js"></script> ②
        <script>
            Elm.PhotoGroove.embed(document.getElementById("elm-area"));
        </script>
    </body>
</html>
```

- ① Our Elm application will render into this div
- ② `PhotoGroove.elm` will get compiled into `elm.js`
- ③ The Elm object comes from `elm.js`

The markup we put in this file covers things like:

- The standard `<!doctype>`, `<html>`, and `<body>` tags
- Importing the `elm.js` file we generated with `elm-make`
- Whatever `<head>` inclusions we need—styles, metadata, `<title>`, and so on

The line `Elm.PhotoGroove.embed(document.getElementById("elm-area"));` starts our Elm code running in the `<div id="elm-area"></div>` element we set up earlier in `index.html`. Let's break down what happens when our Elm code starts running.

THE ELM RUNTIME AND 'MAIN'

When Elm compiles your code into JavaScript, it includes an extra bit of JavaScript known as the Elm Runtime. The Elm Runtime is behind-the-scenes logic that silently handles things like:

- Adding and removing event listeners for any events your code depends on
- Efficiently scheduling tasks like HTTP requests and DOM updates

- Storing and managing application state

When we called `Elm.PhotoGroove.embed`, we told the Elm Runtime to open the module called `PhotoGroove` and use its exposed `main` value as the entry point for the application. If we did not have a module called `PhotoGroove`, or if it did not expose a `main` value, we'd get an error.

This means when the browser runs our compiled code, view "no model yet" will be the first line of code executed, because that's what we assigned to `main`. If we renamed the `PhotoGroove` module to `CubeDraft`, we'd have to call `Elm.CubeDraft.embed` instead, but otherwise everything would still work. If the `CubeDraft` module did not expose a value called `main`, however, the application would not start. There's no renaming `main`!

If you open `index.html`, you should see our application displaying as it does in Figure 2.5.

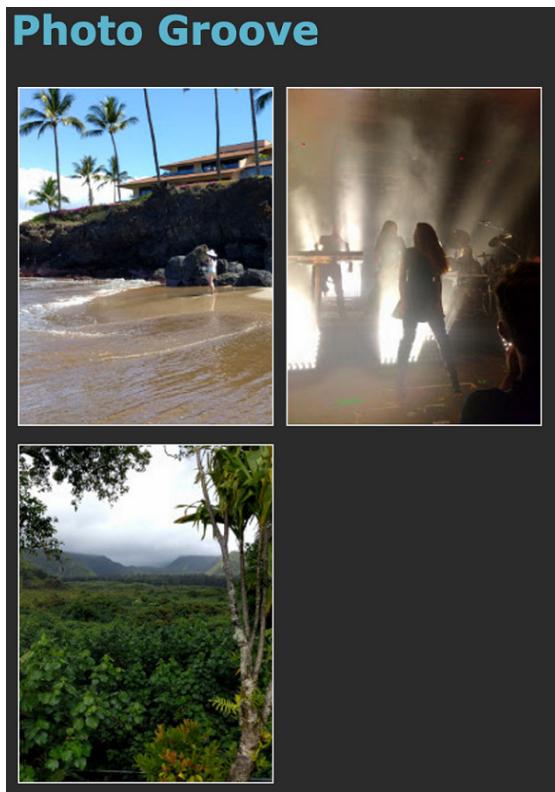


Figure 2.5 Rendering the application

Fantastic! Next we'll make it interactive.

2.2 Handling User Input with the Elm Architecture

So far we haven't had much data flowing through our application. Okay, really we haven't had *any*—we just generated some `Html` and rendered it. That will soon change, as we're about to start handling user input! This brings us to a common question: how will we keep data flow manageable as our code scales?

JavaScript offers a staggering selection of data flow architectures to choose from, but Elm has just one. It's called the Elm Architecture, and the Elm Runtime is optimized for applications that follow it. We'll learn about the Elm Architecture as we add interactivity to Photo Groove.

Figure 2.6 shows a preview of the architecture we'll be building toward in this chapter. Don't worry if this does not make sense yet! We will get there, one step at a time.

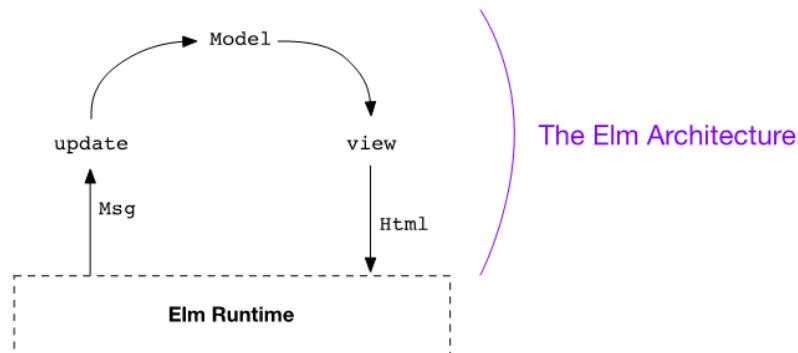


Figure 2.6 The Elm Runtime uses the Elm Architecture to manage data flow

Let's begin where data flow naturally begins: with application state as a whole.

2.2.1 Representing Application State with a Model

Back in the Wild West days of the Web, it was common to store application state primarily in the DOM itself. Is that menu expanded or collapsed? Check whether one of its DOM nodes has `class="expanded"` or `class="collapsed"`. Need to know what value a user has selected in a dropdown? Query it out of the DOM at the last possible instant.

This approach turned out not to scale very well, especially as applications grew more complex and unit testing became increasingly important. Today it's common practice to store application state completely outside the DOM, and to propagate changes from that independent state over to the DOM as necessary. This is what we do in the Elm Architecture.

DECLARING A MODEL

We're going to store our application state separately from the DOM, and will refer to that state as our *model*.

DEFINITION A *model* stores the current state of an application. Any value necessary to render a user interface should be stored in its model.

Remember how earlier we wrote this code?

```
main =
  view "no model yet"
```

Let's replace this code with the contents of Listing 2.5.

Listing 2.5 Adding a Model

```
initialModel =
  [ { url = "1.jpeg" } ①
  , { url = "2.jpeg" } ①
  , { url = "3.jpeg" } ①
  ]
```



```
main =
  view initialModel ②
```

- ① We'll add more fields beyond url later
- ② Pass our new initialModel record to view

Excellent! Now we have an initial model to work with. So far it contains a list of photos, each of which is represented by a record containing a `url` string.

WRITING A VIEW FUNCTION

Next we'll render a thumbnail for each photo in our list. A typical Elm application does this through a *view function*, which describes how the DOM should look based on its arguments.

At the top of a typical Elm application is a single `view` function, which accepts our current model as an argument and then returns some `Html`. The Elm Runtime takes the `Html` returned by this `view` function and renders it.

By pure coincidence, we've already written just such a `view` function—it's the function we had the foresight to name `view`. Unfortunately, our current `view` implementation ignores the `model` argument it receives, which means changing our model won't result in a visible change to the end user. Let's fix that! `view` should base its return value on its `model` argument.

It'll be easier to do this if we first write a separate `viewThumbnail` function, which renders a single thumbnail as `Html`. We can set the stage for that design by replacing our `view` implementation with the following:

Listing 2.6 Splitting out viewThumbnail

```
urlPrefix = ①
  "http://elm-in-action.com/" ①

view model =
  div [ class "content" ]
    [ h1 [] [ text "Photo Groove" ]
    , div [ id "thumbnails" ] []
    ]

viewThumbnail thumbnail =
  img [ src (urlPrefix ++ thumbnail.url) ] [] ②
```

- ① We'll prepend this to strings like "1.jpeg"
 ② Prepend urlPrefix to get a complete URL like "http://elm-in-action.com/1.jpeg"

Figure 2.7 illustrates how our current `model` and `view` connect to the Elm Runtime.

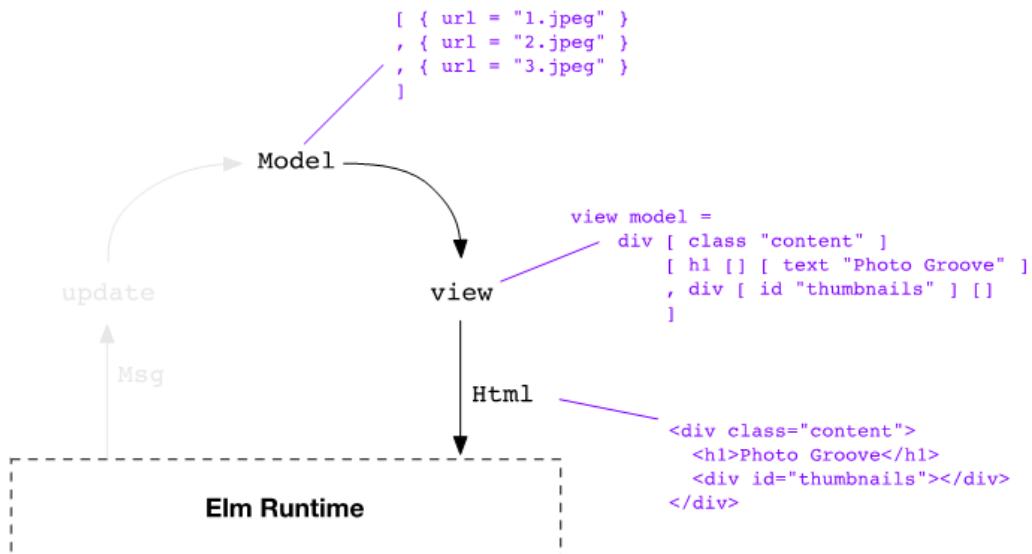


Figure 2.7 Model and view connecting with the Elm Runtime

Next, we'll iterate over our list of photo records and call `viewThumbnail` on each one, in order to translate it from a dusty old record to a vibrant and inspiring `img`.

Fortunately, the `List.map` function does exactly this!

LIST.MAP

`List.map` is another higher-order function similar to the `List.filter` function we used in Chapter 1. You pass `List.map` a translation function and a list, and it runs that translation function on each value in the list. Once that's done, `List.map` returns a new list containing the translated values.

Take a look at Figure 2.8 to see `List.map` do its thing for `viewThumbnail`.



Figure 2.8 Using `List.map` to transform photo records into `img` nodes

Since `div` is just a plain Elm function that accepts two lists as arguments—first a list of attributes, followed by a list of child nodes—we can swap out our entire hardcoded list of child `img` nodes with a single call to `List.map`! Let's go ahead and do that now.

```

view model =
  div [ class "content" ]
    [ h1 [] [ text "Photo Groove" ]
      , div [ id "thumbnails" ] (List.map viewThumbnail model)
    ]

viewThumbnail thumbnail =
  img [ src (urlPrefix ++ thumbnail.url) ] []
  
```

If you run `elm-make PhotoGroove.elm --output elm.js` again to recompile this code, you should see the same result as before. The difference is that now we have a more flexible internal representation, setting us up to add interactivity in a way that was impossible before we connected `model` and `view`.

EXPANDING THE MODEL

Now let's add a feature: when the user clicks on a thumbnail, it will become selected—indicated by a blue border surrounding it—and we'll display a larger version of it beside the thumbnails.

To do this, we first need to store which thumbnail is selected. That means we'll want to convert our model from a list to a record, so we can store both the list of photos and the current `selectedUrl` value at the same time.

Listing 2.7 Converting the model to a record

```
initialModel =
  { photos =
    [ { url = "1.jpeg" }
    , { url = "2.jpeg" }
    , { url = "3.jpeg" }
    ]
  , selectedUrl = "1.jpeg" ①
  }
```

① Select the first photo by default

Next let's update `viewThumbnail` to display the blue border for the selected thumbnail.

That's easier said than done! Being a lowly helper function, `viewThumbnail` has no way to access the model—so it can't know the current value of `selectedUrl`. But without knowing which thumbnail is selected, how can it possibly know whether to return a selected or unselected `img`?

It can't! We'll have to pass that information along from `view` to `viewThumbnail`.

Let's rectify this situation by passing `selectedUrl` into `viewThumbnail` as an additional argument. Armed with that knowledge, it can situationally return an `img` with the "selected" class—which our CSS has already styled to display with a blue border—if the `url` of the given thumbnail matches `selectedUrl`.

```
viewThumbnail selectedUrl thumbnail =
  if selectedUrl == thumbnail.url then
    img
      [ src (urlPrefix ++ thumbnail.url)
      , class "selected"
      ]
    []
  else
    img
      [ src (urlPrefix ++ thumbnail.url) ]
```

Comparing our `then` and `else` cases, we see quite a bit of code duplication. The only thing different about them is whether `class "selected"` is present. Can we trim down this code?

Absolutely! We can use the `Html.classList` function. It builds a `class` attribute using a list of tuples, with each tuple containing first the desired class name, and second a boolean for whether to include the class included in the final class string.

Let's refactor our above code to the following, which accomplishes the same thing:

```
viewThumbnail selectedUrl thumbnail =
    img
        [ src (urlPrefix ++ thumbnail.url)
        , classList [ ( "selected", selectedUrl == thumbnail.url ) ]
        ]
    []
```

Now all that remains is to pass in `selectedUrl`, which we can do with an anonymous function. While we're at it, let's also add another `img` to display a larger version of the selected photo.

Listing 2.8 Rendering Selected Thumbnail via anonymous function

```
view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , div [ id "thumbnails" ]
            (List.map (\photo -> viewThumbnail model.selectedUrl photo)
            model.photos
            )
        , img ①
            [ class "large"
            , src (urlPrefix ++ "large/" ++ model.selectedUrl)
            ]
        []
    ]
```

① Display a larger version of the selected photo

If you recompile with the same `elm-make` command as before, the result should now look like Figure 2.9.

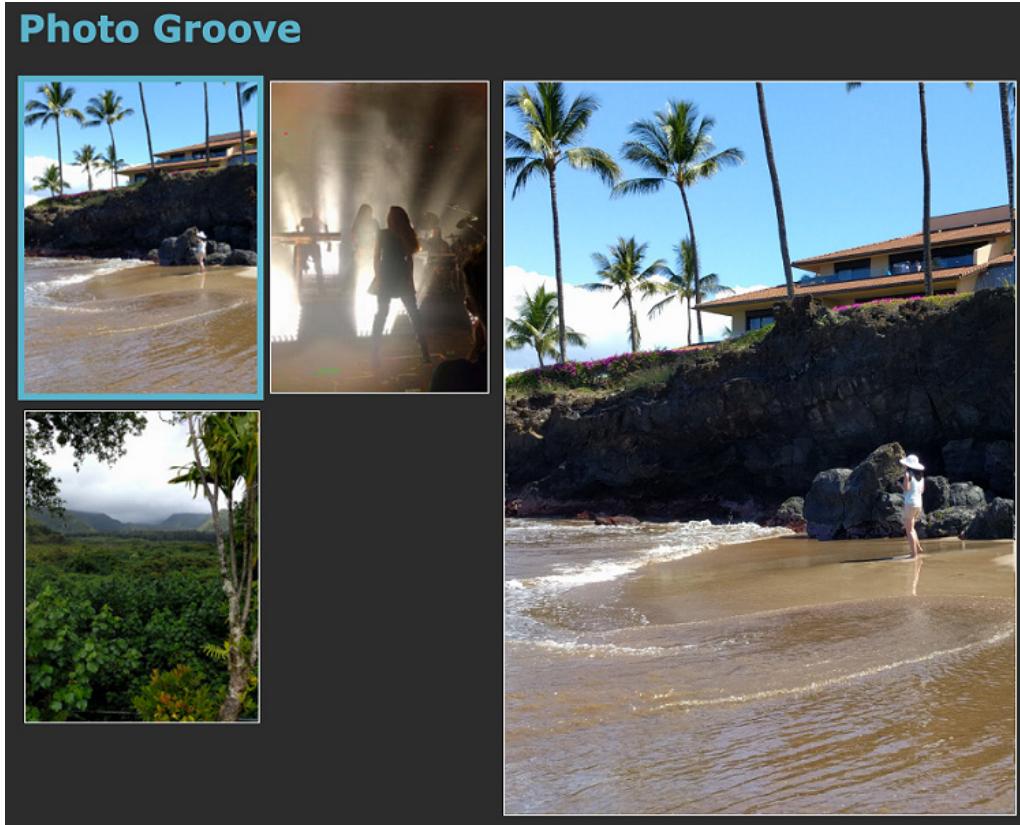


Figure 2.9 Rendering the selected thumbnail alongside a larger version.

Looking good!

REPLACING ANONYMOUS FUNCTIONS WITH PARTIAL APPLICATION

Although the way we've written this works, it's not quite idiomatic Elm code. The idiomatic style would be to remove the anonymous function like so:

```
Before: List.map (\photo -> viewThumbnail model.selectedUrl photo) model.photos
After:  List.map           (viewThumbnail model.selectedUrl)      model.photos
```

Whoa! Does the revised version still work? Do these two lines somehow do the same thing?

It totally does, and they totally do! This is because calling `viewThumbnail` without passing all of its arguments is an example of *partially applying* a function.

DEFINITION *Partially applying* a function means providing one or more of its arguments, and getting back a new function which accepts the remaining arguments and finishes the job.

When we called `viewThumbnail model.selectedUrl photo`, we provided `viewThumbnail` with both of the arguments it needed to return some `Html`. If we call it without that second `photo` argument, what we get back is not `Html`, but rather a function—specifically a function that accepts the missing `photo` argument and *then* returns some `Html`.

Let's think about how this would look in JavaScript, where functions aren't set up to support partial application by default. If we'd written `viewThumbnail` in JavaScript, and wanted it to support partial application, it would have to look like this:

```
function viewThumbnail(selectedUrl) {
    return function(thumbnail) {
        if (selectedUrl === thumbnail.url) {
            // Render a selected thumbnail here
        } else {
            // Render a non-selected thumbnail here
        }
    };
}
```

Functions that can be partially applied, such as the one in this JavaScript code, are known as *curried* functions.

DEFINITION A *curried* function is a function that can be partially applied.

All Elm functions are curried. That's why when we call `(viewThumbnail model.selectedUrl)` we end up partially applying `viewThumbnail`, not getting an `undefined` argument or an error.

In contrast, JavaScript functions are not curried by default. They are instead *tupled*, which is to say they expect a complete “tuple” of arguments. (In this case, “tuple” refers to “a fixed-length sequence of elements,” not specifically one of Elm’s Tuples.)

Elm and JavaScript both support either curried or tupled functions. The difference is which they choose as the default:

- In JavaScript, functions are tupled by default. If you'd like them to support partial application, you can first curry them by hand—like we did in our JavaScript `viewThumbnail` implementation above.
- In Elm, functions are curried by default. If you'd like to partially apply them...go right ahead! They're already set up for it. If you'd like a tupled function, write a curried function that accepts a single Tuple as its argument, then destructure that tuple.

Table 2.1 shows how to define and use both curried and tupled functions in either language.

Table 2.1 Curried functions and Tupled functions in Elm and JavaScript

	Elm	JavaScript
Curried Function	splitA separator str = String.split separator str	<code>function splitA(sep) { return function(str) { return str.split(sep); } }</code>
Tupled Function	splitB (separator, str) = String.split separator str	<code>function splitB(sep, str) { return str.split(sep); }</code>
Total Application	splitB ("-", "867-5309")	splitB("-", "867-5309")
Total Application	splitA "-" "867-5309"	splitA("-")("867-5309")
Partial Application	splitA "-"	splitA("-")

We can use our newfound powers of partial application to make `view` more concise! We now know we can replace our anonymous function with a partial application of `viewThumbnail`.

```
Before: List.map (\photo -> viewThumbnail model.selectedUrl photo) model.photos
After:  List.map      (viewThumbnail model.selectedUrl)      model.photos
```

TIP In Elm, an anonymous function like `(\foo -> bar baz foo)` can always be rewritten as `(bar baz)` by itself. Keep an eye out for this pattern; it comes up surprisingly often.

Here's how our updated `view` function should look.

Listing 2.9 Rendering Selected Thumbnail via partial application

```
view model =
  div [ class "content" ]
    [ h1 [] [ text "Photo Groove" ]
    , div [ id "thumbnails" ]
      (List.map (viewThumbnail model.selectedUrl) model.photos) ①
    , img
      [ class "large"
      , src (urlPrefix ++ "large/" ++ model.selectedUrl)
      ]
    []
  ]
```

① Partially apply `viewThumbnail` with `model.selectedUrl`

Since all Elm functions are curried, it's common to give a helper function more information by adding an argument to the **front** of its arguments list.

For example, when `viewThumbnail` needed access to `selectedUrl`, we made this change:

```
Before: List.map  viewThumbnail           model.photos
After:  List.map (viewThumbnail model.selectedUrl) model.photos
```

Because we added the new `selectedUrl` argument to the front, we could pass it in using partial application instead of an anonymous function. This is a common technique in Elm code!

TIP Since operators are functions, you can partially apply them too! `List.map ((*) 2) [1, 2, 3]` evaluates to `[2, 4, 6]`.

Incidentally, currying is named after acclaimed logician Haskell Brooks Curry. The Haskell programming language is also named after his first name, and whether the Brooks Brothers clothing company is named after his middle name is left as an exercise to the reader.

2.2.2 Handling Events with Messages and Updates

Now that we can properly render which thumbnail is selected, we need to change the appropriate part of the model whenever the user clicks a different thumbnail.

If we were writing JavaScript, we might implement this logic by attaching an event listener to each thumbnail like so:

```
thumbnail.addEventListener("click", function() { model.selectedUrl = url; });
```

Elm wires up event handlers a bit differently. Similarly to how we wrote a `view` function that used Virtual DOM nodes to describe our desired page structure, we're now going to write an `update` function that uses *messages* to describe our desired model.

DEFINITION A *message* is a value used to pass information from one part of the system to another.

When the user clicks a thumbnail, a message will be sent to an `update` function as follows:



Figure 2.10 Handling the event when a user clicks a thumbnail

The format of our message is entirely up to us. We could represent it as a string, or a list, or a number, or anything else we please. Here's a message implemented as a record:

```
{ operation = "SELECT_PHOTO", data = "2.jpeg" }
```

This record is a message which conveys the following information:

"We should update our model to set 2.jpeg as the selectedUrl."

The `update` function receives this message and does the following:

1. Looks at the message it received.
2. Looks at our current model.
3. Uses these two values to determine a new model, then returns it.

We can implement our "select photo" logic by adding this `update` function right above `main`:

```
update msg model =
  if msg.operation == "SELECT_PHOTO" then
    { model | selectedUrl = msg.data }
  else
    model
```

Notice how if we receive an unrecognized message, we return the original model unchanged. This is important! Whatever else happens, the `update` function must always return a new model, even if it happens to be the same as the old model.

ADDING ONCLICK TO VIEWTHUMBNAIL

We can request that a `SELECT_PHOTO` message be sent to `update` whenever the user clicks a thumbnail, by adding an `onClick` attribute to `viewThumbnail`:

```
viewThumbnail selectedUrl thumbnail =
  img
    [ src (urlPrefix ++ thumbnail.url)
    , classList [ ( "selected", selectedUrl == thumbnail.url ) ]
    , onClick { operation = "SELECT_PHOTO", data = thumbnail.url }
    ]
  []
```

The Elm Runtime takes care of managing event listeners behind the scenes, so this one-line addition is the only change we need to make to our view. We're ready to see this in action!

THE MODEL-VIEW-UPDATE LOOP

To wire our Elm application together, we're going to change `main = view model` to the following, which incorporates `update` according to how we've set things up so far.

```
main =
  Html.beginnerProgram
    { model = initialModel
    , view = view
    , update = update
    }
```

The `Html.beginnerProgram` function takes a record with three fields:

- `model` - A value that can be anything you please.
- `view` - A function that takes a model and returns a `Html` node.
- `update` - A function that takes a message and a model, and returns a new model.

It uses these arguments to return a description of a program, which the Elm Runtime sets in motion when the application starts up. Before we got `beginnerProgram` involved, `main` could only render static views. `beginnerProgram` lets us specify how we want to react to user input!

Figure 2.11 demonstrates how data flows through our revised application.

```
update msg model =
  if msg.operation == "SELECT_PHOTO" then
    { model | selectedUrl = msg.data }
  else
    model

{ operation = "SELECT_PHOTO"
, data = "2.jpeg"
}
```

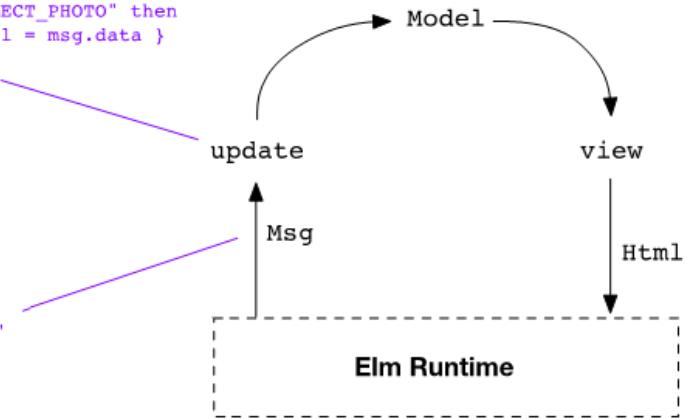


Figure 2.11 Data flowing from the start of the program through the Model-View-Update loop

Notice that `view` builds fresh `Html` values after every `update`. That might sound like a lot of performance overhead, but in practice, it's almost always a performance *benefit*!

This is because Elm doesn't actually recreate the entire DOM structure of the page every time. Instead, it compares the `Html` it got this time to the `Html` it got last time and updates only the parts of the page that are different between the two requested representations.

This approach to "Virtual DOM" rendering, popularized by the JavaScript library React, has several benefits over manually altering individual parts of the DOM:

- Updates are automatically batched to avoid expensive repaints and layout reflows
- It becomes far less likely that application state will get out of sync with the page
- Replying application state changes effectively replays user interface changes

Since `onClick` lives in the `Html.Events` module, we'll need to import it:

```
import Html.Events exposing (onClick)
```

And with that final touch...it's *alive!* You've now written an interactive Elm application!

The complete `PhotoGroove.elm` file should look like this:

Listing 2.10 PhotoGroove.elm with complete Model-View-Update in place

```

module PhotoGroove exposing(..)

import Html exposing(..)
import Html.Attributes exposing(..)
import Html.Events exposing(onClick)

urlPrefix =
    "http://elm-in-action.com/"

view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , div [ id "thumbnails" ]
            (List.map (viewThumbnail model.selectedUrl) model.photos)
        , img
            [ class "large"
            , src (urlPrefix ++ "large/" ++ model.selectedUrl)
            ]
        []
    ]

viewThumbnail selectedUrl thumbnail =
    img
        [ src (urlPrefix ++ thumbnail.url)
        , classList [ ( "selected", selectedUrl == thumbnail.url ) ]
        , onClick { operation = "SELECT_PHOTO", data = thumbnail.url }
        ]
    []

initialModel =
    { photos =
        [ { url = "1.jpeg" }
        , { url = "2.jpeg" }
        , { url = "3.jpeg" }
        ]
    , selectedUrl = "1.jpeg"
    }

update msg model =
    if msg.operation == "SELECT_PHOTO" then
        { model | selectedUrl = msg.data }
    else
        model

main =
    Html.beginnerProgram
        { model = initialModel
        , view = view
        , update = update
        }

```

Let's compile it once more with `elm-make PhotoGroove.elm --output elm.js`. If you open `index.html`, you should now be able to click a thumbnail to select it. Huzzah!

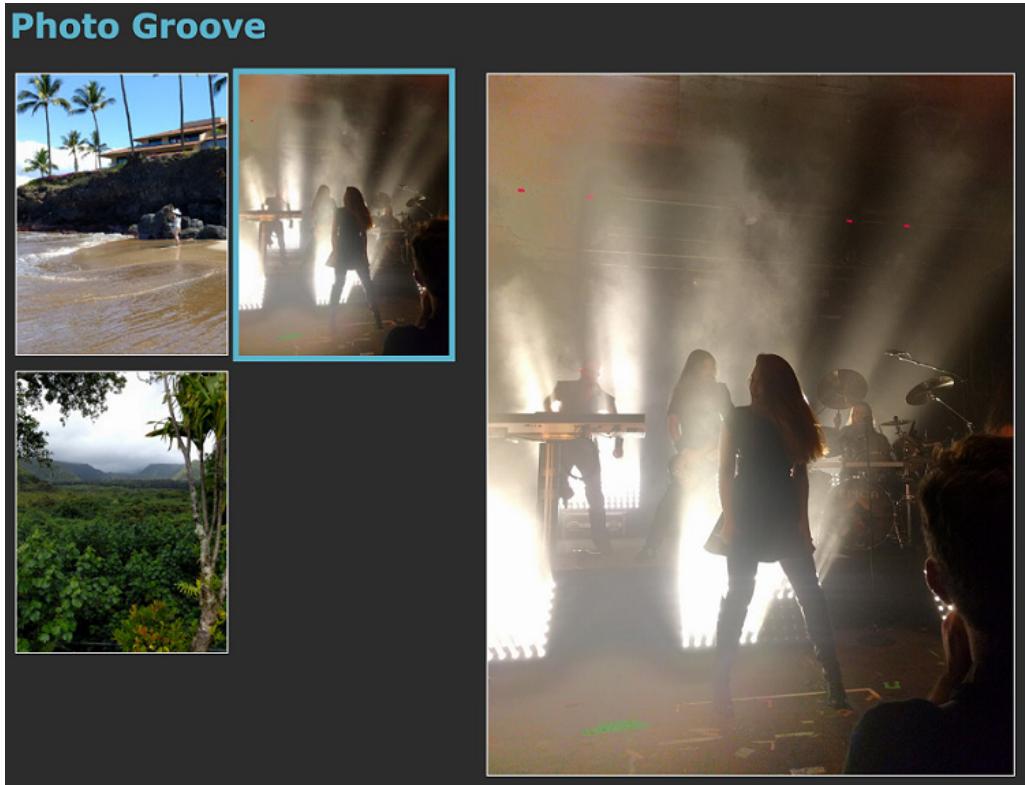


Figure 2.12 Our final Photo Groove application

At this point we've also worked our way through the complete Elm Architecture diagram from the beginning of the chapter. Figure 2.13 shows where things ended up.

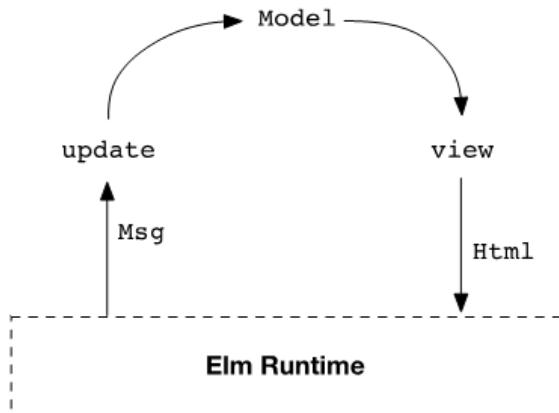


Figure 2.13 Our final Elm Architecture setup

Congratulations on a job well done!

2.3 Summary

In this chapter we learned about three ways to handle interactions, each of which differs from the JavaScript way of handling the same. Table 2.2 summarizes these differences.

Table 2.2 Handling interactions in JavaScript compared to Elm

Interaction	JavaScript approach	Elm approach
Changing the DOM	Directly alter DOM nodes	Return some <code>Html</code> from a view function
Reacting to user input	Attach a listener to an element	Specify a message to send to update
Changing application state	Alter an object in place	Return a new model in update

We covered many other concepts in the course of building our first Elm application, including:

- A model represents our application state
- A view function takes a model and returns a list of `Html` nodes
- User events such as clicks get translated into message values
- Messages get run through the `update` function to produce a new model
- After an `update`, the new model is sent to the view function to determine the new DOM
- `Html.beginnerProgram` wires together `model`, `view`, and `update`
- `List.map` is a higher-order function that translates one list into another
- All Elm functions are curried, which means they can be partially applied

In the next chapter we'll get into ways to improve upon the application we've made so far, both adding features and refactoring to make it easier to maintain.

Onward!

3

Compiler as Assistant

This chapter covers

- Documenting guarantees with type annotations
- Implementing multi-way conditionals with *case-expressions*
- Storing flexible data with *union types*
- Using `Array` and `Maybe` for positional element access
- Generating random numbers using *commands*

In Chapter 2 we built our first Elm application. It doesn't do much yet, but it has potential! So far it displays thumbnails of three photos, and lets users click one to view a larger version.

We showed it to our manager, who was thrilled with what we'd made. "Wow, this is looking *incredible*. The part where you click the thumbnail and it shows the bigger version? Just brilliant. I'm going to get some more team members working with you on this."

Nice! Quite a vote of confidence. Sure, we have no tests or documentation to help get these new teammates up to speed, but there's no time like the present to clean up our code!

Our manager also had a couple of feature requests. "Let's give users the ability to choose between viewing small, medium, or large thumbnails. Also, for this next version I want to kick the fun factor into overdrive. Let's add a button that says *Surprise Me!* and when you click it, it selects one of the photos...**at random.**"

We resist the urge to tell our manager "whoa there, that might be too much fun" and instead review the tasks we've just received:

1. Improve code quality to help new team members get up to speed
2. Let users choose between small, medium, and large thumbnails
3. Add a *Surprise Me!* button that randomly selects a photo

Improving code quality *while* adding new features is often a tall order. Even more so because these particular features involve aspects of Elm we have not yet encountered.

Fortunately, we have an assistant to help us out: Elm’s compiler. In this chapter we’ll learn how it can help us improve documentation, refactor without introducing regressions, and accomplish both while introducing new features. Let’s get to it!

3.1 Documenting Guarantees with Type Annotations

One of the quickest ways we can make our code nicer for incoming teammates is to add some comments that document what our code does. Although comments are simple and flexible, they’re also notoriously unreliable. Sometimes they’re written inaccurately. Other times they start out accurate but become inaccurate as the code base changes out from under them.

“Code never lies. Comments sometimes do.”

—Ron Jeffries

Elm gives us access to a genuinely trustworthy form of documentation: *type annotations*.

3.1.1 Adding Optional Type Annotations

In Chapter 2 we assigned the constant `urlPrefix` to "http://elm-in-action.com". Let’s edit our `PhotoGroove.elm` file to add a type annotation on top of that `urlPrefix` assignment:

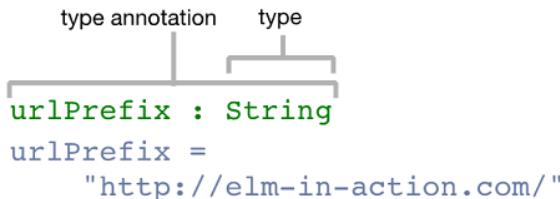


Figure 3.1 Adding a type annotation to `urlPrefix`

This annotation is saying “`urlPrefix` is a `String`,” and it’s not kidding around with that. Elm’s compiler will check our entire code base to verify this claim. If finds even one incident where `urlPrefix` is not used as a `String`, it will give us an error at compile time. This means if our code compiles, we can be certain `urlPrefix` is a `String` absolutely everywhere it’s used!

TIP Searching for `"urlPrefix : "` in your editor is now a quick way to jump to `urlPrefix`’s definition.

This guarantee is even more useful for annotated functions. When teammates are new to the codebase, being able to reliably tell at a glance what arguments a function takes, and what it returns, can be an incredible time-saver.

ANNOTATING FUNCTIONS

We can annotate functions by writing `->` between their arguments and return values.

```
isEmpty : String -> Bool
isEmpty str = str == ""
```

This annotation says “`isEmpty` takes a `String` and returns a `Bool`.” `Bool` refers to one of those `True` or `False` boolean values that we can use as the condition of an *if-expression*.

ANNOTATING RECORDS

Record annotations use `:` instead of `=` but otherwise look about the same as record values.

```
selectPhoto : { operation : String, data : String }
selectPhoto = { operation = "SELECT_PHOTO", data = "1.jpeg" }
```

This annotation says “`selectPhoto` is a record with an `operation` field and a `data` field, and each of those fields is a `String`.” We can add a similar annotation to our `initialModel`:

```
initialModel =
  { photos = [ { url = "1.jpeg" }, { url = "2.jpeg" }, { url = "3.jpeg" } ]
    , selectedUrl = "1.jpeg"
  }
```

Ah, but this record has a list in it! How do we annotate lists?

ANNOTATING LISTS

In Chapter 1 we learned that all elements in an Elm list must have a consistent type. This is reflected in their type annotations. For example, Elm represents a “list of strings” with the type annotation `List String`, and a “list of booleans” with the type annotation `List Bool`.

Let’s play around with some different lists to see their type annotations in `elm-repl`:

Listing 3.1 Annotating Lists

```
> [ "funk", "soul" ]                                     ①
["funk","soul"] : List String

> [ [ "don't", "forget" ], [ "about", "Dre" ] ]          ②
[["don't","forget"],[{"about","Dre"}]] : List (List String)

> [ { url = "1.jpeg" }, { url = "2.jpeg" } ]             ③
[{ url = "1.jpeg" },{ url = "2.jpeg" }] : List { url : String }
```

① “a list of strings”

② “a list of lists of strings”

③ “a list of records”

Notice anything about the structure of that last example? It looks just like the list of photos in our model! This means we can use it to write an annotation for our model. Let's do that:

```
initialModel : { photos : List { url : String }, selectedUrl : String }
initialModel =
  ...
```

Model annotations are among the most helpful forms of documentation for new teammates, because they concisely (and reliably!) describe the structure of our entire application state.

3.1.2 Annotating Functions with Type Variables

Functions, records, and lists served us well in Chapter 2, but implementing the *Surprise Me!* button will involve a collection we haven't used before.

LISTS AND ARRAYS

We mentioned in Chapter 1 that Elm supports arrays as well as lists. Both are sequential collections of varying length, whose elements share a consistent type. However, there are a few differences between arrays and lists:

- Lists can be created with square bracket literals, whereas arrays have no literal syntax in Elm. We always create arrays by calling functions.
- Lists have better performance in typical Elm use cases. This makes them the standard choice for representing a sequence of values, in contrast to JavaScript where arrays are standard. In Chapter 10 we'll see why lists are a better standard choice in Elm.
- Arrays are better for arbitrary positional access. That means it's quick and easy to say "give me the 3rd element in this array." Lists have no first-class way to do this.

Randomly choosing a photo involves arbitrary positional access. As such, we've found ourselves in exactly the sort of situation where an array will serve us better than a list!

ARRAY.FROMLIST

The most common way to obtain an array is by calling `Array.fromList` on an existing list. Table 3.1 shows the results of a few different calls to `Array.fromList`.

Table 3.1 Translating a list into an array using `Array.fromList`

Expression	Result
<code>Array.fromList [2, 4, 6]</code>	Array containing 3 elements: 2, 4, and 6
<code>Array.fromList ["foo"]</code>	Array containing 1 element: "foo"
<code>Array.fromList []</code>	Array containing 0 elements

TYPE VARIABLES

How would we write a type annotation for `Array.fromList`? It's bit different from the functions we've annotated so far, as it returns a different type depending on what you pass it.

Table 3.2 shows some examples of the different types of values it can return.

Table 3.2: Array.fromList's return type depends on its input

Given one of these	Array.fromList returns one of these
List <code>Float</code>	Array <code>Float</code>
List <code>String</code>	Array <code>String</code>
List { <code>url</code> : <code>String</code> }	Array { <code>url</code> : <code>String</code> }

Notice the pattern? Whatever type of list we pass in, that's what type of array we get out.

We can capture this pattern in a type annotate like so:

```
fromList : List elementType -> Array elementType
```

In this annotation, `elementType` is a *type variable*.

DEFINITION A *type variable* represents more than one possible type. Type variables have lowercase names, making them easy to tell them apart from *concrete types* like `String`, which are always capitalized.

When you call `Array.fromList` passing some list, the type variable `elementType` gets replaced by that list's element type. Table 3.3 shows this process in action.

Table 3.3 Replacing Array.fromList's type variable

When we pass this in...	...elementType becomes...	...and we get this back
List <code>Float</code>	<code>Float</code>	Array <code>Float</code>
List <code>String</code>	<code>String</code>	Array <code>String</code>
List { <code>url</code> : <code>String</code> }	{ <code>url</code> : <code>String</code> }	Array { <code>url</code> : <code>String</code> }

CHOOSING NAMES FOR TYPE VARIABLES

When annotating `fromList`, we can choose a different name for its type variable besides `elementType`. Since a type variable essentially serves as a placeholder for a concrete type like `String` or `List Int`, it can have just about any name, so long as it's lowercase and we use it consistently. For example, we could annotate `Array.fromList` in any of the following ways:

- List `elementType` -> Array `elementType`
- List `foo` -> Array `foo`

- List `a` -> Array `a`

In contrast, **none** of the following would work as annotations for `Array.fromList`:

- List `elementType` -> Array `blah`
- List `foo` -> Array `bar`
- List `a` -> Array `b`

These three annotations are saying “`Array.fromList` takes a list containing elements of one type and returns an array that potentially contains elements of another type.”

That just ain’t so! `Array.fromList` can only return an array whose elements have the same type as the list it receives. Its type annotation must reflect that, by having the same type variable for both the `List` it accepts and the `Array` it returns.

NOTE Elm has three type variable names that have special meanings—`number`, `appendable`, and `comparable`—which we’ll dive into later. For now, avoid choosing any of those as your type variable names!

By the way, you’ll often encounter type variables in documentation with single-letter names like `a`, `b`, `c`, and so on. (For example, the official documentation for `Array.fromList` annotates it as `List a -> Array a`.) You are by no means obliged to do the same! Feel free to use the most self-descriptive names you can think of for your own type variables.

CREATING AN ARRAY OF PHOTOS

Since we want to access our photos by position, we’ll be using `Array.fromList` on our `model.photos` list to translate it into an array.

Let’s add this code below our `initialModel` definition in `PhotoGroove.elm`:

```
photoArray : Array { url : String }
photoArray =
    Array.fromList initialModel.photos
```

If we tried to compile this right now, we’d get an error because we have not imported the `Array` module! Let’s add it to our imports at the top of the file:

```
import Array exposing (Array)
```

Nice. Let’s verify that our code still compiles by running `elm-make` on it before proceeding.

Exposing Imported Types

Notice how we wrote `exposing (Array)` there? If we’d written `import Array` without the `exposing (Array)` part, we’d need to refer to the `Array` type in a qualified style, like so:

```
photoArray : Array.Array { url : String }
```

`Array.Array` refers to “the `Array` type from the `Array` module” just like how `Array.fromList` refers to “the `fromList` function from the `Array` module.” Kinda verbose, right?

By writing exposing (`Array`) after import `Array`, we get to skip the “`Array.`” prefix and use the more concise annotation:

```
photoArray : Array { url : String }
```

Much better!

3.1.3 Reusing Annotations with Type Aliases

Now that we’ve created `photoArray`, we can see that our annotations for `initialModel` and `photoArray` have a bit of code duplication going on: they both include `{ url : String }`.

```
initialModel : { photos : List { url : String }, selectedUrl : String }
photoArray : Array { url : String }
```

We can replace this duplication with shared code using a *type alias*.

DEFINITION A *type alias* assigns a name to a type. Anywhere you would refer to that type, you can substitute this name instead.

Let’s create a type alias called `Photo`, and then use that in place of `{ url : String }`.

Listing 3.2 Creating a type alias for Photo

```
type alias Photo =
    { url : String }

initialModel : { photos : List Photo, selectedUrl : String }
initialModel =
    ...
photoArray : Array Photo
photoArray =
    Array.fromList initialModel.photos
```

① “whenever I say `Photo`, I mean `{ url : String }`”

Not only does this make our `initialModel` annotation more concise, it also makes it easier to maintain! Now if we add a new field to our `Photo` record, we can change the type alias in one place instead of having to hunt down several individual annotations.

ANNOTATING INITIALMODEL USING A TYPE ALIAS

Since `initialModel`, `view`, and `update` all involve our model record, let’s add a type alias for `Model` as well, and then revise `initialModel`’s type annotation to use it:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

```
type alias Model =
    { photos : List Photo
    , selectedUrl : String
    }

initialModel : Model
```

That takes care of `initialModel`. What about `view`?

HTML'S TYPE VARIABLE

We know `view` takes a `Model` and returns `Html`, but we can't just write "`view returns Html`" and call it a day. This is because, just like `List` and `Array`, the `Html` type has a type variable!

`Html`'s type variable reflects the type of message it sends to `update` in response to events from handlers like `onClick`.

Table 3.4 Comparing type variables for List and Html

Value	Type	Description
<code>["foo"]</code>	<code>List String</code>	<code>List of String elements</code>
<code>[3.14]</code>	<code>List Float</code>	<code>List of Float elements</code>
<code>div [onClick "foo"] []</code>	<code>Html String</code>	<code>Html producing String messages</code>
<code>div [onClick 3.14] []</code>	<code>Html Float</code>	<code>Html producing Float messages</code>
<code>div [onClick { x = 3.3 }] []</code>	<code>Html { x : Float }</code>	<code>Html producing { x : Float } messages</code>

Since our `onClick` handler produces messages in the form of records that have an `operation` string and a data string, our `view` function's return type is:

```
Html { operation : String, data : String }
```

That's pretty verbose! Even though it won't remove any code duplication yet, it's perfectly fine to introduce a type alias to make `view`'s annotation more concise. Let's add one above `view`:

```
type alias Msg =
    { operation : String, data : String }

view : Model -> Html Msg
```

Excellent! Next we'll look into annotating `view`'s helper function, `viewThumbnail`.

3.1.4 Annotating Longer Functions

We've now seen a few type annotations for functions that take a single argument, but none for functions that take multiple arguments. Our `viewThumbnail` function will be the first.

To learn how to annotate a multi-argument function, let's play around with one that has a simpler return type: the humble `padLeft` function. `padLeft` makes sure strings meet a certain minimum length. You give it a minimum length, a "filler character," and a string. If the string is not at least the given length, `padLeft` adds filler characters to its left until the string reaches that length.

We can try it out in `elm-repl`, since `padLeft` is included in Elm's core `String` module:

```
> String.padLeft 9 '.' "not!"
"....not!" : String

> String.padLeft 2 '.' "not!"
"not!" : String
```

We can see that `String.padLeft` takes three arguments—an `Int`, a `Char`, and a `String`—and then returns another `String`. How can we annotate a function like that?

Believe it or not, we've already seen the answer! It's just one of those answers that likes to hide in plain sight.

ANNOTATING A PARTIALLY APPLIED FUNCTION

Let's think back to Chapter 2, where we learned that all functions in Elm are *curried*. That means they all support *partial application*—the practice of calling a function without passing all the arguments it requires. If we call `padLeft` passing only one of its three arguments, the result will be a function that takes the remaining arguments and "finishes the job."

Here's a partial application of `padLeft`. What would the annotation for this be?

```
padNine = String.padLeft 9
```

`padLeft 9` is a partial application, so we know `padNine` must be a function. That's a start!

We also know that `padLeft` takes a `Char` as its next argument after the `9`. That gives us enough information to rough out some pseudocode:

```
padNine : Char -> (a function of some sort)
padNine = String.padLeft 9
```

Now let's suppose we gave this `padNine` function the `Char` it wants. Passing a `'.'` character to `padNine` would be the same as passing `9` and then `'.'` to the original `padLeft`, like so:

```
padNineDots = String.padLeft 9 .'
```

How would we annotate `padNineDots`? Since it gives `padLeft` two of three arguments, that means only one more argument is needed to finish the job. That last argument is a `String`, so **padNineDots must take a String and return a String**. We know how to write that one!

```
padNineDots : String -> String
padNineDots = String.padLeft 9 .'
```

We now know that:

- `padNine` takes a `Char` and returns a function (of some type)
- If you pass `padNine` a `Char`, it returns a `String -> String` function

Putting those two together, our mystery function must be a `String -> String`!

That tells us we can annotate `padNine` like so:

```
padNine : Char -> (String -> String)
padNine = String.padLeft 9
```

Now let's unwind our original partial application. If we cease partially applying the `9`, we can follow the pattern we saw here to arrive at a valid annotation for the original `padLeft` itself:

```
padLeft : Int -> (Char -> (String -> String))
```

This is a perfectly accurate annotation, but it's a bit heavy on the parentheses, yeah? Fortunately, we can be more concise. Elm's syntax lets us omit the parentheses here, so the following two annotations are equivalent:

```
padNine : Int -> (Char -> (String -> String))
padNine : Int -> Char -> String -> String
```

Either style works just as well as the other, but the latter is considered best practice. You can see `elm-repl` use this annotation style in Listing 3.3, as we partially apply `String.padLeft` step by step, until it has no more arguments to accept and finally returns a `String`.

Listing 3.3 Multi-Argument function annotations

```
> String.padLeft
<function:padLeft> : Int -> Char -> String -> String

> String.padLeft 9
<function> : Char -> String -> String

> String.padLeft 9 '.'
<function> : String -> String

> String.padLeft 9 '.' "not!"
"....not!" : String
```

Notice how each time we partially applied `padLeft`, our annotation got shorter:

Function	Type Annotation
<code>String.padLeft</code>	<code>Int -> Char -> String -> String</code>
<code>String.padLeft 9</code>	<code>Char -> String -> String</code>
<code>String.padLeft 9 '.'</code>	<code>String -> String</code>
<code>String.padLeft 9 '.' "not!"</code>	<code>String</code>

If you think about it, this means technically **every Elm function takes only one argument**.

After all, any function that *appears* to take multiple arguments is ultimately calling out to single-argument functions behind the scenes. The fact that Elm lets you omit the parentheses for these nested calls is just a syntactic convenience.

TIP You can now tell your friends that in Chapter 2 of this book, you wrote an entire working Elm application where every function took only a single argument. This sounds like some really hardcore programming unless your friends are familiar with currying!

ANNOTATING VIEWTHUMBNAIL

Armed with this knowledge, we can follow this pattern to annotate `viewThumbnail` like so:

```
viewThumbnail : String -> Photo -> Html Msg
viewThumbnail selectedUrl thumbnail =
```

Splendid! Our code is getting easier and easier for teammates to pick up at a glance.

3.2 Case-expressions and Union Types

Now that our documentation situation is looking better, let's shift gears to work on those two new features: the *Surprise Me!* button and the thumbnail size chooser.

We'll introduce these iteratively, first adding the visual elements with only a token level of interactivity, and then circling back to make them work properly afterwards.

3.2.1 Using case-expressions

Let's start by adding the *Surprise Me!* button to our `view`, right above the thumbnails `div`.

```
, button
  [ onClick { operation = "SURPRISE_ME", data = "" } ]
  [ text "Surprise Me!" ]
, div [ id "thumbnails" ]
```

Next let's revise our `update` function to add a quick `else if` branch for our new `SURPRISE_ME` operation, and a shiny new type annotation while we're at it.

```
update : Msg -> Model -> Model
update msg model =
  if msg.operation == "SELECT_PHOTO" then
    { model | selectedUrl = msg.data }
  else if msg.operation == "SURPRISE_ME" then
    { model | selectedUrl = "2.jpeg" }
  else
    model
```

This implementation always selects the second photo. That's not much of a *surprise*, granted, but we'll get there. For now, let's smooth out that code duplication we just introduced!

REFACTORING AN IF-EXPRESSION INTO A CASE-EXPRESSION

We now have two conditionals that do nothing more than compare `msg.operation` to a string. Let's express this more cleanly by rewriting these *if-expressions* as a *case-expression*:

```
case msg.operation of
  "SELECT_PHOTO" ->
    { model | selectedUrl = msg.data }

  "SURPRISE_ME" ->
    { model | selectedUrl = "2.jpeg" }

  _ ->
    model
```

equivalent *if-expression*

```
if msg.operation == "SELECT_PHOTO" then
  { model | selectedUrl = msg.data }

else if msg.operation == "SURPRISE_ME" then
  { model | selectedUrl = "2.jpeg" }

else
  model
```

Figure 3.2 Refactoring an if-expression into a case-expression

Whereas an *if-expression* is a two-way conditional, a *case-expression* is a multi-way conditional. It lets us compare something to a wider range of values than just `True` or `False`.

Just like with a JavaScript *switch-statement*, we begin a *case-expression* by providing a value which will be run through a series of comparisons. Here we wrote `case msg.operation of` because we want to run `msg.operation` through these comparisons.

Following the `case` are a series of *branches* such as the following:

```
"SELECT_PHOTO" ->
  { model | selectedUrl = msg.data }
```

This says that if `msg.operation` is equal to `"SELECT_PHOTO"`, then the branch after the `->` will be evaluated, and the entire *case-expression* will evaluate to that branch's result.

If `msg.operation` is not equal to `"SELECT_PHOTO"`, then the next branch (in this case the `"SURPRISE_ME"` `->` branch) will be checked in the same way, and so on.

Just like *if-expressions*, *case-expressions* must always evaluate to a single value, meaning exactly one branch must always be chosen for evaluation. When writing *if-expressions* we have `else` to ensure we end up with a value no matter what, whereas in *case-expressions* we can use the default branch of `_ ->` for the same purpose.

NOTE Elm has no equivalent of JavaScript's `break` statement. This is no problem, since *case-expression* branches don't "fall through" like branches in JavaScript's *switch-statements* do.

Our code is really shaping up! Let's copy the *case-expression* from Figure 3.2 into `PhotoGroove.elm`, then run `elm-make PhotoGroove.elm --output elm.js` to recompile it.

TRYING OUT THE ‘SURPRISE ME!’ BUTTON

Before we open up `index.html` and view our progress, we have one last change to make. While we’ve been hard at work on code quality and new functionality, one of our new teammates has been busy improving our stylesheets.

Let’s edit `index.html` to make use of their work, by changing its `<head>` to the following:

```
<head>
  <link rel="stylesheet" href="http://elm-in-action.com/styles.css">
</head>
```

Lovely! Now we can open our revised `index.html` to see the new button in all its glory.

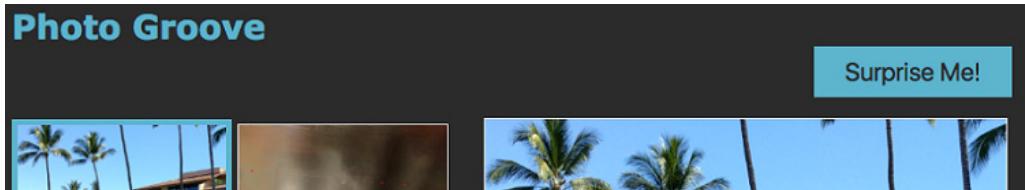


Figure 3.3 The new Surprise Me! button

In addition to styling the *Surprise Me!* button, our helpful teammate went ahead and added styles for the Thumbnail Size Chooser too.

How about we introduce some logic to go with those styles?

3.2.2 Enumerating possibilities with Union Types

Our second feature lets users choose one of three thumbnail sizes: small, medium, or large.

The first step toward implementing it will be storing the current `chosenSize` in our model. In JavaScript we might represent `chosenSize` as a String—perhaps setting it to either “`SMALL`”, “`MEDIUM`”, or “`LARGE`”—but in Elm we can do better with *union types*.

DEFINITION A union type is a custom type you create by specifying the values it can contain.

One use for a union type is as an enumeration of values. Figure 3.4 illustrates how we can define a union type to represent the different thumbnail size choices we’ll support:

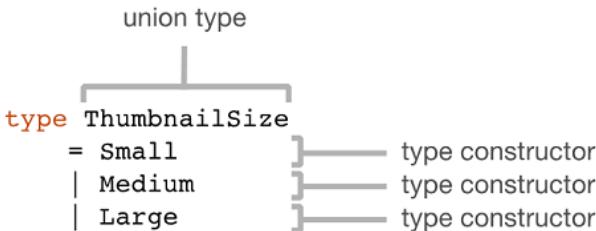


Figure 3.4 Defining an enumeration of values with a union type called `ThumbnailSize`

This is saying “define a new type called `ThumbnailSize` with three possible values: `Small`, `Medium`, and `Large`.” Let’s add this code to `PhotoGroove.elm`, right above `type alias Msg`.

We can now define `chosenSize` in one of the following ways:

```

chosenSize : ThumbnailSize
chosenSize = Small

chosenSize : ThumbnailSize
chosenSize = Medium

chosenSize : ThumbnailSize
chosenSize = Large
  
```

Notice that in each of these examples, the type of `chosenSize` is `ThumbnailSize`. These new values we’ve created—`Small`, `Medium`, and `Large`—aren’t “actually integers under the hood” or “just strings behind the scenes.” We really have scratch-built a brand new type here! Trying to compare a `ThumbnailSize` to a number, string, or any other type (using `=` or any other comparison) will yield an error at build time.

This is different from `type alias`, which gives a name to an existing type—much like how a variable gives a name to an existing value.

These values are also unique. The expression `Medium == Medium` is `True`, but put absolutely any other value (besides `Medium`) on either side of that `=` and the expression will no longer be `True`. The same can be said of `Small == Small` and `Large == Large`.

NOTE Boolean values in Elm are capitalized because `Bool` is a union type: `type Bool = True | False`

Let’s add the `chosenSize` field to our model and initialize it to `Medium` in `initialModel`:

```

type alias Model =
  { photos : List Photo
  , selectedUrl : String
  , chosenSize : ThumbnailSize
  }

initialModel : Model
initialModel =
  { photos = ...
  , selectedUrl = "1.jpeg"
  }
  
```

```
, chosenSize = Medium
}
```

Great! Next we'll render some radio buttons to let users change it.

RENDERING A THUMBNAIL SIZE RADIO BUTTON

We can use our newfound knowledge of union types and *case-expressions* to write a helper function that takes a `ThumbnailSize` and renders a radio button which chooses that size. Let's add the following to `PhotoGroove.elm`, right below `viewThumbnail`:

Listing 3.4 Rendering a thumbnail size radio button

```
viewSizeChooser : ThumbnailSize -> Html Msg
viewSizeChooser size =
    label []
        [ input [ type_ "radio", name "size" ] []
        , text (sizeToString size)
        ]

sizeToString : ThumbnailSize -> String
sizeToString size =
    case size of
        Small -> ①
            "small" ①

        Medium -> ②
            "med" ②

        Large -> ③
            "large" ③
```

- ① evaluates to "small" if `size == Small`
- ② evaluates to "med" if `size == Medium`
- ③ evaluates to "large" if `size == Large`

NOTE The underscore in that `type_` attribute is very important! As we've seen, `type` is a reserved keyword in Elm (used to define union types), which means it can't be used for an `Html` attribute name. The `Html.Attributes` module names the attribute `type_` to work around this.

Notice anything missing from that *case-expression*? It has no default branch! We accounted for the cases where `size` is `Small`, `Medium`, and `Large`...but what if `size` has some other value, such as `null`, `undefined`, or `"halibut"`?

Can't happen! `sizeToString`'s type annotation tells us that `size` is guaranteed to be a `ThumbnailSize`, and a `ThumbnailSize` can only be `Small`, `Medium`, or `Large`. It can't be `null` or `undefined` because those don't exist in Elm, and it can't be `"halibut"` because that's a `String`, not a `ThumbnailSize`.

Elm's compiler knows we've covered every possibility here, so it doesn't require a default branch. In fact, if we tried to add one, the compiler would politely inform us that we'd written unreachable code.

RENDERING THREE RADIO BUTTONS

We can now call this `viewSizeChooser` function three times, passing `Small`, `Medium`, and `Large`, to render the three radio buttons on the page. Let's do that right above `thumbnails`:

```
, h3 [] [ text "Thumbnail Size:" ]
, div [ id "choose-size" ]
  [ viewSizeChooser Small, viewSizeChooser Medium, viewSizeChooser Large ]
, div [ id "thumbnails" ]
  (List.map (viewThumbnail model.selectedUrl) model.photos)
```

Once again we see some code duplication: we're calling the same function (`viewSizeChooser`) on every element in a list. Any time we're calling the same function on every element in a list, there's a good chance we can make our code cleaner with `List.map`!

This is absolutely one of those times. Let's refactor our `choose-size` element to this:

```
, div [ id "choose-size" ]
  (List.map viewSizeChooser [ Small, Medium, Large ])
```

Not only does this refactor reduce duplication, it also makes our code more concise.

RENDERING DIFFERENT THUMBNAIL SIZES

Rendering the interface for *choosing* a thumbnail size is great, but it's only half the battle! We still need to actually render our thumbnails differently based on what the user has chosen.

Thanks to our coworker having written a stylesheet for us, we can implement this by adding the class `"small"`, `"med"`, or `"large"` to our `thumbnails` container. Since those classes conveniently correspond to the results of our `sizeToString` function, we can go ahead and replace our current `div [id "thumbnails"]` with the following:

```
div [ id "thumbnails", class (sizeToString model.chosenSize) ]
```

TIP It would be best practice to write a separate `sizeToClass` function and use it here in place of `sizeToString`, even if their implementations were identical for the time being. Having separate functions means if someone later changes the text on the radio buttons, our thumbnail classes won't accidentally break.

Let's see how it looks! Since we haven't implemented the radio button logic yet, try changing the `chosenSize` value in `initialModel` to `Large` and then recompile to see how that affects the way the page looks. Then try recompiling again with `Small`, and finally back to `Medium`.

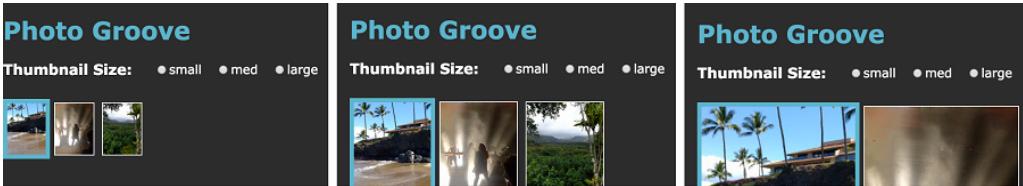


Figure 3.5 Recompiling with `chosenSize = Small`, `chosenSize = Medium`, and `chosenSize = Large`

Cool! Now that we have our interface rendering, it's time to start making the logic work.

3.2.3 Holding Data in Union Types

Making the *Surprise Me!* button change our model's `selectedUrl` field to a random photo's `url` means we'll need to access a random photo from our `photoArray`.

In JavaScript we might implement reading a photo from an array as follows:

```
var photos = [ ...maybe there are photos here, but maybe not... ];
var selectedId = photos[2].url; // Select the third photo
```

What happens if `photos` doesn't have at least three elements? In that case, `photos[2]` would evaluate to `undefined`, and `photos[2].url` would throw an angry runtime exception such as:

```
Uncaught TypeError: Cannot read property 'url' of undefined
```

Elm avoids this runtime exception by handling this situation differently in three ways:

1. In Elm you call `Array.get` to get an element in an Array; there is no `[2]` accessor.
2. `Array.get` never returns `undefined` or `null`, because Elm has neither of these.
3. `Array.get` **always** returns a container value called a `Maybe`.

What's a `Maybe`? Glad you asked!

DEFINITION A `Maybe` is a container like a `List`, except it can hold at most one element.

`Maybe` is implemented as a union type, but a flavor of union type we haven't seen yet: one that holds data. Its definition looks like this:

```
type Maybe value
  = Just value
  | Nothing
```

MAYBE VERSUS UNDEFINED

The most common use for `Maybe` is to represent the potential absence of a value. It provides a **container-based** alternative to JavaScript's **drop-in replacements** of `null` and `undefined`.

Figure 3.6 compares Elm's `Array.get` and `Maybe` to JavaScript's `[2]` and `undefined`.

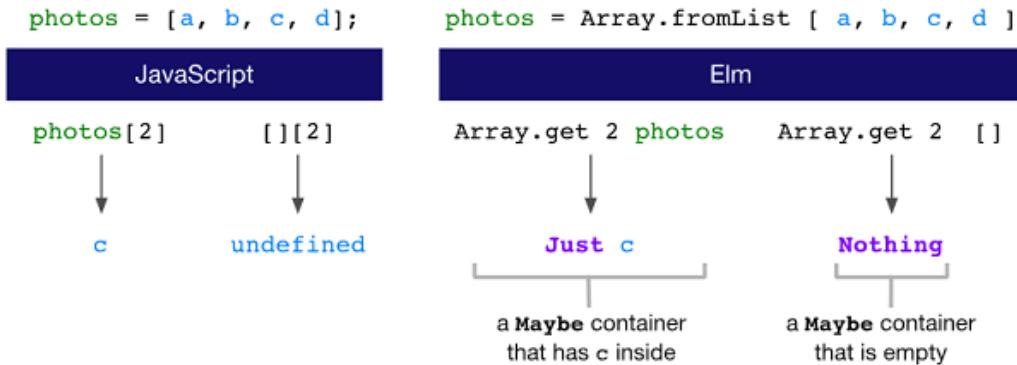


Figure 3.6 Comparing Elm's `Maybe` to JavaScript's `undefined`

`Array.get 2 photos` expresses the desire to access the third element of an array called `photos`. However, the `photos` array may have fewer than three elements in it! If the index 2 is outside the bounds of `photos`, then `Array.get 2 photos` will return `Nothing`.

On the other hand, if 2 is inside its bounds, then `Array.get 2 photos` will return `Just c`—where `c` is the element at index 2 in the `photos` array.

MAYBE'S TYPE VARIABLE

As we've seen before with `List`, `Array`, and `Html`, the `Maybe` type has a type variable!

We chose to name it `value` in the definition we gave for `Maybe`'s union type earlier:

```
type Maybe value
  = Just value
  | Nothing
```

Similarly to how you can have a `List String` or a `List Photo`, you can also have a `Maybe String` or `Maybe Photo`. The difference is that whereas `List Photo` means “a list of photos,” `Maybe Photo` means “either a `Photo` or nothing at all.” Put another way, `Maybe` is a container that can hold at most one element.

Now that we know `Just "1.jpeg"` is a `Maybe String`, what is the type of `Just` itself?

"JUST" IS A FUNCTION

Argh, that heading totally gave away the answer. Oh well.

As the heading suggests, whereas `Nothing` is a `Maybe` value, `Just` is a function that returns a `Maybe` value. We can confirm this in `elm-repl`:

```
> Nothing
Nothing : Maybe.Maybe a

> Just
<function> : a -> Maybe.Maybe a
```

```
> Just "dance"
Just ("dance") : Maybe.Maybe String
```

NOTE elm-repl calls the type `Maybe.Maybe` because the `Maybe` type lives in a module called `Maybe`. This is similar to `Array.Array` from earlier, except that unlike the `Array` module, the `Maybe` module is imported by default. (The `Maybe.Maybe` type is not exposed by default.)

In Chapter 1 we noted that when you write a function by hand, its name must be lowercase...yet here we see that union type constructors like `Just` are functions with capitalized names. What makes them so special?

A capital question! The answer is right around the corner.

DESTRUCTURING UNION TYPES

In Chapter 1 we learned how we can *destructure* tuples to both extract and name their values.

```
multiply3d ( x, y, z ) = x * y * z
```

As it turns out, we can also destructure union type constructors such as `Just` in the branches of *case-expressions*. Let's use what we've learned about `Array.get` and `Maybe` to add this `getPhotoUrl` function right above our `update` function.

Listing 3.5 Selecting a photo by index

```
getPhotoUrl : Int -> String
getPhotoUrl index =
    case Array.get index photoArray of
        Just photo -> ①
            photo.url
        Nothing ->
            "" ②
```

- ① Destructuring `Just` and naming its contained value `photo`
- ② Fall back on `""` if there was no photo at that index

Here, `Just photo ->` is saying two things:

1. This branch matches a `Maybe` value that was created using the `Just` type constructor.
2. We're extracting the value that was passed to `Just` and naming it `photo`.

NOTE This is where the distinction between capitalized and uncapitalized functions matters. By comparing their capitalizations, Elm's compiler can tell that `Just photo ->` refers to a type constructor called `Just` which holds a value we've chosen to name `photo`. If we'd instead written `Just True ->` the compiler would know we meant "the `Just` type constructor holding exactly the value `True`."

Notice that this *case-expression* did not need a `_ ->` branch! This is because its branches already cover all possible values of `Maybe`. One branch covers `Nothing`, and the other covers

any values created with `Just`. Since `Nothing` and `Just` are the only ways to obtain a `Maybe`, these branches have every possibility covered.

Also notice how because `Array.get` returns a `Maybe`, you can't help but remember to handle the case where `index` is outside the bounds of `photoArray`. This API design means that you'll always be doing the minimal amount of coding necessary to handle error cases. You don't need to defensively check whether something is `null` or `undefined`, because Elm's compiler has your back in cases where the value you want might not be available.

3.2.4 Representing Flexible Messages with Union Types

We've made so much progress! We've updated our views to render the new interfaces, expanded our model to hold our new `chosenSize` field, and introduced helper functions for `photoArray`. Now it's time to wire all that together with changes to `update`.

Let's start by updating that `chosenSize` field, which currently gets initialized to `Medium` and then never changes. We'll want to update it whenever the user clicks a radio button.

This might seem like the obvious change to make to `viewSizeChooser`:

```
viewSizeChooser : ThumbnailSize -> Html Msg
viewSizeChooser size =
    ...
        onClick { operation = "SET_SIZE", data = size }
```

There's just one problem...this is a type mismatch! Remember our type alias for `Msg`?

```
type alias Msg =
    { operation : String, data : String }
```

This says the `data` field of a `Msg` needs to be a `String`, not a `ThumbnailSize`! To preserve backwards compatibility with existing code, we need a `Msg` type that can accommodate both `ThumbnailSize` and `String`. One way we can do this is by adding a new field called `sizeData`:

```
type alias Msg =
    { operation : String, data : String, sizeData : ThumbnailSize }
```

This solves one problem while creating another. Now our existing `onClick` handler will no longer compile, because it's missing a field. Here's how we wrote it back in Chapter 2:

```
onClick { operation = "SELECT_PHOTO", data = thumbnail.url }
```

Without specifying a value for that new `sizeData` field, this record is no longer a valid `Msg`. We'd need to change it to something like this:

```
onClick { operation = "SELECT_PHOTO", data = thumbnail.url, sizeData = Small }
```

Yuck. Having to set `sizeData = Small` for an operation that doesn't care about `sizeData`? That doesn't smell right. And are we really planning to add a field to every `Msg` in our program whenever we need it to support a new data type? This doesn't seem like it will scale well.

Is there a better way?

IMPLEMENTING MSG AS A UNION TYPE

There totally is! Implementing `Msg` as a union type will work much better. Here's the plan:

1. Replace our type alias `Msg` declaration with a type `Msg` declaration.
2. Revise `update` to use a *case-expression* which destructures our new union type.
3. Change our `onClick` handler to pass a type constructor instead of a record.

Listing 3.6 shows the first two changes: replacing type alias `Msg` with type `Msg`, and revising `update` accordingly. Let's edit `PhotoGroove.elm` to incorporate these changes.

Listing 3.6 Implementing `Msg` as a union type

```
type Msg
    = SelectByUrl String      1
    | SurpriseMe              2
    | SetSize ThumbnailSize   3
                                4

update : Msg -> Model -> Model
update msg model =
    case msg of
        SelectByUrl url ->      5
            { model | selectedUrl = url }

        SurpriseMe ->           6
            { model | selectedUrl = "2.jpeg" }
```

- 1 Replaces our earlier declaration of type alias `Msg`
- 2 Replaces our "SELECT_PHOTO" operation
- 3 Replaces our "SURPRISE_ME" operation
- 4 Our new operation: set thumbnail size
- 5 Previous condition: `msg.operation == "SELECT_PHOTO"`
- 6 Nothing to destructure for `SurpriseMe`!

This change means our `onClick` handlers now expect type constructors instead of records. We'll need to make this change in `view`:

```
Old: onClick { operation = "SURPRISE_ME", data = "" }
New: onClick SurpriseMe
```

Let's also make this change in `viewThumbnail`:

```
Old: onClick { operation = "SELECT_PHOTO", data = thumbnail.url }
New: onClick (SelectByUrl thumbnail.url)
```

THE MISSING PATTERNS ERROR

If we recompile, we'll see a type of error we haven't seen before: a *missing patterns* error.

```
-- MISSING PATTERNS ----- PhotoGroove.elm

This `case` does not have branches for all possibilities.
...
You need to account for the following values:

  PhotoGroove.SetSize _

Add a branch to cover this pattern!
```

Oops! The compiler noticed we only handled `SelectByUrl` and `SurpriseMe` in our update function's *case-expression*; we never implemented the logic for `SetSize`. Fortunately, that deficiency will never reach our end users because the compiler *didn't let us forget*.

TIP Anytime you use the default case `_ ->` in a *case-expression*, you cannot get this error. That is not a good thing! The *missing patterns* error is your friend. When you see it, it's often saving you from a bug you would have had to hunt down later. Try to use `_ ->` only as a last resort, so you can benefit from as many *missing pattern* safeguards as possible.

Let's add this to the end of our `update` function:

```
SurpriseMe index ->
  { model | selectedUrl = "2.jpeg" }

SetSize size ->
  { model | chosenSize = size }
```

Finally, we need to add a new `onClick` handler to our radio button in `viewSizeChooser`:

<i>Old:</i> <code>input [type_ "radio", name "size"] []</code> <i>New:</i> <code>input [type_ "radio", name "size", onClick (SetSize size)] []</code>
--

If we recompile and click the radio buttons, we can now see the thumbnail sizes change!

TIP There are at least two ways we could improve user experience here. One way would be making the "medium" option display as selected on page load. Another would be to use a broader event handler than `onClick`—one that detects whenever the radio state changes, even if it's not from a click. Try implementing these improvements sometime for practice!

Using union types for `Msg` has two major advantages over records.

1. The compiler can save us from typos. Before, if we wrote "`SELECT_PHOTOA`" instead of "`SELECT_PHOTO`", our code would silently fail and we'd have to hunt down the cause. Now, if we write `SelectByUrl` instead of `SelectByUrl`, Elm's compiler will give us an error at build time—including the line number of the typo. No bug hunting necessary!

- Each flavor of `Msg` now holds only the minimum amount of data it needs. `SelectByUrl` holds only a `String`, and `SetSize` holds only a `ThumbnailSize`. Best of all, `SurpriseMe` doesn't need to hold any data at all...and so it doesn't!

To sum up, implementing `Msg` as a union type has made our code more reliable, more concise, and easier to scale. These advantages make union types the typical choice for representing messages in production Elm applications, and they're what we will use for the rest of the book.

Now all that remains is to incorporate random number generation into `SurpriseMe`.

3.3 Generating Random Numbers with Commands

Currently our *Surprise Me!* button always selects the second photo, but we want it to select a random photo instead. Here's how we're going to do that:

- Generate a random integer between 0 and 2. This will be our index.
- Ask `photoArray` for the photo it's storing at that index.
- Set our model's `selectedUrl` to be that photo's `url`.

We'll start by generating the random integer.

3.3.1 Describing Random Values with Random.Generator

In JavaScript, `Math.random()` is the typical starting point for generating random values. Calling `Math.random()` gives you a random `Float` between 0 and 1. To randomly generate an `Int` or a `String`, you manually convert that `Float` into the value you actually want.

In Elm we start with a `Random.Generator`, which specifies the type of value we want to randomly generate. For example, the `Random.int` function takes a lower bound and an upper bound, and returns a `Random.Generator` that generates a random `Int` between those bounds.

Here's a generator that randomly generates integers between 0 and 2:

```
randomPhotoPicker : Random.Generator Int
randomPhotoPicker =
    Random.int 0 2
```

Notice how the type annotation says `Random.Generator Int`. This means "we have a `Random.Generator` that produces `Int` values." If we'd used `Random.bool` instead of `Random.int`, it would have returned a `Random.Generator Bool` instead.

GENERATING A RANDOM PHOTO INDEX

Since `randomPhotoPicker` generates random integers between 0 and 2, and `photoArray` has 3 photos in it, we could use this generator to pick an index within `photoArray` to select. That would work, but it would be brittle since it relies on `photoArray` having exactly 3 photos. We can improve it by replacing the hardcoded 2 with `(Array.length photoArray - 1)`.

Let's add the following after our `getPhotoUrl` declaration.

```
randomPhotoPicker : Random.Generator Int
```

```
randomPhotoPicker =
    Random.int 0 (Array.length photoArray - 1)
```

That'll do nicely. We also need to add the `Random` module to our imports.

```
import Random
```

Now that we've described the random value we want, it's time to actually generate one!

3.3.2 Introducing Commands to the Elm Architecture

If you call JavaScript's `Math.random()` five times, you'll probably get back five different numbers. Elm functions are more consistent. If you call any Elm function five times with the same arguments, you can expect to get the same return value each time. This is no mere guideline, but a language-level guarantee! Knowing that all Elm functions have this useful property makes bugs easier to track down and reproduce.

Since Elm forbids functions from returning different values when they receive the same arguments, it's not possible to write an inconsistent function like `Math.random` as a plain Elm function. Instead, Elm implements random number generation using a *command*.

DEFINITION A *command* is a value that describes an operation for the Elm Runtime to perform. Unlike functions, running the same command multiple times can have different results.

When the user clicks *Surprise Me!* we'll use a command to translate our `Random.Generator Int` into a randomly-generated `Int`, which will represent our new selected photo index.

RETURNING COMMANDS FROM UPDATE

Remember how we specified what `onClick` should do in terms of a message that got sent to our `update` function? We didn't say "add this click event listener to the DOM right away," we said "I want this `Msg` to get sent to my `update` function whenever the user clicks here."

Commands work similarly. We don't say "generate a random number right this instant," we say "I want a random number, so please generate one and send it to my `update` function gift-wrapped in a `Msg`." As with `onClick`, we let `update` take it from there—and if we so desire, `update` can return new commands which trigger new calls to `update` when they complete.

Figure 3.7 shows how commands flow through the Elm Architecture.

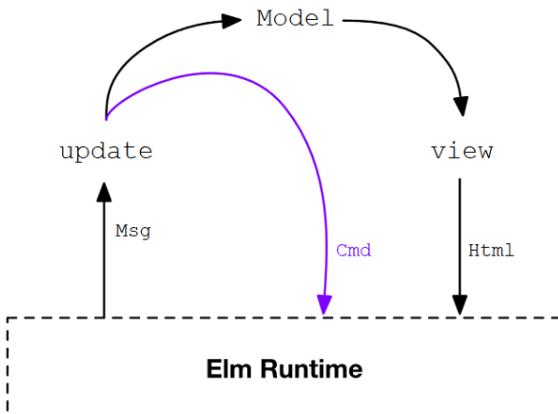


Figure 3.7 How commands flow through the Elm Architecture

Importantly, the addition of commands has not altered the fundamental structure of the Elm Architecture we learned in Chapter 2. All of the following still hold true.

- Our `Model` value is still the single source of truth for the application state.
- Our `update` function is still the only way to alter that model.
- Sending a `Msg` to `update` remains the only way to tell `update` what we want done.

All we've done is to give `update` some new powers: the ability to run logic that can have inconsistent results, such as generating a random number. As we will soon see, the key to unlocking these powers is graduating from `Html.beginnerProgram` to `Html.program`.

RUNNING COMMANDS IN RESPONSE TO USER INPUT

We want to generate our random number in response to a click event, so we'll need to alter our `update` function a bit. Specifically we'll have `update` return a tuple containing not only the new `Model` we want, but also whatever commands we want the Elm Runtime to execute.

Before we edit `update`'s implementation, let's revise its type annotation to guide our work.

```
update : Msg -> Model -> ( Model, Cmd Msg )
```

As you can see, tuple annotations look just like tuple values. `(Model, Cmd Msg)` means "a tuple where the first element is a `Model` and the second is a `Cmd Msg`."

NOTE We write `Cmd Msg` instead of just `Cmd` for the same reason that we write `Html Msg` instead of `Html`.

Like `Html`, the `Cmd` type also has a type variable, and for a familiar reason: whenever a `Cmd` finishes, a `Msg` value is what it sends to `update` for further processing.

Let's revise our `update` implementation to reflect this new reality.

```
update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
  case msg of
    SelectByUrl url ->
      ( { model | selectedUrl = url }, Cmd.none )

    SurpriseMe ->
      ( { model | selectedUrl = "2.jpeg" }, Cmd.none )

    SetSize size ->
      ( { model | chosenSize = size }, Cmd.none )
```

A fine start! Next we'll replace that `Cmd.none` in `SurpriseMe`'s branch with a command to randomly generate the index of the photo we want to select.

3.3.3 Generating Random Values with `Random.generate`

The `Random.generate` function returns a command that generates random values wrapped up in messages. It's just what the doctor ordered here! `Random.generate` takes two arguments:

1. A `Random.Generator`, which specifies the type of random value we want.
2. A function which can wrap the resulting random value in one of our `Msg` values.

The `randomPhotoPicker` we created earlier will do nicely for the `Random.Generator`, so all we're missing is a way to wrap these randomly-generated integers in a `Msg`. We can introduce one by adding a new type constructor to `Msg` and a new branch to `update`:

```
type Msg
  = SelectByUrl String
  | SelectByIndex Int
  ...

update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
  case msg of
    SelectByIndex index ->
      ( { model | selectedUrl = getPhotoUrl index }, Cmd.none )

  ...
```

Nice! Now we just need `Random.generate` to produce one of these `SelectByIndex` messages.

CALLING RANDOM.GENERATE

When a user clicks the *Surprise Me!* button, the `SurpriseMe ->` branch of our `update` function's *case-expression* gets run. Currently that returns a `Cmd.none`, but now we want it to return a different command: one that generates a random `Int` and sends it back to `update` wrapped in a `SelectByIndex` message. We'll get this command by calling `Random.generate`.

Earlier we noted that `Random.generate` needs two ingredients:

- A `Random.Generator`, which specifies what type of random value we want. (We'll use `randomPhotoPicker`, our `Random.Generator Int`, for this.)

- A function which can wrap the resulting random value in one of our `Msg` values. (We'll use our `SelectByIndex` type constructor for this, as it is a function that returns a `Msg`.)

Let's revise our `SurpriseMe` branch to return an unchanged model, and to call `Random.generate` instead of returning `Cmd.none`.

```
SurpriseMe ->
  ( model, Random.generate SelectByIndex randomPhotoPicker )
```

Figure 3.8 illustrates how this would cause data to flow through our application, assuming the Elm Runtime randomly generated a **2** after the user clicked our *Surprise Me!* button.

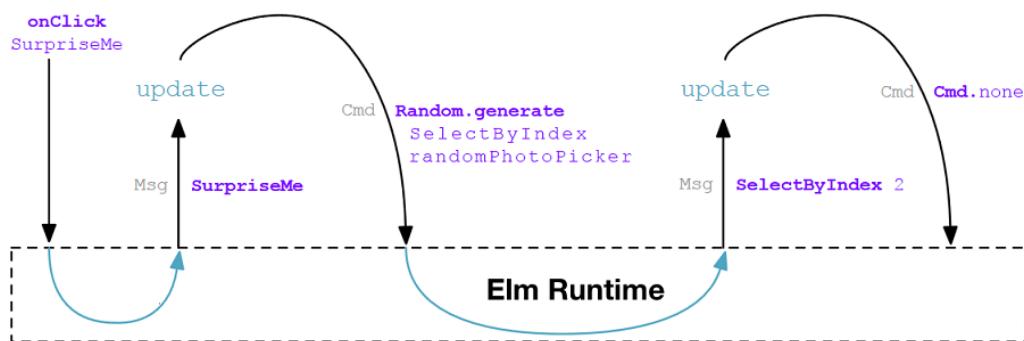


Figure 3.8 Data flowing through our application, assuming the Elm Runtime randomly generated a **2**

Calling `Random.generate SelectByIndex randomPhotoPicker` returns a command which:

1. Randomly generates an `Int` between 0 and 2, because that's the type of random value `randomPhotoPicker` specified it should generate.
2. Takes that randomly-generated `Int`, passes it to `SelectByIndex` (an `Int -> Msg` function), and runs the resulting `Msg` through `update`.

We could write a type annotation for `Random.generate` like so:

```
generate : (randomValue -> msg) -> Random.Generator randomValue -> Cmd msg
```

NOTE Capitalization is important! There is a big difference between `Msg` and `msg`. A function that returns `Msg` returns an instance of the exact `Msg` union type we defined in `PhotoGroove.elm`. In contrast, `msg` is a type variable. A function that returns `msg` could return anything at all!

Table 3.5 shows how the expressions involved in our call to `Random.generate` relate to the parameters and type variables in this annotation.

Table 3.5 Types involved in calling Random.generate

Expression	Type	Type in Original Annotation
SelectByIndex randomPhotoPicker	(Int -> Msg) Random.Generator Int	(randomValue -> msg) Random.Generator randomValue
Random.generate SelectByIndex randomPhotoPicker	Cmd Msg	Cmd msg

NOTE Because it takes a function as an argument, `Random.generate` is another *higher-order function* like `List.map` or `String.filter`.

UPGRADING TO HTML.PROGRAM

We're almost done, but if we compile our code right now, we'll get an error. This is because `Html.beginnerProgram` wants `update` to return a `Model`, whereas our `update` function now returns a `(Model, Cmd Msg)` tuple. Fortunately `beginnerProgram`'s older sibling, `Html.program`, expects just such a tuple! Let's replace our `main` declaration with this:

```
main =
  Html.program
    { init = ( initialModel, Cmd.none )
    , view = view
    , update = update
    , subscriptions = (\model -> Sub.none)
    }
```

We're passing two things differently from what we passed to `Html.beginnerProgram`.

1. The `model` record has been replaced by an `init` tuple. The first element in the tuple is our `initial Model`, and the second is a command to run when the application loads. (`init` is the only place besides `update` that can specify commands to run.) Since we have nothing to run on startup, we write `Cmd.none`.
2. We've added a `subscriptions` field. We'll make use of that in Chapter 5, but for now let's disregard it.

Next we'll give `main` a type annotation. As we've been going through adding annotations, you may have wondered why the compiler hasn't been helping out with that. It clearly knows what the types are, since it infers them for us in `elm-repl`, so why can't it tell us the annotations?

GENERATING ANNOTATIONS WITH THE --WARN FLAG

Actually, it totally can! Try running `elm-make` with the `--warn` flag:

```
elm-make PhotoGroove.elm --output elm.js --warn
```

You should see this warning:

```
Top-level value `main` does not have a type annotation.
```

```
139| main =
     ^^^^
I inferred the type annotation so you can copy it into your code:
main : Program Never Model Msg
```

Thanks, compiler! Go ahead and copy this final annotation right into your code. (Chapter 5 will get into what `Program Never Model Msg` means.) Generated annotations won't always match what you would have written yourself, but you can at least expect them to compile.

TIP This is presented as a compiler warning because it's best practice to document your top-level functions with type annotations. When the compiler can generate these for you, documentation takes much less effort!

At this point you may be wondering something else: why did we go through this whole chapter writing annotations by hand, when we could have had the compiler generate immediately usable ones for us all along? To that I say...*hey look, a distraction!*

THE FINAL BUILD

If you recompile with `elm-make` and open `index.html`, you should experience a gloriously functional application. Try clicking *Surprise Me!* to watch it randomly pick a different photo!

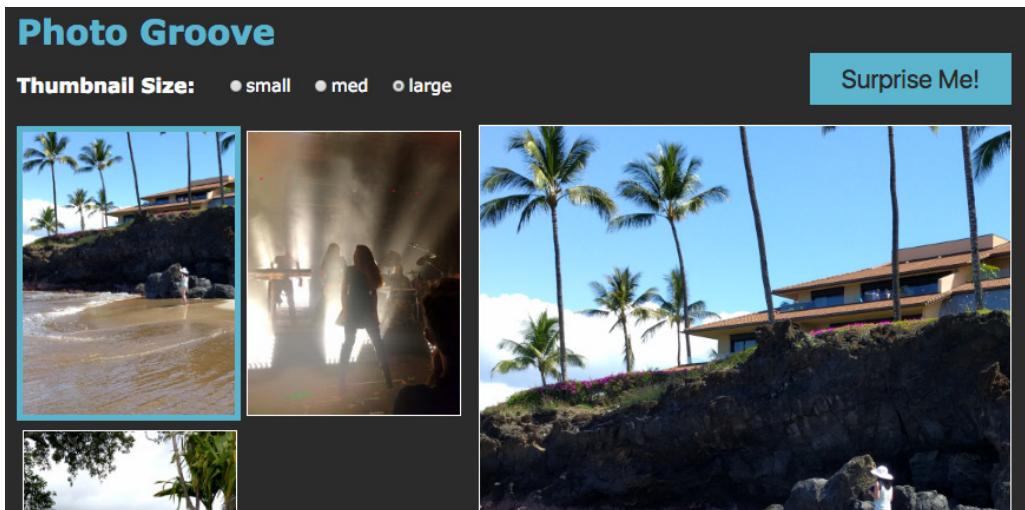


Figure 3.9 Randomly choosing the initially selected photo

Ain't it grand? Even after all the changes we made, once it compiled it just worked—no regressions! Not only that, but our code never came anywhere near causing a runtime exception. This is a normal experience with Elm, yet it never seems to get old.

Compared to where we were at the end of Chapter 2, our revised application is not only more feature-rich, it's also more reliable and better-documented. It does more, it does it better, and the code is cleaner. Huzzah!

3.4 Summary

We've now improved on our application from Chapter 2 in several ways.

- We added documentation in the form of type annotations
- Users can select from one of three thumbnail sizes
- It has a *Surprise Me!* button which selects a thumbnail at random
- From now on whenever we add a new `Msg` value, the compiler will give us a *Missing Patterns* error if we forget to handle it (like it did when we didn't account for `SetSize`)

Along the way we covered many new concepts, such as:

- Type variables represent concrete types that have not been specified yet.
- A `type alias` declaration assigns a name to a type, much like how a constant assigns a name to a value.
- A `type` declaration defines a new union type, a custom type that did not exist before.
- Union types can hold more flexible data than records or tuples can.
- Union type constructors can be either values which are instances of that union type, or functions which return instances of that union type.
- You can destructure union type constructors in *case-expressions* to extract their data.
- If you don't write a fallback `_ ->` branch in a *case-expression*, you'll get a compiler error unless your code handles all possible cases.
- `Array.get` prevents runtime crashes by returning a `Maybe` instead of a normal element.

We also learned some differences between comments and type annotations.

Table 3.6 Documenting Code with Comments and Type Annotations

Comment	Type Annotation
Arbitrary string, can describe anything	Can only describe a thing's type
Can be inaccurate or get out of date	Compiler guarantees it's accurate and up-to-date
Can be written just about anywhere	Always goes on the line above the thing it documents

Finally, we learned about `Maybe`, the container value which `Array.get` returns to represent the potential absence of a value. Table 3.7 compares `Maybe` and `List`.

Table 3.7 Comparing Maybe and List

Container Contents	List	Maybe
Empty	<code>[] : List a</code>	<code>Nothing : Maybe a</code>
One value	<code>["foo"] : List String</code>	<code>Just "foo" : Maybe String</code>
Two values	<code>["foo", "bar"] : List String</code>	<code>not possible!</code>

In Chapter 4 we'll take our application **to the extreme**, by talking to a server to obtain our list of photos to display. Stay tuned!

Listing 3.7 The complete PhotoGroove.elm

```
module PhotoGroove exposing(..)

import Html exposing(..)
import Html.Attributes exposing(..)
import Html.Events exposing(onClick)
import Array exposing(Array)
import Random

urlPrefix : String
urlPrefix =
    "http://elm-in-action.com/"

type ThumbnailSize
    = Small
    | Medium
    | Large

view : Model -> Html Msg
view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , button
            [ onClick SurpriseMe ]
            [ text "Surprise Me!" ]
        , h3 [] [ text "Thumbnail Size:" ]
        , div [ id "choose-size" ]
            (List.map viewSizeChooser [ Small, Medium, Large ])
        , div [ id "thumbnails", class (sizeToString model.chosenSize) ]
            (List.map (viewThumbnail model.selectedUrl) model.photos)
        , img
            [ class "large"
            , src (urlPrefix ++ "large/" ++ model.selectedUrl)
            ]
        []
    ]

viewThumbnail : String -> Photo -> Html Msg
```

```

viewThumbnail selectedUrl thumbnail =
    img
        [ src (urlPrefix ++ thumbnail.url)
        , classList [ ( "selected", selectedUrl == thumbnail.url ) ]
        , onClick (SelectByUrl thumbnail.url)
        ]
    []

viewSizeChooser : ThumbnailSize -> Html Msg
viewSizeChooser size =
    label []
        [ input [ type_ "radio", name "size", onClick (SetSize size) ] []
        , text (sizeToString size)
        ]

sizeToString : ThumbnailSize -> String
sizeToString size =
    case size of
        Small ->
            "small"

        Medium ->
            "med"

        Large ->
            "large"

type alias Photo =
    { url : String }

type alias Model =
    { photos : List Photo
    , selectedUrl : String
    , chosenSize : ThumbnailSize
    }

initialModel : Model
initialModel =
    { photos =
        [ { url = "1.jpeg" }
        , { url = "2.jpeg" }
        , { url = "3.jpeg" }
        ]
    , selectedUrl = "1.jpeg"
    , chosenSize = Medium
    }

photoArray : Array Photo
photoArray =
    Array.fromList initialModel.photos

```

```

getPhotoUrl : Int -> String
getPhotoUrl index =
    case Array.get index photoArray of
        Just photo ->
            photo.url
        Nothing ->
            ""

randomPhotoPicker : Random.Generator Int
randomPhotoPicker =
    Random.int 0 (Array.length photoArray - 1)

type Msg
    = SelectByUrl String
    | SelectByIndex Int
    | SurpriseMe
    | SetSize ThumbnailSize

update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
    case msg of
        SelectByIndex index ->
            ( { model | selectedUrl = getPhotoUrl index }, Cmd.none )

        SelectByUrl url ->
            ( { model | selectedUrl = url }, Cmd.none )

        SurpriseMe ->
            ( model, Random.generate SelectByIndex randomPhotoPicker )

        SetSize size ->
            ( { model | chosenSize = size }, Cmd.none )

main =
    Html.program
        { init = ( initialModel, Cmd.none )
        , view = view
        , update = update
        , subscriptions = (\model -> Sub.none)
        }

```

4

Talking to Servers

This chapter covers

- Using Decoders to validate and translate JSON
- Handling descriptive errors with Results
- Communicating over HTTP with Requests

We've made great progress! Users of our *Photo Groove* application can now do quite a bit:

- Select a thumbnail to view a larger version
- Choose Small, Medium, or Large thumbnail sizes
- Click *Surprise Me!* to select a random thumbnail

Our manager is impressed with our progress, but has even more in mind. "It's time to take *Photo Groove* to the next level. The highest level. The *cloud* level. That's right, we're going to have *Photo Groove* start getting its photo information from our very own servers!"

As with last time, our manager has one more minor feature request. "There's also going to be some metadata associated with each photo—specifically download size and an optional caption. We can have it show those on top of the big photo, right?"

Sure, since you asked so nicely.

In this chapter we'll teach our application how to talk to servers. We'll validate and translate JSON data, communicate over HTTP using the Elm Architecture, and reliably handle errors in client-server communication.

Let's get to it!

4.1 Preparing for Server-Loaded Data

Since we're going to be loading our initial list of photos from the server, our data modeling needs will change.

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

4.1.1 Modeling Incremental Initialization

Right now our `initialModel` looks like this:

```
initialModel : Model
initialModel =
  { photos = [ { url = "1.jpeg" }, { url = "2.jpeg" }, { url = "3.jpeg" } ]
  , selectedUrl = "1.jpeg"
  , chosenSize = Medium
  }
```

We'll be loading our photos from the server, meaning `initialModel.photos` will be `[]` and we'll update that list once we hear back from the server. This will have a few consequences:

1. Because we have no initial photos, we can't have an initial `selectedUrl` either.
2. The initial request to load the data might fail, for example because the server is down. In that case, we should display an error.
3. Our `photoArray` value references `initialModel.photos`, so it will be affected.
4. Both `getPhotoUrl` and `randomPhotoPicker` reference `photoArray`, so they will be affected as well.

Let's translate these new considerations into data model changes!

MODELING ABSENT VALUES WITH MAYBE

The `Maybe` container we learned about in Chapter 3 seems like a nice way to implement the first consideration—that we won't have an initial `selectedUrl`. Since `Maybe` represents the potential absence of a value—by holding either the value `Nothing` or some other value wrapped in a `Just`—it's an intuitive way to represent the data that is absent until we hear back from the server.

We can also use a `Maybe` to represent a potential loading error message. It will begin as `Nothing`, but might become something like `Just "Server unavailable"` if we get an error.

Go ahead and change `Model` and `initialModel` to look like the following.

```
type alias Model =
  { photos : List Photo
  , selectedUrl : Maybe String
  , loadingError : Maybe String
  , chosenSize : ThumbnailSize
  }

initialModel : Model
initialModel =
  { photos = []
  , selectedUrl = Nothing
  , loadingError = Nothing
  , chosenSize = Medium
  }
```

Now that we've revised our data model, we can lean on the compiler to tell us what parts of our code base were affected by this revision. This will let us reliably propagate these changes throughout our code base!

PROPAGATING CHANGES

If we recompile we'll see two type mismatches: one in `view` and one in `update`.

Since `selectedUrl` is now a `Maybe String` instead of a `String`, our `SelectByUrl` logic inside `update` now needs a `Just`. Let's add it:

```
SelectByUrl url ->
  ( { model | selectedUrl = Just url }, Cmd.none )
```

Our `SelectByIndex` logic has a similar type mismatch, but there we can actually do better! One way to fix the mismatch is to wrap what `getPhotoUrl` returns in a `Just`...but another way is to propagate the `Maybe` itself! Let's rewrite `getPhotoUrl` to return a `Maybe`:

```
getPhotoUrl : Int -> Maybe String
getPhotoUrl index =
  case Array.get index photoArray of
    Just photo ->
      Just photo.url

    Nothing ->
      Nothing
```

Notice how this not only resolves the type mismatch, but also means we are no longer returning the clearly invalid URL of "" in the case where the index is out of bounds. One less potential bug to worry about!

FIXING VIEWTHUMBNAIL

Now that we've resolved our type mismatches in `update`, we can move on to resolving the ones in `view`. One of the problems in `view` is that the `viewThumbnail` function it calls has an inaccurate type annotation. It says it's taking a `String` called `selectedUrl`, but `selectedUrl` is now a `Maybe String`. Let's fix that annotation real quick:

```
viewThumbnail : Maybe String -> Photo -> Html Msg
viewThumbnail selectedUrl thumbnail =
```

There's another problem in `viewThumbnail`: the code `selectedUrl == thumbnail.url` is now comparing a `String` to a `Maybe String`, which is a type mismatch! The easiest way to fix this is by calling `Just` on `thumbnail.url`, converting it to a `Maybe String`. Now we're comparing a `Maybe String` to a `Maybe String` and balance has been restored to the galaxy.

```
classList [ ( "selected", selectedUrl == Just thumbnail.url ) ]
```

FIXING THE PREVIEW IMAGE

With those “quick fixes” out of the way, we can recompile to see one final error: the code that renders the large preview image has a type mismatch.

```
src (urlPrefix ++ "large/" ++ model.selectedUrl)
```

This case is not a “quick fix” like the others were. If `model.selectedUrl` is `Nothing` here, we really don’t want to render this `img` at all!

Let’s pull the logic to view the large photo into its own function, and handle the `Maybe` in there by declining to render the `img` at all.

```
view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        ... (everything in between stays the same) ...
        , div [ id "thumbnails", class (sizeToString model.chosenSize) ]
            (List.map (viewThumbnail model.selectedUrl) model.photos)
        , viewLarge model.selectedUrl
    ]

viewLarge : Maybe String -> Html Msg
viewLarge maybeUrl =
    case maybeUrl of
        Nothing ->
            text ""
        Just url ->
            img [ class "large", src (urlPrefix ++ "large/" ++ url) ] []
```

Notice that we’re returning `text ""` if `maybeUrl` is `Nothing`. This “render an empty text node” technique is a quick way to obtain some `Html` that has no visual impact on the page.

Let’s recompile everything with `elm-make --output=elm.js PhotoGroove.elm` and open `index.html` to see our progress.

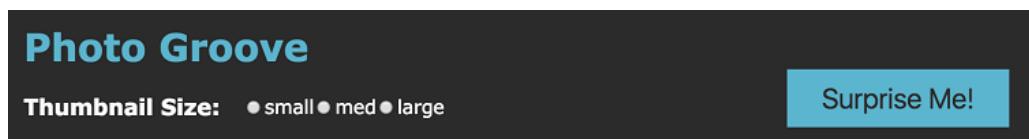


Figure 4.1 The updated application, with no photos displaying

Excellent! We’ve achieved that peculiar flavor of success where we smile at our desk...while hoping nobody walks by and sees how the previously-working application looks right now.

4.1.2 Resolving Data Dependencies

Now that we’ve resolved our type mismatches, let’s turn to our business logic.

`photoArray` is still based on `initialModel.photos`. This is perfectly reasonable if `initialModel.photos` has a bunch of photos in it...but now that it's hardcoded to `[]`? Not so much. That means `photoArray` will always be empty, which means even once we've loaded some photos to display, our *Surprise Me!* button won't be able to randomly select one.

FIXING SELECTBYINDEX

To fix this, we need to make our `SelectByIndex` logic reference the *current* photos, rather than the hardcoded ones we had at startup. We can do this by deleting `photoArray` and `getPhotoUrl`, and changing the `SelectByIndex` branch of our `update` function like so:

```
SelectByIndex index ->
let
    newSelectedPhoto : Maybe Photo
    newSelectedPhoto =
        Array.get index (Array.fromList model.photos)

    newSelectedUrl : Maybe String
    newSelectedUrl =
        case newSelectedPhoto of
            Just photo ->
                Just photo.url

            Nothing ->
                Nothing
in
    ( { model | selectedUrl = newSelectedUrl }, Cmd.none )
```

Our `newSelectedUrl` implementation fits a pattern that comes up fairly often:

1. Run a *case-expression* on a `Maybe`
2. In the `Nothing` branch, evaluate to `Nothing`
3. In the `Just` branch, evaluate to another `Just`

In other words, if we have a `Nothing`, we leave it alone, but if we have a `Just`, we transform it into a different `Just`.

REFACTORING TO USE MAYBE.MAP

The `Maybe.map` function provides a way to perform this transformation more concisely. Let's refactor our `getPhotoUrl` implementation to use it.

```
newSelectedUrl : Maybe String
newSelectedUrl =
    Maybe.map (\photo -> photo.url) newSelectedPhoto
```

You can read this code as "Look at the result of `Array.get index photoArray`. If it's `Nothing`, then return `Nothing`, but if it's `Just photo`, return `Just photo.url` instead."

`Maybe.map` does this by accepting a function which transforms the old `Just` value into the new one. (If the `Maybe` is `Nothing`, this `Just`-transforming function does not get called.)

Table 4.1 illustrates the similarities between `Maybe.map` and our old friend `List.map`.

Table 4.1 Comparing List.map and Maybe.map

List.map	Maybe.map
List.map negate [39] == [-39]	Maybe.map negate (Just 39) == Just -39
List.map negate [5] == [-5]	Maybe.map negate (Just 5) == Just -5
List.map negate [] == []	Maybe.map negate Nothing == Nothing

This code is now significantly more concise than when we used an entire *case-expression* to implement the same logic...*but wait! There's more!*

REFACTORING TO USE .URL

Anytime you see a function that takes a record and immediately return one of its fields—as we're doing with `(\photo -> photo.url)` above—you can take advantage of a special syntax Elm offers for these occasions. If we replace `(\photo -> photo.url)` with `.url` instead, it will do exactly the same thing.

That means we can delete `newSelectedPhoto` and refactor `newSelectedUrl` to this:

```
newSelectedUrl : Maybe String
newSelectedUrl =
    Maybe.map .url (Array.get index (Array.fromList model.photos))
```

TIP You can think of `.url` as “a function which takes a record and returns its `.url` field.”

We've now replaced a *case-expression* with one function call, but the result is not the nicest to read. Fortunately, Elm provides an excellent way to make expressions like this more readable!

USING THE PIPELINE OPERATOR

We can express `newSelectedUrl` as a “pipeline” of sequential operations by using the `|>` operator like so:

```
newSelectedUrl : Maybe String
newSelectedUrl =
    model.photos
        |> Array.fromList
        |> Array.get index
        |> Maybe.map .url
```

This expression is exactly equivalent to writing `Maybe.map .url (Array.get index (Array.fromList model.photos))` except it's expressed in a different style.

Figure 4.2 compares these two ways to write the same expression.

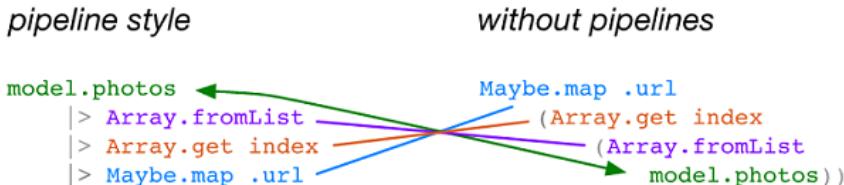


Figure 4.2 The same expression, with and without pipelines

The pipelined code is saying the following:

1. Start with `model.photos`
2. Pass it to `Array.fromList`
3. Pass the previous step's return value as the final argument to `Array.get index`
4. Pass *that* previous step's return value as the final argument to `Maybe.map .url`

Pipelines make it easier to read and to modify sequences of transformations like this. We'll be using them quite often throughout the rest of the book!

NOTE Before Elm's compiler generates JavaScript code, it quietly rewrites pipelines into normal function calls. So `"foo" |> String.reverse` compiles to the same thing as `String.reverse "foo"`. Since they compile to the same thing, there's no performance cost to choosing one over the other. Choose whichever makes for nicer code!

Go ahead and revise our `SelectByIndex` logic to use this pipeline style, like so:

```

SelectByIndex index ->
let
    newSelectedUrl : Maybe String
    newSelectedUrl =
        model.photos
            |> Array.fromList
            |> Array.get index
            |> Maybe.map .url
in
    ( { model | selectedUrl = newSelectedUrl }, Cmd.none )

```

INLINE RANDOMPHOTOPICKER

Now we have one remaining problem: `randomPhotoPicker` is still referencing the `photoArray` value we deleted.

Fortunately, it was only using `photoArray` to get its length inside the `SurpriseMe` branch of `update`'s *case-expression*. We can get that same length value using `List.length model` at that point in the code, so we can delete the top-level `randomPhotoPicker` in favor of this:

```

SurpriseMe ->
let
    randomPhotoPicker =
        Random.int 0 (List.length model.photos - 1)

```

```
in
( model, Random.generate SelectByIndex randomPhotoPicker )
```

Great! Now that we've installed the requisite package and prepared our data model for the transition, we're ready to get some HTTP going!

4.2 Fetching Data from a Server

One of our helpful coworkers has set up a simple server endpoint which returns a comma-separated list of photo filenames. The list doesn't yet contain the metadata that we'll ultimately need, but it'll give us a good starting point toward that goal.

Let's send a `HTTP GET` request to fetch that comma-separated string from a server!

4.2.1 Describing HTTP Requests

We'll be using the `Http` module to send HTTP requests to our server. Since `Http` is not one of Elm's core modules, we need to install the `elm-lang/http` package to gain access to it:

```
elm-package install elm-lang/http
```

In Chapter 6 we'll learn more about package management, but for now let's keep moving!

MANAGED EFFECTS INSTEAD OF SIDE EFFECTS

In Chapter 3 we saw how generating random numbers in Elm must be done with a `Cmd` rather than a plain function call.

Whereas JavaScript's `Math.random()` can return a different random number each time you call it, Elm functions must be more consistent. When you pass an Elm function the same arguments, it's guaranteed to return the same value.

There's another rule that applies to all Elm functions, and this one affects the HTTP requests we're about to make. The rule is that Elm functions cannot have *side effects*.

DEFINITION An *effect* is an operation that modifies external state. A function which modifies external state when it executes has a *side effect*.

HTTP requests can always modify external state—since even a `GET` request can result in a server changing values in a database—so performing an HTTP request is an effect. This means if we execute a function and it performs an HTTP request, that function has a side effect.

Listing 4.1 shows a few JavaScript functions which consistently return the same value, but which have side effects because they modify external state along the way. None of these could be written as plain Elm functions.

Listing 4.1 JavaScript functions which have side effects

```
function storeStuff() {
  localStorage.stuff = "foo";
```

```

    return 1;          1
}

function editField(object) { 2
  object.foo = "bar";
  return object;
} 2

var something = 1;

function editVar() { 3
  something = 2; 3
  return 0; 3
}
  
```

- ① modifies the contents of localStorage
- ② modifies the object it receives
- ③ modifies an external variable

If you look through all the code we've written for *Photo Groove*, you won't find a single function that has a side effect. All effects are performed by the Elm Runtime itself; our code only **describes** which effects to perform, by returning values from `update`.

NOTE Calling `update` does not directly alter any state! All `update` does is return a tuple. If you wanted to, you could call `update` a hundred times in a row, and all it would do is give you a hundred tuples.

This system of *managed effects*, where the Elm Runtime is in charge of performing all effects, means that Elm programs can be written entirely in terms of data transformations. In Chapter 6 we'll see how nice this makes Elm code to test.

This means that in order to perform a HTTP request, we'll do what we did with random number generation in Chapter 3: use a `Cmd` to tell the Elm runtime to perform that effect.

DESCRIBING REQUESTS WITH `HTTP.GETSTRING`

The `Http.send` function returns a `Cmd` representing the HTTP request we want to make.

To configure this request, we pass `Http.send` an `Http.Request` value describing what the request should do. This `Request` value serves a similar purpose to the `Random.Generator` value we passed to `Random.generate` in Chapter 3—except where `Random.Generator` described the random value we wanted, `Request` describes the HTTP request we want.

One way we can obtain a `Request` value is with the `Http.getString` function, like so:

```
Http.getString "http://elm-in-action.com/photos/list"
```

The `getString` function has this type:

```
getString : String -> Request String
```

We pass it a URL string, and it returns a `Request` that asks the Elm Runtime to send a `HTTP GET` request to that URL. If the request succeeds, we'll end up with a plain `String` containing the response body the server sent back.

TRANSLATING REQUESTS INTO COMMANDS WITH `HTTP.SEND`

Now that we have our `Request` `String`, we can pass it to `Http.send` to get a `Cmd` `Msg`. Just like in Chapter 3 with random number generation, once this `Cmd` completes, it will send a `Msg` to update telling us what happened.

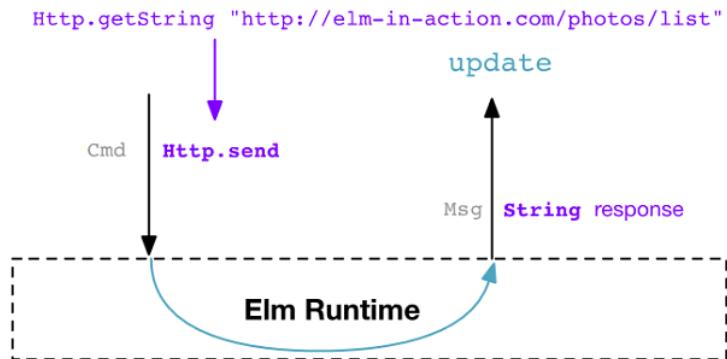


Figure 4.3 The Elm Runtime executing a `Cmd` from `Http.send`

Unlike random number generation, a lot can go wrong with a HTTP request! What if we send the server an invalid URL? Or an unrecognized one? What if the server has been unplugged by a mischievous weasel? If the request fails for any reason, we'll want to know went wrong so we can inform the user.

`Http.send` produces a value that accounts for this possibility: a `Result`.

4.2.2 Handling HTTP Responses

To get a feel for how `Result` works, let's compare it to its cousin `Maybe`.

In Chapter 3 we saw how `Array.get` returned a `Maybe` in order to account for the possibility of a missing value. The definition of `Maybe` looks like this:

```
type Maybe value
  = Just value
  | Nothing
```

Elm has a similar union type for representing the *result* of an operation that can fail, such as performing a HTTP request. It's called `Result`, and it looks like this:

```
type Result errValue okValue
  = Err errValue
  | Ok okValue
```

As we can see, `Result` has two type variables:

- `okValue` represents the type of value we'll get if the operation succeeded. For example, if we are trying to obtain a `String` from the server, then `okValue` would be `String`.
- `errValue` represents the type of error we'll get if the operation failed. In our case, `errValue` will refer to a union type that enumerates various mishaps that can befall HTTP requests. That union type is called `Http.Error`.

Similarly to how `Maybe` requires that we write logic to handle both the `Just` and `Nothing` cases, so too does `Result` require that we handle both the `Ok` and `Err` cases.

HANDLING THE RESULT

Let's start by importing the `Http` module we're about to begin using.

```
import Http
```

Next let's add a `Msg` type constructor to handle the `Result` value we'll get from `Http.send`:

```
type Msg
  = SelectByUrl String
  | SelectByIndex Int
  | SurpriseMe
  | SetSize ThumbnailSize
  | LoadPhotos (Result Http.Error String)
```

We'll learn more about that `Http.Error` type in a bit. For now, let's add a branch to our `update` function's *case-expression* for this new `Msg`. We'll start with something like this:

```
LoadPhotos result ->
  case result of
    Ok responseStr ->
      ...translate responseStr into a List of Photos for our Model...
    Err httpError ->
      ...record an error message here...
```

With that basic skeleton in place, we can implement these two cases one at a time. We'll start with the `Ok` `responseStr` case, where the server gave us a valid response and all is well.

READING PHOTO FILENAMES WITH STRING.SPLIT

Remember earlier when we mentioned how we've got a server endpoint which returns a comma-separated list of photo filenames? If not, here it is again:

One of our helpful coworkers has set up a simple server endpoint which returns a comma-separated list of photo filenames.

—This book, earlier.

Sure enough, if we visit this endpoint at <http://elm-in-action.com/photos/list> we can see the list of filenames. It's a plain old string:

```
"1.jpeg,2.jpeg,3.jpeg,4.jpeg"
```

When we wrote `Ok responseStr` a moment ago, that `responseStr` value will refer to exactly this string! Now we can split `responseStr` into a list of individual filenames using Elm's `String.split` function. It has this type:

```
split : String -> String -> List String
```

We give `String.split` a separator (in this case a comma) and a string to split, and it gives us back a list of all the strings it found between those separators.

Table 4.2 shows what `String.split ",,"` will do to the string we get from the server.

Table 4.2 Calling `String.split ",,"` on a string

stringFromServer	<code>String.split ",," stringFromServer</code>
<code>"1.jpeg,2.jpeg,3.jpeg,4.jpeg"</code>	<code>["1.jpeg", "2.jpeg", "3.jpeg", "4.jpeg"]</code>

Once we have this list of filename strings, we can use it to set the `photos : List Photo` field in our `Model` like in Listing 4.2.

Listing 4.2 Incorporating `responseStr` into the Model

```
Ok responseStr ->
    let
        urls =
            String.split ",," responseStr           ①

        photos =
            List.map (\url -> { url = url }) urls  ②
    in
        ( { model | photos = photos }, Cmd.none ) ③
```

- ① Split the String into a List String
- ② Translate the List String into a List Photo
- ③ Set the photos in the model

Now when we successfully receive a `String` response from the server, we'll translate it into a list of `Photo` records and store them in our `Model`.

USING TYPE ALIASES TO CREATE RECORDS

Declaring type alias `Photo = { url : String }` does more than giving us a `Photo` type we can use in type annotations. It also gives us a convenience function whose job is to build `Photo` record instances. This function is also called `Photo`. Here it is in action:

```
Photo "1.jpeg" == { url = "1.jpeg" }
```

This also works with record type aliases involving multiple fields, like the one for `Model`:

```
type alias Model =
    { photos : List Photo
    , selectedUrl : String
    , chosenSize : ThumbnailSize
    }
```

This declaration gives us a convenience function called `Model` which builds a record and returns it. Since this record has three fields where `Photo` had only one, the `Model` function accepts three arguments instead of one. The type of the `Model` function looks like this:

```
Model : List Photo -> String -> ThumbnailSize -> Model
```

The order of arguments matches the order of the fields in the `type alias` declaration. So if you were to move the `photos : List Photo` declaration to the end of the `type alias`, then the `Model` function would look like this instead:

```
Model : String -> ThumbnailSize -> List Photo -> Model
```

We can use this knowledge to perform a quick refactor of our `Photo` construction:

```
Old: List.map (\url -> { url = url }) urls
New: List.map Photo urls
```

Lovely!

HANDLING ERRORS (OR NOT)

What about when things don't go as planned? We'll add some nice error handling later, but first we want to get something up and running. Elm won't let us forget to handle the error case—if we omitted the `Err` branch, we'd get a `MISSING PATTERNS` compile error!—but the compiler won't stop us from choosing to do nothing when we receive an `Err`.

```
Err httpError ->
    ( model, Cmd.none )
```

Since we aren't using our `httpError` value, let's rename it to an underscore:

```
Err _ ->
    ( model, Cmd.none )
```

The underscore parameter

The underscore is a special parameter name whose purpose is to signify that there is a parameter here, but we're choosing not to use it. Unlike a normal parameter, attempting to reference `_` in our logic would be a compile error.

One handy feature of `_` is that you can use it multiple times in the same parameter list. For example:

```
functionThatTakesThreeArguments _ _ _ =
    "I ignore all three of my arguments and return this string!"
You can use _ either in function parameter lists or in case-expression branches.
```

Later on we'll circle back to this branch of our *case-expression* and use it to show the user a nice error message when an HTTP error happens.

PATTERN MATCHING

Before we move on, let's pause for a tasty refactor. If you zoom out a bit, you'll notice we have a *case-expression* nested directly inside another *case-expression*.

```
case msg of
  ...
  LoadPhotos result ->
    case result of
      Ok responseStr ->
        ...
        Err _ ->
          ( model, Cmd.none )
```

In situations like this, we can use concise *pattern matching* to express the same logic:

```
case msg of
  ...
  LoadPhotos (Ok responseStr) ->
  ...
  LoadPhotos (Err _) ->
    ( model, Cmd.none )
```

DEFINITION *Pattern matching* is a way of destructuring values based on how their containers look. In the example above, if we have a `LoadPhotos` containing an `Ok` containing a value, that value will go into a variable called `responseStr`.

This refactored code is equivalent to what we had before. The difference is that now each of our branches is expressing two conditions at once:

- The `LoadPhotos (Ok responseStr)` branch runs if `msg` is a `LoadPhotos` type constructor which contains an `Ok` value.
- The `LoadPhotos (Err _)` branch runs if `msg` is a `LoadPhotos` type constructor and it contains an `Err` value.

You can nest patterns like these as much as you want, even assembling elaborate creations like `NestAllTheThings (Just (Ok (Listen (This (IsReally "great" _ _)))))` - but try not to overdo it.

We'll encounter other pattern matching techniques as we progress through the book.

4.2.3 Sending HTTP Requests

At this point we've assembled all the ingredients necessary to launch our HTTP rocketship in the direction of our server:

- The Request String returned by `Http.getString "http://elm-in-action.com/photos/list"`
- The `LoadPhotos` message that will handle the response

The `Http.send` function will wire these up for us! It has this type:

```
send : (Result Http.Error val -> msg) -> Request val -> Cmd msg
```

`Http.send` has two type variables: `msg` and `val`. We can infer what they'll become in our case:

- Since we're passing a `Request String`, the `val` type variable must become `String`
- Since we need it to return a `Cmd Msg`, the `msg` type variable must become `Msg`

Knowing this, we can tell that the first argument to `send` is a function that looks like this:

```
(Result Http.Error String -> Msg)
```

This means we could declare a new value called `initialCmd` using this call to `Http.send`:

```
initialCmd : Cmd Msg
initialCmd =
    Http.send
        (\result -> LoadPhotos result)
        (Http.getString "http://elm-in-action.com/photos/list")
```

We'll use this `initialCmd` value to run our HTTP request when the program starts up.

This will compile, but we can simplify it! Back in Section 2.2.1 of Chapter 2, we noted that an anonymous function like `(\foo -> bar baz foo)` can always be rewritten as `(bar baz)` by itself. This means we can replace `(\result -> LoadPhotos result)` with `LoadPhotos` like so:

```
initialCmd : Cmd Msg
initialCmd =
    Http.send LoadPhotos
        (Http.getString "http://elm-in-action.com/photos/list")
```

USING PIPELINE STYLE FOR INITIALCMD

By the way, remember how earlier we refactored `newSelectedUrl` to use a pipeline style instead of multiple constants? We can use the same technique to express `initialCmd` in terms of the steps needed to build it up:

```
initialCmd : Cmd Msg
initialCmd =
    "http://elm-in-action.com/photos/list"
        |> Http.getString
        |> Http.send LoadPhotos
```

Nice! Let's add this version of `initialCmd` to `PhotoGroove.elm` right above our `main` declaration.

RUNNING A COMMAND ON INIT

Now we have a `Cmd` to run, but we want to kick this one off a bit differently than last time. We ran our random number generation command from `update` in response to a user click, but we want to run this command right when the program starts up.

We can do that by updating `main` to use `initialCmd` instead of `Cmd.none` on `init`, like so:

```
main : Program Never Model Msg
main =
    Html.program
        { init = ( initialModel, initialCmd )
        , view = view
        , update = update
        , subscriptions = _ -> Sub.none ①
        }
```

① We don't use this argument, so name it `_`

While we're at it, we also updated `subscriptions` to use `_` for its anonymous function's argument, since that argument never gets used.

TIP An anonymous function which ignores its argument (like the `_ ->` above) looks like a hockey stick.

Figure 4.4 illustrates how our revised `init` kicks off the sequence of events leading up to our `update` function receiving a `LoadPhotos` message.

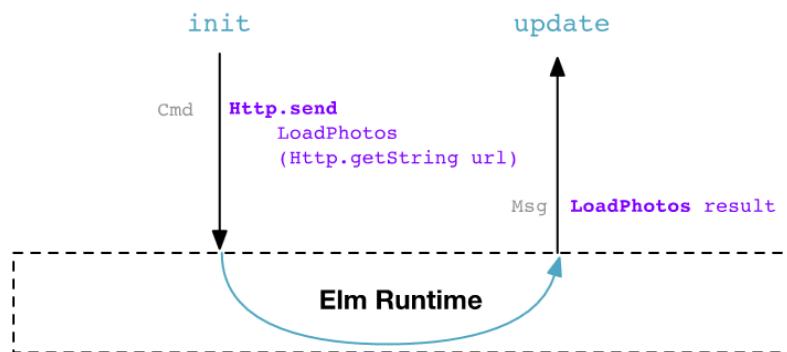


Figure 4.4 Running a `LoadPhotos` command on init

Now let's recompile and open `index.html`, so we can bask in the glory of our new server-loaded photos!

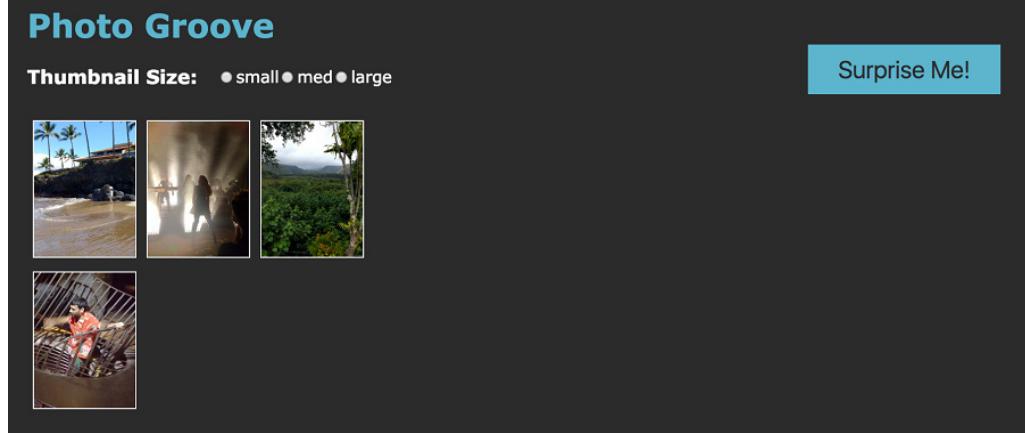


Figure 4.5 The updated application, with an additional photo loaded

Excellent! However, we've lost something along the way: we no longer select the first photo like we did before.

SELECTING AN INITIAL PHOTO USING LIST.HEAD

We'd like to select the first URL that comes back from the server, whatever that URL may be. The `List.head` function is perfect for this! Its type looks like this:

```
head : List elem -> Maybe elem
```

`List.head` someList and `Array.get 0` someArray serve essentially the same purpose: returning `Nothing` if run on an empty collection, or `Just whateverElementWasFirst` if the collection had anything in it. This means we can grab the first URL out of the list like so:

```
List.head urls
```

Since we're going to set `selectedUrl` to this value, and `selectedUrl` is already a `Maybe String`, we don't need run a *case-expression* on this `Maybe String` to do any special handling. We can drop it in and let whoever works with `selectedUrl` handle the `Nothing` case.

Let's expand the record update in the `LoadPhotos (Ok responseStr)` branch to set not only the `photos` field but the `selectedUrl` field as well:

```
LoadPhotos (Ok responseStr) ->
let
  urls =
    String.split "," responseStr
  photos =
    List.map Photo urls
in
  ( { model
```

```

        | photos = photos
        , selectedUrl = List.head urls
    }
, Cmd.none
)
}

```

Now we can try it out and see how this looks. Recompile and...

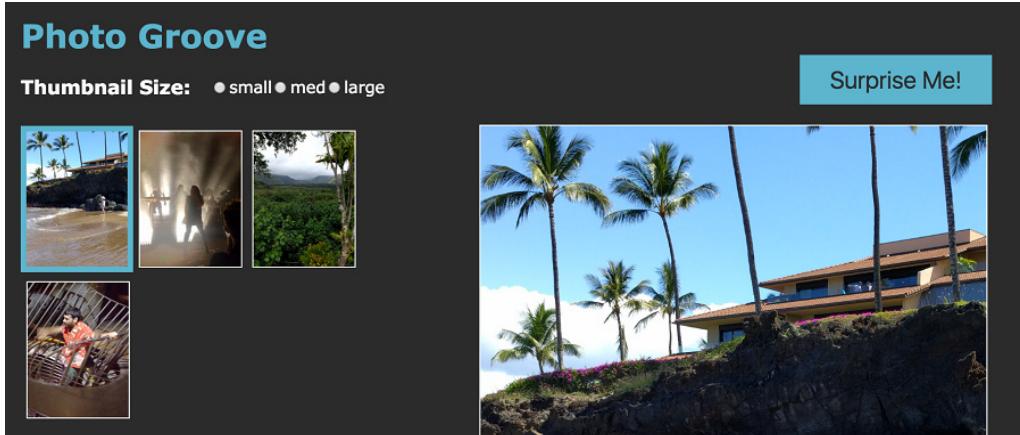


Figure 4.6 Selecting the first loaded photo by default.

Beautiful! Now that we're loading photos from the server, we're ready to take the final step: obtaining the metadata and complete URLs from the server rather than just the filenames.

Now that we've gotten things working in the case where the response successfully comes back, let's revisit the case where everything is not all sunshine and kittens. It's time for some proper error handling!

4.2.4 Gracefully Handling Errors

Let's make it so that if there's an error, we render a helpful message for our users. To do this, we'll need to make two changes:

1. Have `update` set this error message when it receives a `ReportError` message.
2. Have `view` render this error message if it's present.

Let's start with the change to `update`:

```

LoadPhotos (Err _) ->
( { model
  | loadingError = Just "Error! (Try turning it off and on again?)"
  }
, Cmd.none
)

```

NOTE We need the `Just` here because we're storing `loadingError` in our Model as a `Maybe String`, whereas `ReportError` holds a plain `String`.

Next let's introduce a brand-new function which renders the error message, and then modify `Html.program` to use it instead of using `view` directly:

```
viewOrError : Model -> Html Msg
viewOrError model =
  case model.loadingError of
    Nothing ->
      view model

    Just errorMessage ->
      div [ class "error-message" ]
        [ h1 [] [ text "Photo Groove" ]
        , p [] [ text errorMessage ]
        ]

main : Program Never Model Msg
main =
  Html.program
    { init = ( initialModel, initialCmd )
    , view = viewOrError }
```

Whoa! Can we do that? Can we pass a function other than `view` to `Html.program` as the `view` function to use?

Well, sure! All `Html.program` needs there is a function of type `Model -> Html Msg`. As long as we give it a function of that type, `Html.program` could care less what that function is called.

We've done several refactors where we split smaller functions out of `view`, but this is the first time we've wrapped `view` in another function. And it was no big deal! One of the nice things about working with plain old Elm functions all the time is that they're easily interchangeable as long as the types line up.

VERIFYING THAT ERROR HANDLING WORKS

Let's start by making sure the error handling works properly. Change `initialCmd` to fetch data from an invalid URL:

```
initialCmd : Cmd Msg
initialCmd =
  "http://elm-in-action.com/breakfast-burritos/list"
    |> Http.getString
    |> Task.perform handleLoadFailure handleLoadSuccess
```

Now when you recompile and open `index.html`, you'll see our new error message:

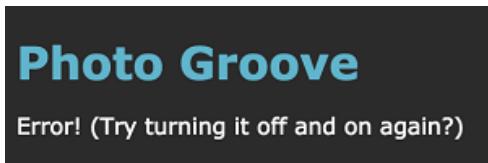


Figure 4.7 Displaying an error as a status message

Now let's change `initialCmd` back to `"http://elm-in-action.com/photos/list"` and press on toward receiving that photo metadata!

4.3 Decoding JSON

Browsers can talk to servers using many different formats. JSON, XML, RPC, BBQ...the list goes on. Regardless of the format, when the browser sends off a HTTP request to a server, the response comes back as a raw string. It's up to the program running in the browser to decode that raw string into a useful representation.

The server with the metadata will be sending us that information via JSON. This means we'll need some way to decode a JSON string into something more useful!

4.3.1 Decoding JSON Strings into Results

Plenty of things can go wrong during the process of decoding JSON. Suppose the server messes up and sends us XML instead of JSON. Our JSON-decoding logic will certainly fail! Or suppose we get back a response that's valid JSON, but a required field is missing. What then?

THE DECODESTRING FUNCTION

The `Json.Decode.decodeString` function returns a `Result`. Its complete type is:

```
decodeString : Decoder val -> String -> Result String val
```

Notice that the `Result` we get back from `decodeString` has one concrete type—namely that `errValue` is always a `String`—and one that matches the type of the given decoder's type variable. Table 4.3 shows how the decoder passed to `decodeString` affects its return type.

Table 4.3 How the Decoder passed to `decodeString` affects its return type

Decoder passed in	<code>decodeString</code> returns	Example success value	Example failure value
<code>Decoder Bool</code>	<code>Result String Bool</code>	<code>Ok True</code>	<code>Err "nope"</code>
<code>Decoder String</code>	<code>Result String String</code>	<code>Ok "Win!"</code>	<code>Err "nah"</code>
<code>Decoder (List Int)</code>	<code>Result String (List Int)</code>	<code>Ok [1, 2, 3]</code>	<code>Err "argh!"</code>

DECODING PRIMITIVES

The `Json.Decode` module has a Decoder `Bool` called `bool`, which translates a JSON boolean string (either "true" or "false") into an Elm `Bool` (`True` or `False`). Let's try it in `elm-repl`!

Listing 4.1 Using `decodeString bool`

```
> import Json.Decode exposing(..)

> decodeString Bool "true"
Ok True : Result String Bool

> decodeString Bool "false"
Ok False : Result String Bool

> decodeString Bool "42"
Err "Expecting a Bool but instead got: 42" : Result String Bool

> decodeString Bool "@&!*/%?"
Err "Given an invalid JSON: Unexpected token @" : Result String Bool
```

NOTE `elm-repl` prints types in fully-qualified style to avoid ambiguity, meaning it always prints `Result.Result` instead of `Result`. These examples will use the unqualified `Result` for brevity, which means they won't have the quite same output as what you'll see in your own `elm-repl`.

Besides `bool`, the `Json.Decode` module offers other primitive decoders like `string`, `int`, and `float`. They work similarly to `bool`, as we can see in `elm-repl`:

Listing 4.2 Using `decodeString` with `int`, `float`, and `string`

```
> import Json.Decode exposing(..)

> decodeString float "3.33"
Ok 3.33 : Result String Float

    > decodeString string "\"backslashes escape quotation marks\""
Ok "backslashes escape quotation marks" : Result String String

> decodeString int "76"
Ok 76 : Result String Int

> decodeString int "3.33"
Err "Expecting an Int but instead got: 3.33" : Result String Int
```

The only primitive decoders are `bool`, `int`, `float`, `string`, and `null`, since JavaScript's `undefined` is not allowed in valid JSON. As we will soon see, there are more popular ways of handling `null` than with this primitive decoder.

4.3.2 Decoding JSON Collections

We've now seen how to decode primitives like booleans, integers, and strings from JSON into their Elm counterparts. This is a good start, but JSON also supports arrays and objects, and we'll need to decode both of these types in order to receive our list of photos with metadata.

DECODING JSON ARRAYS INTO LISTS

Suppose we have the JSON string "[true, true, false]". To decode a list of booleans from this JSON array, we can write `list Bool`. This will give us a `Decoder (List Bool)` value.

Here's how `Json.Decode.bool` and `Json.Decode.list` compare:

```
bool : Decoder Bool
list : Decoder value -> Decoder (List value)
```

Whereas `bool` is a decoder, `list` is a **function** that takes a decoder and returns a new one. We can use it in `elm-repl` to make decoders for lists of primitives, or even lists of lists!

Listing 4.2 Using Json.Decode.list

```
> import Json.Decode exposing (list, bool, string, int)

> list
<function> : Decoder a -> Decoder (List a) ①

> list bool
Decoder (List Bool)

> list string
Decoder (List String)

> list (list int)
Decoder (List (List Int))
```

① elm-repl chooses "a" for the type variable we called "value"

DECODING OBJECTS

The simplest way to decode an object is with the `field` function. Suppose we write this:

```
decoder : Decoder String
decoder =
    field "email" string
```

When this decoder runs, it performs three checks:

1. Are we decoding an Object?
2. If so, does that Object have a field called `email`?
3. If so, is the Object's `email` field a String?

If all three are true, then decoding succeeds with the value of the Object's `email` field.

Table 4.4 shows how this decoder would work on a variety of inputs.

Table 4.4. Decoding objects with field decoders

Decoder	JSON	Result
field "email" string	5	Err "Expected object, got 5"
	{"email": 5}	Err "Expected string for 'email', got 5"
	{"email": "cate@nolf.com"}	Ok "cate@nolf.com"

DECODING MULTIPLE FIELDS

Building a decoder for a single `field` is all well and good, but typically when decoding objects, we care about more than one of their fields. How do we do that?

The simplest is with a function like `map2`, which we can see in Table 4.5.

Table 4.5. Decoding objects with field decoders

Decoder	JSON	Result
<code>map2</code> (\x y -> (x, y)) (field "x" int) (field "y" int)	{"width": 5}	Err "Missing 'y' field"
	{"x": 5, "y": null}	Err "Expected int for 'y', got null"
	{"x": 5, "y": 12}	Ok (5, 12)

DECODING MANY FIELDS

The photo information we'll be getting back from our server will be in the form of JSON which looks like this:

```
{"url": "1.jpeg", "size": 36, "title": "Beachside"}
```

This is an object with three fields: two strings and one int. Let's update our Photo type alias to reflect this:

```
type alias Photo =
    { url : String
    , size : Int
    , title : String
    }
```

We could decode this using the technique we just learned, with one slight difference. We'd have to use `map3` instead of `map2`, because we have three fields instead of two.

```
photoDecoder : Decoder Photo
photoDecoder =
    map3
        (\url size title -> { url = url, size = size, title = title })
        (field "url" string)
        (field "size" int)
        (field "title" string)
```

How far can this approach take us? If we added a fourth field, we'd change `map3` to `map4`. The `Json.Decode` module also includes `map5`, `map6`, `map7`, and `map8`, but `map8` is as high as it goes.

From there we can either combine the decoders we've already used...or we can introduce a library designed for larger-scale JSON decoding!

PIPELINE DECODING

The `Json.Decode.Pipeline` is designed to make life easier when decoding large objects. It comes from a popular third-party package called `NoRedInk/elm-decode-pipeline` - so let's install it real quick before we proceed:

```
elm-package install NoRedInk/elm-decode-pipeline
```

Let's also write a decoder for it. We'll do that by deleting the following import:

```
import Html.Attributes exposing (→)
```

...and then adding all of this to the end of our imports list:

```
import Html.Attributes exposing (id, class, classList, src, name, type_, title)
import Json.Decode exposing (string, int, list, Decoder)
import Json.Decode.Pipeline exposing (decode, required, optional)
```

Then the coast is clear to replace `photoDecoder` with this:

Listing 4.2 photoDecoder

```
photoDecoder : Decoder Photo
photoDecoder =
    decode buildPhoto
        |> required "url" string
        |> required "size" int
        |> optional "title" string "(untitled)"
```

```
buildPhoto : String -> Int -> String -> Photo
buildPhoto url size title =
    { url = url, size = size, title = title }
```

- ➊ Pass decoded values to the `buildPhoto` function
- ➋ "url" is required, and must be a string
- ➌ "size" is required, and must be an integer
- ➍ "title" is optional, and defaults to "(untitled)"

NOTE We changed our `import Html.Attributes` line because if it were still exposing `(..)`, then we'd be exposing two functions called `required`: both `Json.Decode.Pipeline.required` and `Html.Attributes.required`. Referencing a value that's exposed by two modules at once is an error!

Let's break down what's happening here.

1. `decode buildPhoto` begins the pipeline. It says that our decoder will decode the arguments to `buildPhoto`, one by one, and ultimately the whole decoder will succeed

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

unless any of the steps in this pipeline fails. Since `buildPhoto` accepts three arguments, we'll need three pipeline steps after this. otherwise, the compiler will give an error.

2. required "url" string says that we need what we're decoding to be a JSON object with the `String` field "url". We're also saying that if decoding succeeds, we should use this first result as the first argument to `buildPhoto`. Decoding could fail here, because either the "url" field was missing, or because the field was present...but was not a `String`.
3. required "size" int does the same thing required "url" string except that it decodes to an integer instead of a string.
4. optional "title" string "(untitled)" - this is similar to the required steps, but with one important difference: in this example, if the "title" field were either missing or null, instead of failing decoding, this decoder will default on that final argument—that is, the title string would default to "(untitled)".

Figure 4.8. shows how the `buildPhoto` arguments and `photoDecoder` arguments match up.

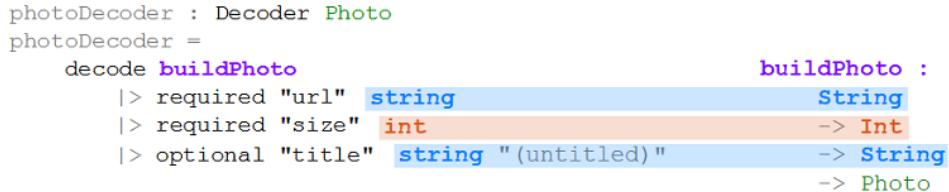


Figure 4.8 The relationship between `photoDecoder` and `buildPhoto`

By the way, notice anything familiar about what `buildPhoto` does?

All it does is take one argument for each of the fields in `Photo`, and then assign them uncritically without altering them in any way. We already have a function that does this! It's the `Photo` function, which we got it for free because we defined the `Photo` type alias.

Let's delete `buildPhoto` and replace `photoDecoder = decode buildPhoto` with this:

```

photoDecoder =
    decode Photo
  
```

Much easier!

WARNING Reordering any function's arguments can lead to unpleasant surprises. Since reordering the fields in the Model type alias has the consequence of reordering the Model function's arguments, you should be exactly as careful when reordering a type alias as you would be when reordering any function's arguments!

4.3.3 Decoding JSON HTTP Responses

We've already seen how we can use `Http.getString` to obtain a `String` from a server, and how we can use `decodeString` and a `Decoder` to translate a `String` into a list of `Photo` records for our `Model`.

Although we could use `Http.getString` and `decodeString` to populate our `Model` in this way, there is another function in the `Http` module which will take care of both for us.

`HTTP.GET`

The `Http.get` function requests data from a server and then decodes it. Here is how the types of `Http.getString` and `Http.get` match up:

```
getString : String -> Request String
get      : Decoder value -> String -> Request value
```

Comparing types like this suggests how these functions are similar and how they differ.

They both accept a URL string and return a `Request`. However, where `getString` takes no other arguments and produces a `String` on success, `get` additionally accepts a `Decoder` value, and on success produces a `value` instead of a `String`.

As you might expect, if we give `Http.get` a decoder of `(list int)` and the response it gets back is the JSON payload `"[1, 2, 3]"`, then `Http.get` will successfully decode that into `Ok (List Int)`. If decoding fails, we will instead get `Err UnexpectedPayload`.

What's `UnexpectedPayload`? It's a type constructor for that `Http.Error` type we said we'd get back to! (Sure enough, here we are getting back to it.) `Http.Error` is a union type which describes various ways a HTTP request can fail. It looks like this:

```
type Error
  = BadUrl String
  | Timeout
  | NetworkError
  | BadStatus (Response String)
  | BadPayload String (Response String)
```

Since it's a union type, we can run a *case-expression* on any `Http.Error` value—with different branches for `Timeout`, `NetworkError`, and so on—to do custom error handling based on what went wrong.

TIP The lower-level `Http.request` function lets you customize requests in more depth, including performing error handling based on raw status code.

Since we want to decode the JSON from our server into a `List Photo`, and we have a `Decoder` `Photo`, we can use `Json.Decode.list` to end up with the information we want. Let's change `initialCmd` to do just that while it references our shiny new URL:

```
initialCmd : Cmd Msg
initialCmd =
  list photoDecoder
```

```
|> Http.get "http://elm-in-action.com/photos/list.json"
|> Http.send LoadPhotos
```

This means we'll be sending a `List Photo` to `LoadPhotos` instead of a `String`, so we'll need to update its definition to match:

```
type Msg
  = SelectByUrl String
  | SelectByIndex Int
  | SurpriseMe
  | SetSize ThumbnailSize
  | LoadPhotos (Result Http.Error (List Photo))
```

This lets us simplify the `LoadPhotos (Ok ...)` branch of update's *case-expression* quite a bit! Having direct access to photos means we no longer need to build that value up using a *let-expression*. We can also use our new friends `Maybe.map` and `.url` to determine `selectedUrl` right on the same line!

```
LoadPhotos (Ok photos) ->
  { model
    | photos = photos
      , selectedUrl = Maybe.map .url (List.head photos)
    }
  , Cmd.none
)
```

RENDERING THE METADATA

Now all that remains is to have our view render the new caption and download size metadata!

We can do this by adding one line to `viewThumbnail`:

```
viewThumbnail selectedUrl thumbnail =
  img
    [ src (urlPrefix ++ thumbnail.url)
    , title (thumbnail.title ++ " [" ++ toString thumbnail.size ++ " KB]")
    , classList [ ( "selected", selectedUrl == Just thumbnail.url ) ]
    , onClick (SelectByUrl thumbnail.url)
    ]
  []
```

Great! At this point everything should compile, and we can open `index.html` to see the result:

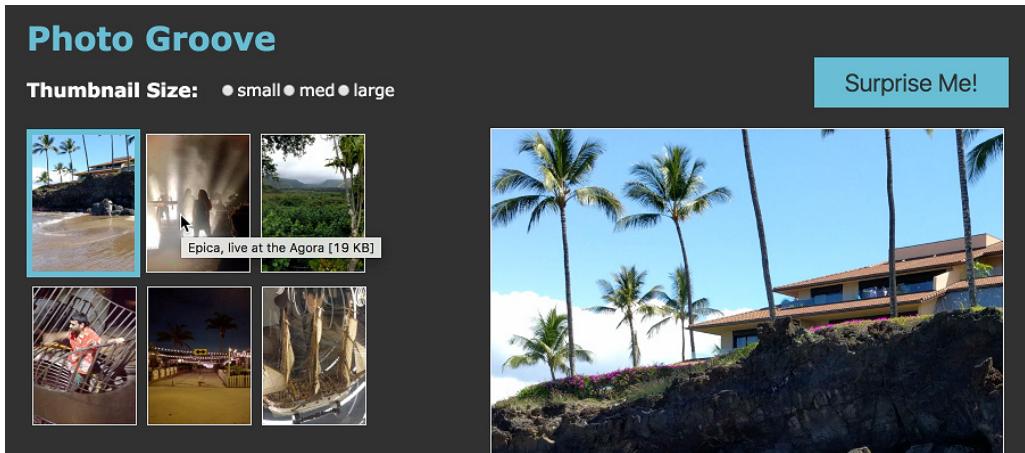


Figure 4.9 The final application

Now we're reading our list of photos from the server rather than hardcoding them, and we've given ourselves a nice foundation on which to build an even richer application!

4.4 Summary

We learned quite a few things in the course of making Photo Groove talk to servers!

- A `Decoder` can validate and translate JSON into an Elm value.
- The `Json.Decode` module provides primitive decoders like `float`, `int`, and `string`.
- The `Json.Decode.list` function turns a `Decoder Bool` into a `Decoder (List Bool)`.
- The `Json.Decode.Pipeline` module offers functions to decode objects in pipeline style.
- *Pattern matching* lets us trade nested *case-expressions* for longer branch conditions.
- A `Result` is either `Ok okValue` in case of success, or `Err errValue` in case of failure.
- `String.split` splits a string around a given separator, resulting in a list of strings.
- A `Request` describes a request we want to make. `Http.send` turns it into a `Cmd`.
- `Http.getString` requests a plain `String` from the server at the given URL.
- `Http.get` works like `Http.getString`, except it runs the given `Decoder` on the result.
- The `init` field passed to `Html.program` lets us specify a `Cmd` to run on startup.

We also saw how the pipeline operator (`|>`) lets us write expressions by starting with a value and then running a series of transformations on it. Table 4.6 shows an example of this.

Table 4.6 The same expression with and without pipelines

Pipeline Style	Without Pipelines
model.photos > Array.fromList > Array.get index	Array.get index (Array.fromList model.photos)

Now that we've gotten *Photo Groove* talking to a server, we'll get it talking to JavaScript. This will let us tap into the enormous ecosystem of JavaScript libraries out there. Let's see what that can get us!

Listing 3.7 The complete PhotoGroove.elm

```
module PhotoGroove exposing(..)

import Html exposing(..)
import Html.Events exposing (onClick)
import Array exposing (Array)
import Random
import Http
import Html.Attributes exposing (id, class, classList, src, name, type_, title)
import Json.Decode exposing (string, int, list, Decoder)
import Json.Decode.Pipeline exposing (decode, required, optional)

photoDecoder : Decoder Photo
photoDecoder =
    decode Photo
        |> required "url" string
        |> required "size" int
        |> optional "title" string "(untitled)"

urlPrefix : String
urlPrefix =
    "http://elm-in-action.com/"

type ThumbnailSize
    = Small
    | Medium
    | Large

view : Model -> Html Msg
view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , button
            [ onClick SurpriseMe ]
            [ text "Surprise Me!" ]
        , h3 [] [ text "Thumbnail Size:" ]
        , div [ id "choose-size" ]
            (List.map viewSizeChooser [ Small, Medium, Large ])
        , div [ id "thumbnails", class (sizeToString model.chosenSize) ]
            (List.map (viewThumbnail model.selectedUrl) model.photos)
```

```

        , viewLarge model.selectedUrl
    ]

viewLarge : Maybe String -> Html Msg
viewLarge maybeUrl =
    case maybeUrl of
        Nothing ->
            text ""
        Just url ->
            img [ class "large", src (urlPrefix ++ "large/" ++ url) ] []

viewThumbnail : Maybe String -> Photo -> Html Msg
viewThumbnail selectedUrl thumbnail =
    img
        [ src (urlPrefix ++ thumbnail.url)
        , title (thumbnail.title ++ " [" ++ toString thumbnail.size ++ " KB]")
        , classList [ ( "selected", selectedUrl == Just thumbnail.url ) ]
        , onClick (SelectByUrl thumbnail.url)
        ]
    []

viewSizeChooser : ThumbnailSize -> Html Msg
viewSizeChooser size =
    label []
        [ input [ type_ "radio", name "size", onClick (SetSize size) ] []
        , text (sizeToString size)
        ]

sizeToString : ThumbnailSize -> String
sizeToString size =
    case size of
        Small ->
            "small"
        Medium ->
            "med"
        Large ->
            "large"

type alias Photo =
    { url : String
    , size : Int
    , title : String
    }

type alias Model =
    { photos : List Photo
    , selectedUrl : Maybe String
    , loadingError : Maybe String
    , chosenSize : ThumbnailSize
    }

```

```

    }

initialModel : Model
initialModel =
{ photos = []
, selectedUrl = Nothing
, loadingError = Nothing
, chosenSize = Medium
}

photoArray : Array Photo
photoArray =
    Array.fromList initialModel.photos

getPhotoUrl : Int -> Maybe String
getPhotoUrl index =
    case Array.get index photoArray of
        Just photo ->
            Just photo.url

        Nothing ->
            Nothing

type Msg
= SelectByUrl String
| SelectByIndex Int
| SurpriseMe
| SetSize ThumbnailSize
| LoadPhotos (Result Http.Error (List Photo))

randomPhotoPicker : Random.Generator Int
randomPhotoPicker =
    Random.int 0 (Array.length photoArray - 1)

update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
    case msg of
        SelectByIndex index ->
            let
                newSelectedUrl : Maybe String
                newSelectedUrl =
                    model.photos
                        |> Array.fromList
                        |> Array.get index
                        |> Maybe.map .url
            in
                ( { model | selectedUrl = newSelectedUrl }, Cmd.none )

        SelectByUrl url ->
            ( { model | selectedUrl = Just url }, Cmd.none )

        SurpriseMe ->

```

```

let
    randomPhotoPicker =
        Random.int 0 (List.length model.photos - 1)
in
    ( model, Random.generate SelectByIndex randomPhotoPicker )

SetSize size ->
    ( { model | chosenSize = size }, Cmd.none )

LoadPhotos (Ok photos) ->
    ( { model
        | photos = photos
        , selectedUrl = Maybe.map .url (List.head photos)
    }
    , Cmd.none
)

LoadPhotos (Err _) ->
    ( { model
        | loadingError = Just "Error! (Try turning it off and on again?)"
    }
    , Cmd.none
)

initialCmd : Cmd Msg
initialCmd =
    list photoDecoder
        |> Http.get "http://elm-in-action.com/photos/list.json"
        |> Http.send LoadPhotos

viewOrError : Model -> Html Msg
viewOrError model =
    case model.loadingError of
        Nothing ->
            view model

        Just errorMessage ->
            div [ class "error-message" ]
                [ h1 [] [ text "Photo Groove" ]
                , p [] [ text errorMessage ]
                ]

main : Program Never Model Msg
main =
    Html.program
        { init = ( initialModel, initialCmd )
        , view = viewOrError
        , update = update
        , subscriptions = (\_ -> Sub.none)
        }

```

5

Talking to JavaScript

This chapter covers

- **Rendering Custom Elements**
- **Sending data to JavaScript**
- **Receiving data from JavaScript**
- **Initializing our Elm application using data from JavaScript**

Now we've gotten *Photo Groove* loading photos from a server. It's looking better and better! However, our manager has a concern. "We have the photos, and they certainly look sharp. But where is the *groove*? This is Photo Groove, not Photo Browse! You know what we need? *Filters*. You know, the kind that make normal photos look all wacky and messed-up? Those!"

Sounds groovy.

We collaborate with our team's visual designer to create a mockup of the filtering feature. The final design calls for a status bar to report on the filtering process, and sliders to control the filters' settings. And not just any sliders! Our designer wants each to have the specific look and feel of a Paper Slider, a Web Component built on Google's Material Design guidelines.

As luck would have it, we won't have to code either the filter effects or the sliders from scratch. They're both available as open-source libraries, but there's a catch: these libraries are written in JavaScript, not Elm.

Fortunately, Elm applications are not limited to using Elm libraries alone! In this chapter we'll expand *Photo Groove* to include a filtering feature. We'll add sliders which let users apply filters in varying degrees to the large photo. We'll create a status bar that displays information from the JavaScript libraries. Along the way we'll learn two different ways to incorporate JavaScript code into an Elm application: *custom elements* and *ports*.

Here we go!

5.1 Using Custom Elements

We'll start with the sliders. We want to create three sliders in total: **Hue**, **Ripple**, and **Noise**. Adjusting each value will transform how the large photo appears to the end user.

Figure 5.1 shows how the sliders fit into what we'll be building toward in this chapter.

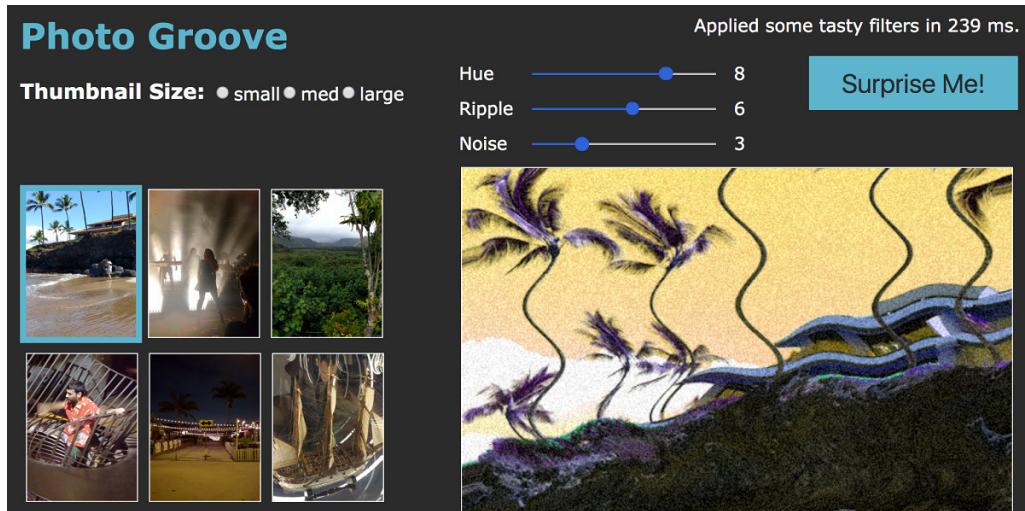


Figure 5.1 Adding sliders which control filtering for the large photo

Each slider controls an integer value which the filtering library will use for the magnitude of the filter effect. Sliding all three to 0 would disable filtering—leaving the photos unaltered, like how they look now—whereas setting all three to the maximum value of 11 would apply each filter at full power.

Implementing the sliders using Custom Elements will involve three steps:

1. Run some JavaScript code before our Elm code executes, to register custom behavior for all `<paper-slider>` elements on the page.
2. Write Elm code to render these `<paper-slider>` elements using plain `Html` values.
3. Write custom event handlers to update our model based on the sliders' states.

Along the way, we'll learn a new use for the JSON Decoders we saw in Chapter 4, and compare two alternative ways of integrating the sliders into our data model. Let's get to it!

5.1.1 Importing Custom Elements

The Paper Slider we'll be using is implemented as a Custom Element, which is a part of the Web Components specification. Here's how Custom Elements work:

- Someone writes some JavaScript which defines a new element type. In this case the

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

new element is a `<paper-slider>`.

- We run that JavaScript code when our page loads, to register this custom element with the browser.
- From now on, whenever any code—in JavaScript or in Elm—creates an element whose tag name is "paper-slider", that element will behave according to this custom logic.

WARNING Custom Elements are implemented in JavaScript, so they may throw runtime exceptions!

REGISTERING CUSTOM ELEMENTS

To register the custom `<paper-slider>` element on our page, we'll first need to import the code for it. Let's add a `<link rel="import">` to the `<head>` of our `index.html` file, right below the `<link rel="stylesheet">` for our stylesheet:

```
<link rel="stylesheet" href="http://elm-in-action.com/styles.css">
<link rel="import" href="http://elm-in-action.com/wc/slider/paper-slider.html">
```

This `<link>` tag loads the JavaScript code, CSS styles, and markup necessary to describe a custom `<paper-slider>` element, and then registers that `<paper-slider>` on the page. From now on, whenever an element with the tag name "paper-slider" appears on this page, it will behave according to the implementation we've imported here.

You can find more Custom Elements at <http://webcomponents.org/>, or create and serve your own from any URL you like.

Polyfilling Custom Elements for Older Browsers

Since custom elements and `<link rel="import">` tags are relatively new features, not all older browsers support them.

We can fix this for browsers as old as Internet Explorer 11 by loading some JavaScript code known as a *polyfill*, which patches support for custom elements into these browsers. To polyfill support for Custom Elements, add this `<script>` between your two `<link>` tags:

```
<link rel="stylesheet" href="http://elm-in-action.com/styles.css">
<script src="http://elm-in-action.com/polyfill/webcomponents.js"></script>
<link rel="import" href="http://elm-in-action.com/wc/slider/paper-slider.html">
```

If your users only run modern browsers which have already implemented Custom Elements, you don't need the polyfill and can omit this `<script>`. (That said, the polyfill has no effect on browsers which support Custom Elements, so it should be harmless aside from being more for users to download.)

Custom elements from the Polymer Project use something called Shady DOM to improve performance in browsers that do not natively support certain parts of the Web Components specification. However, if you use Polymer elements inside other Polymer elements, they may not work properly in their stock configuration due to how Shady DOM interacts with virtual DOM systems such as Elm's.

A workaround is to disable Shady DOM by adding this before your first `<link rel="import">` tag:

```
<script>window.Polymer = {dom: "shadow"};</script>
```

For more information on configuring Shady DOM, see <https://www.polymer-project.org/1.0/docs/devguide/settings>

In practice, polyfills are never as good as proper support for a browser feature. This polyfill often incurs significant performance penalties for some Custom Elements, and users have reported some polyfilled Custom Elements not working on Mobile Safari 10. Be sure to cross-browser test! See <http://webcomponents.org/polyfills> to learn more.

ADDING A <PAPER-SLIDER> TO THE PAGE

Back at the beginning of Chapter 2 we saw the `Html.node` function. It takes three arguments:

1. Tag name
2. List of attributes
3. List of children

Here are two ways to create the same button, one using `node` and the other using `button`:

```
node "button" [ class "large" ] [ text "Send" ]
button [ class "large" ] [ text "Send" ]
```

Functions that create elements like `button`, `label` and `input` have tiny implementations—they do nothing more than call `node`. For example, `label` and `input` can be implemented like so:

```
label attributes children =
  node "label" attributes children

input attributes children =
  node "input" attributes children
```

TIP Because of partial application, we could also have written these as `label = node "label"` and `input = node "input"`—either way works!

Now that the `<paper-slider>` custom element has been registered on the page with our `<link>` tag, the only Elm code necessary to use it is to call `node "paper-slider"`.

Let's add a function that does this, at the end of `PhotoGroove.elm`:

```
paperSlider =
  node "paper-slider"
```

This `paperSlider` function will work the same way as functions like `button`, `div`, and so on. It also has the same type as they do:

```
paperSlider : List (Attribute msg) -> List (Html msg) -> Html msg
```

Go ahead and add this type annotation right above our `paperSlider` implementation.

VIEWING THE SLIDERS

Next let's invoke `paperSlider` to render the sliders. Add this beneath the `view` function:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

Listing 5.1 viewFilter

```
viewFilter : String -> Int -> Html Msg
viewFilter name magnitude =
    div [ class "filter-slider" ]
        [ label [] [ text name ] ①
        , paperSlider [ max "11" ] [] ②
        , label [] [ text (toString magnitude) ] ③
    ]
```

- ① Display the filter's name
- ② <paper-slider> that goes from 0 to 11
- ③ Display the current magnitude

We'll also need to add `max` to the `Html.Attributes` we're exposing in our imports:

```
import Html.Attributes exposing (id, class, classList, src, name, max, type_, title)
```

If we try to compile this, we'll get a naming error: `max` is ambiguous! The compiler suggests:

Maybe you want one of the following?

```
Basics.max
Html.Attributes.max
```

The cause of this error is that Elm has a built-in math function called `max`. It's in the `Basics` module, and every Elm file has an implicit `import Basics exposing(..)` so that we can use common functions like `negate`, `not`, and `max` without having to qualify them as `Basics.negate`, `Basics.not`, or `Basics.max`.

RESOLVING AMBIGUOUS NAMES

The compiler is reporting an ambiguity in our code: when we wrote `max`, we might have meant either `Basics.max` or `Html.Attributes.max`, since both `max` functions are exposed in this file.

Let's clarify which we meant by fully qualifying `max`:

```
paperSlider [ Html.Attributes.max "11" ] []
```

This resolves the compiler error, but `Html.Attributes.max` is pretty verbose! We can shorten it by giving `Html.Attributes` an alias of `Attr` using the `as` keyword:

```
import Html.Attributes as Attr exposing (id, class, classList, src, name, type_, title)
```

Now anywhere we would write `Html.Attributes.foo` we can write `Attr.foo` instead.

Let's refactor our `input` to look like this:

```
paperSlider [ Attr.max "11" ] []
```

Splendid! This should make everything compile neatly.

RENDERING THREE SLIDERS

Now we can call `viewFilter` in our view function to get a few sliders rendering. Let's do that right below the code for the *Surprise Me!* button.

```
view model =
  div [ class "content" ]
    [ h1 [] [ text "Photo Groove" ]
    , button
      [ onClick SurpriseMe ]
      [ text "Surprise Me!" ]
    , div [ class "filters" ]
      [ viewFilter "Hue" 0
      , viewFilter "Ripple" 0
      , viewFilter "Noise" 0
      ]
    ]
```

Figure 5.2 shows how this looks on the page.

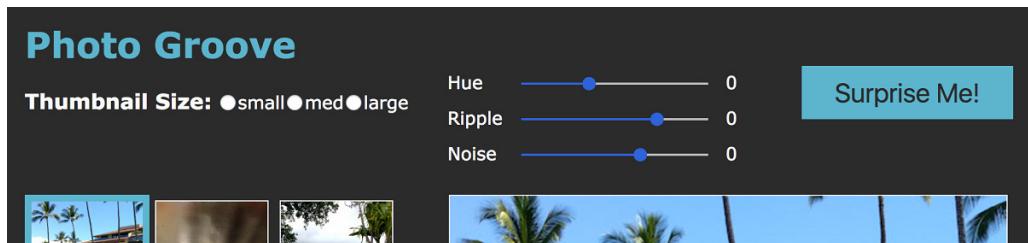


Figure 5.2 Adding paperSliders for Hue, Ripple, and Noise

Those sliders look slick, but they aren't completely operational yet. We want the numbers next to the sliders to change as the user slides their values around. Let's make that happen!

5.1.2 Handling Custom Events

If we take a quick glance at the Paper Slider's official documentation - hosted at <https://www.webcomponents.org/element/PolymerElements/paper-slider/paper-slider> - we can see that the slider fires a custom event, `immediate-value-changed`, as the user slides.

How might we respond to that event? `Html.Attributes` has a variety of built-in event handler functions such as `onClick`, but there's no `onImmediateValueChange`. The `immediate-value-changed` event is a custom event built into `paper-slider`, and the `Html.Attributes` module has no idea it exists. So how can we specify a handler for it?

HTML.ATTRIBUTES.ON

The `Html.Attributes.on` function lets us create a custom event handler, similarly to how the `Html.node` function lets us create a custom element. The `on` function has this type:

```
on : String -> Decoder msg -> Attribute msg
```

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

The `String` is the name of the event, which in this case is `"immediate-value-changed"`. The `Decoder` argument is a `Json.Decode.Decoder`, the same type of decoder we built in Chapter 4 for our HTTP responses. Here we won't be using the `Decoder` on a JSON string coming back from a server, but rather on a JavaScript event object!

DECODING IMMEDIATE-VALUE-CHANGED

The event object JavaScript uses for `"immediate-value-changed"` looks something like this:

```
{target: {immediateValue: 7}}
```

The `target` field refers to the slider object itself, which in turn has a field called `value`—an `Int` representing the slider's current value. We can use the `Json.Decode.field` and `Json.Decode.int` functions to write a decoder for this like so:

```
field "target" (field "immediateValue" int)
```

Table 5.1 compares this decoder to the example `email` decoder we wrote in Chapter 4.

Table 5.1. Decoding objects with field decoders

Decoder	JSON	Result
field "email" string	5	Err "Expected object, got 5"
	{"email": 5}	Err "Expected string for 'email', got 5"
	{"email": "cate@nolf.com"}	Ok "cate@nolf.com"
Decoder	JavaScript	Result
field "target" (field "immediateValue" int)	9 {"target": 9} {"target": {"immediateValue": 9} }	Err "Expected object, got 9" Err "Expected object for 'target', got 9" Ok 9

JSON.decode.at

There's a convenience function in `Json.Decode` for the case where we want to call `field` on another `field` like this: `Json.Decode.at`. It takes a list of field strings and traverses them in order. These two decoders do the same thing:

```
field "target" (field "immediateValue" int)  
at [ "target", "immediateValue" ] int
```

DECODING A MSG

This `Decoder Int` will decode an integer from a JavaScript object such as `{target: {value: 7}}`. But is that what we want? Let's look at the type of `on` again:

```
on : String -> Decoder msg -> Attribute msg
```

Notice that it wants a `Decoder msg` and then returns an `Attribute msg`. That tells us we want it to decode not an integer, but a message. Ah! We have a message type named `Msg`. So how do we convert between the `Decoder Int` that we have and the `Decoder Msg` that `on` expects?

USING JSON.DECODE.MAP

The `Json.Decode.map` function is just what the doctor ordered! It converts between decoded values, as shown in Table 5.2.

Table 5.2 `Json.Decode.map`

Expression	Description
<code>Json.Decode.map negate float</code>	Decode a float, then negate it.
<code>Json.Decode.map (\num -> num * 2) int</code>	Decode an integer, then double it.
<code>Json.Decode.map (_ -> "[[redacted]]") string</code>	Decode a string, then replace it with "[[redacted]]" no matter what it was originally. Note that this will still fail if it attempts to decode a non-string value!

Since `Json.Decode.map` takes a function which converts one decoded value to another, we can use it to convert our decoded `Int` into a `Msg`.

UPDATING IMPORTS

We'll need to expose the `on` and `at` functions in our imports in order to reference them:

```
import Html.Events exposing (onClick, on)
import Json.Decode exposing (string, int, list, Decoder, at)
```

With those changes in place, we're ready to add this `onImmediateValueChange` function to the end of `PhotoGroove.elm`:

Listing 5.2 `onImmediateValueChange`

```
onImmediateValueChange : (Int -> msg) -> Attribute msg
onImmediateValueChange toMsg =
    let
        targetImmediateValue : Decoder Int
        targetImmediateValue =
            at [ "target", "immediateValue" ] int
        msgDecoder : Decoder msg
    in
        1
        1
        1
    2
```

```

msgDecoder =
    Json.Decode.map toMsg targetImmediateValue ②
in
    on "immediate-value-changed" msgDecoder ③

```

- ① Decode the integer located at event.target.immediateValue
- ② Convert that integer to a message using toMsg
- ③ Create a custom event handler using that decoder

Notice how `onImmediateValueChange` takes a `toMsg` function? This is because we need it to be flexible. We plan to have multiple sliders on the page, and we'll want each of them to have a unique `Msg` constructor so that we can tell their messages apart. The `toMsg` argument lets us pass in the appropriate constructor on a case-by-case basis, which will come in handy later.

REFACTORING TO USE PIPELINES

Notice how we assemble this value in three steps (`targetImmediateValue`, `msgDecoder`, and `on`), and each step's final argument is the previous step's return value? That means we can rewrite this to use the pipeline style we learned in Chapter 4! Let's do that refactor:

```

onImmediateValueChange : (Int -> msg) -> Attribute msg
onImmediateValueChange toMsg =
    at [ "target", "immediateValue" ] int
        |> Json.Decode.map toMsg
        |> on "immediate-value-changed"

```

Try walking through each step in the pipeline and finding the equivalent code in Listing 5.2. All the same logic is still there, just reorganized!

ADDING EVENT HANDLING TO VIEWFILTER

Now that we have `onImmediateValueChange`, we can use it in our `viewFilter` function.

Remember how we made `onImmediateValueChange` accept an `Int -> msg` function, so that we could pass it a different `Msg` constructor on a case-by-case basis? We'll want `viewFilter` to have that same flexibility. The only difference will be that since `viewFilter` returns an `Html Msg`, we won't be using a type variable like `msg`; instead, `viewFilter` will accept an `Int -> Msg` function.

Listing 5.3 using onImmediateValueChange in viewFilter

```

viewFilter : String -> (Int -> Msg) -> Int -> Html Msg
viewFilter name toMsg magnitude =
    div [ class "filter-slider" ]
        [ label [] [ text name ]
            , paperSlider [ Attr.max "11", onImmediateValueChange toMsg ] []
            , label [] [ text (toString magnitude) ]
        ]

```

- ① Calling `onImmediateValueChange` just like we did `onClick`

With this revised `viewFilter` implementation at the ready, our display logic is ready to be connected to our model and `update`. Once we've revised those to work with `viewFilter`, our shiny new Custom Elements will be fully integrated into our application!

5.1.3 Responding to Slider Changes

We have three filters, each of which has a name and a magnitude. How should we track their current values in our model? Let's walk through two approaches, comparing the pros and cons of each, and then at the end decide which to use.

The first approach would be to add three fields to the model: one for Hue, one for Ripple, and one for Noise. The alternative would prioritize flexibility, and store the filters as a list of `{ name : String, amount : Int }` records. Table 5.3 shows these approaches side-by-side.

Table 5.3 Storing filter data as three ints versus one list of records

Three Ints	One List of Records
<pre>type alias Model = { photos : List Photo ... , hue : Int , ripple : Int , noise : Int }</pre>	<pre>type alias Model = { photos : List Photo ... , filters : List { name : String, amount : Int } }</pre>
<pre>initialModel : Model initialModel = { photos = [] ... , hue = 0 , ripple = 0 , noise = 0 }</pre>	<pre>initialModel : Model initialModel = { photos = [] ... , filters = [{ name = "Hue", amount = 0 } , { name = "Ripple", amount = 0 } , { name = "Noise", amount = 0 }] }</pre>

Each model design has its own strengths and weaknesses, which become clearer when we write code that references these parts of the model.

Let's consider how our `update` implementations might compare. In either approach, we'd expand our `Msg` type's constructors and then add at least one branch to our `update` function's *case-expression*. Table 5.4 compares how these revisions would look for each approach.

Table 5.4 Updating the model

Three Ints	One List of Records
<pre>type Msg = SelectByUrl String SetHue Int SetRipple Int SetNoise Int</pre> <pre>case msg of SetHue hue -> ({ model hue = hue } , Cmd.none) SetRipple ripple -> ({ model ripple = ripple } , Cmd.none) SetNoise noise -> ({ model noise = noise } , Cmd.none)</pre>	<pre>type Msg = SelectByUrl String SetFilter String Int</pre> <pre>case msg of SetFilter name amount -> let transform filter = if filter.name == name then { name = name , amount = amount } else filter filters = model.filters > List.map transform in ({ model filters = filters } , Cmd.none)</pre>

The “list of records” approach is more flexible. If we decided to add a fourth filter, such as “blur”, we could add it to `initialModel` and boom! It would appear on the page instantly. With the three integers approach, adding a fourth integer would require expanding not only `initialModel`, but also adding a field to our type alias for `Model`, a constructor for our `Msg`, and a branch for update’s `case-expression`. That sure sounds like more work!

Finally let’s compare how rendering would look.

Table 5.5 Viewing filter data from three fields versus one

Three Ints	List of Records (assuming <code>viewFilter</code> tweaked)
<pre>, div [class "filters"] [viewFilter SetHue "Hue" model.hue , viewFilter SetRipple "Ripple" model.ripple , viewFilter SetNoise "Noise" model.noise]</pre>	<pre>, div [class "filters"] (List.map viewFilter model.filters)</pre>

Assuming we tweaked `viewFilter` to accept a `{ name : String, amount : Int }` record, we could save ourselves a few lines of code with the “list of records” approach, using `List.map` as shown in Table 5.5. Once again “list of records” comes out ahead!

CHOOSING AN APPROACH

So far we’ve compared these approaches by conciseness, and also by how much effort it would take to add more filters later. However, our analysis is overlooking a crucial consideration, one which tends to make a much bigger difference in Elm than in JavaScript:

Which approach rules out more bugs?

As we’ve seen, Elm’s compiler gives us certain guarantees that can rule out entire categories of bugs. Some code can better leverage these guarantees than others! We can “help the compiler help us” by taking it into consideration when making decisions like these.

What potential future headaches would each approach let the compiler rule out?

CONSIDERING POTENTIAL BUGS

Suppose we go with the “list of records” approach and one of our coworkers makes an innocuous typo—writing “`Rippl`” instead of “`Ripple`”—such that a `SetFilter "Rippl" 5` message comes through our `update` function. That code will compile, but it won’t work properly because our `filter.name == "Rippl"` condition will never pass. We’ll have to hope our coworker catches this bug in testing!

What if our coworker makes the same typo in the three-ints approach? Attempting to create a `SetRippl 5` message will be a compiler error, because we named our `Msg` constructor `SetRipple`, not `SetRippl`.

Table 5.6 Making a Mistake

	Three Ints	List of Records
Message	<code>SetRippl 5</code>	<code>SetFilter "Rippl" 5</code>
Outcome	Elm’s compiler tells us we have a naming error.	This compiles, but now we have a bug!

CONSIDERING FUTURE CHANGES

What about making changes? Suppose in the future we need to rename “`Ripple`” to “`Swirl`”.

In the “three ints” approach, we can rename our `Msg` constructor from `SetRipple` to `SetSwirl`, and our `Model` field from `ripple` to `swirl`—and then watch Elm’s compiler tell us exactly what other parts of our code base need changing as a result. If we miss anything, we’ll get a helpful type mismatch error telling us every spot we overlooked.

With the “list of records” approach, we have to hope none of our coworkers ever wrote code using the hardcoded magic string “`Ripple`” instead of a constant. If they did, and

something breaks, the only way we'll find out before the bug reaches production is if someone also wrote a test that happens to fail when we change "Ripple" to "Swirl".

CHOOSING RELIABILITY

By using individual fields instead of a list of records, we can rule out the entire category of bugs related to invalid filter names.

Increasing conciseness and saving potential future effort are nice, but preventing bugs in a growing code base tends to be more valuable over time. Verbosity has a predictable impact on a project, whereas the impact of bugs can range from "quick fix" to "apocalyptic progress torpedo." Let's rule them out when we can!

We'll go with the approach that prevents more bugs. Take a moment to look back at Tables 5.3, 5.4, and 5.5, and implement the changes in the first column. Then let's revise our calls to `viewFilter` to use our model's new fields:

```
, div [ class "filters" ]
  [ viewFilter "Hue" SetHue model.hue
  , viewFilter "Ripple" SetRipple model.ripple
  , viewFilter "Noise" SetNoise model.noise
  ]
```

When the dust settles, you should be able to recompile with `elm-make --output=elm.js PhotoGroove.elm` and see the number labels change when the user slides.

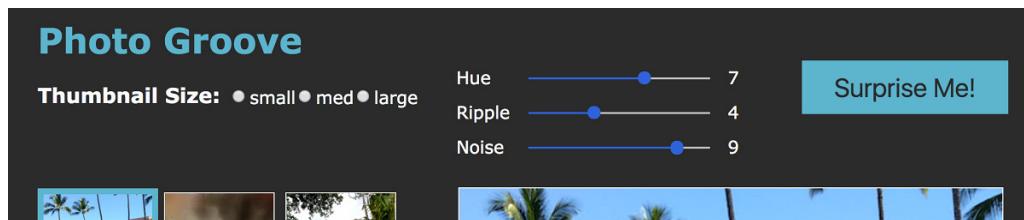


Figure 5.3 Sliding now changes the numbers next to the sliders

Now that we have our sliders set up, we can move on to introducing the filters themselves!

5.2 Sending Data to JavaScript

Now our model has all the data it needs to calibrate the filters. We'll use some JavaScript once again to apply the filters, although this time not Custom Elements. Instead we'll write code that passes configuration options and a `<canvas>` element to a JavaScript function, which will then proceed to draw groovy pictures on it.

First we'll get a basic proof-of-concept working, to confirm that we're successfully communicating across languages, and then we'll smooth out the implementation details until we're satisfied with how things are working.

5.2.1 Creating a Command using a Port

You may recall from Chapter 4 that an *effect* is an operation that modifies external state. You may also recall that if a function modifies external state when it runs, that function has a *side effect*. Elm functions are not permitted to have side effects, but JavaScript functions are.

TALKING TO JAVASCRIPT IS LIKE TALKING TO SERVERS

Since calling any JavaScript function may result in a side effect, Elm functions cannot call JavaScript functions anytime they please; this would destroy the guarantee that Elm functions have no side effects!

Instead, Elm talks to JavaScript the same way it talks to servers: by sending data out through a command, and receiving data in through a message. Figure 5.4 illustrates this.

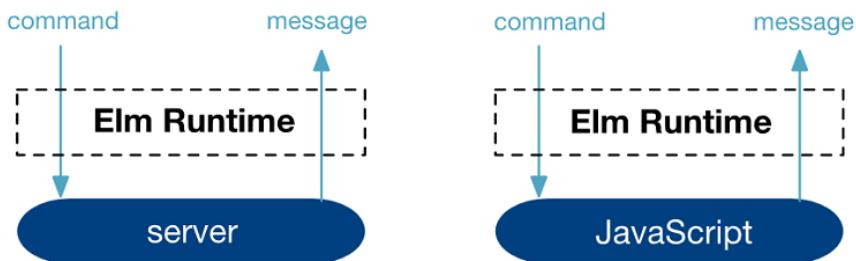


Figure 5.4 Using commands and messages to talk to both servers and JavaScript

This means that talking to JavaScript will have some characteristics in common with what we saw when talking to servers in Chapter 4:

- Data can only be sent using a command
- Data can only be received by `update`, and that data must be wrapped in a message
- We can translate and validate this incoming data using Decoders

NOTE In JavaScript, some effects are performed synchronously, with program execution halting until the effect completes. In contrast, an Elm `Cmd` always represents an *asynchronous* effect. This means that when we send data to JavaScript, it's always possible that other code might run before data gets sent back to Elm!

CREATING A COMMAND

Let's create a `Cmd` to send some data to JavaScript. First we'll define a type alias for the data we're going to send, right above our type alias for `Photo`:

```
type alias FilterOptions =
    { url : String
    , filters : List { name : String, amount : Int }
    }
```

This represents all the information our JavaScript library will need in order to apply the filters: the URL of the photo in question, plus the list of filters and their amounts.

USING A PORT TO DEFINE A FUNCTION

In Chapter 4 we created a `Cmd` using the `Http.send` function, but here we'll instead use a language feature designed specifically for talking to JavaScript: the `port` keyword.

Add this right above the `type alias for FilterOptions`:

```
port setFilters : FilterOptions -> Cmd msg
```

This declares a function called `setFilters`. Its type is `(FilterOptions -> Cmd msg)`. We don't write an implementation for this function, because the `port` keyword automatically writes one for us! `port` only needs to look at the type we requested to decide what the function should do.

All `port` functions that send data to JavaScript are defined using a very specific pattern:

- The `port` keyword must be followed by a function name and a type annotation
- The type annotation must be for a function that takes one argument
- The function must return `Cmd msg`, and nothing else—not even `Cmd Msg`!

TYPE VARIABLES IN VALUES

Notice that our new `setFilters` function returns `Cmd msg`, not `Cmd Msg`. We know from Chapter 3 that the lowercase "m" means `msg` is a *type variable*, like the `val` type variable we saw in `Array.fromList`:

```
fromList : List val -> Array val
```

This `List String -> Cmd msg` function is the first time we've seen a function annotation with a type variable in its return type, but no type variable in its parameters!

What does that mean? And how is `Cmd msg` different from `Cmd Msg`?

PORT COMMANDS NEVER SEND MESSAGES

Believe it or not, we've been using a `Cmd msg` for quite some time now! The official documentation for `Cmd.none` shows that it has this type:

```
none : Cmd msg
```

We'll dive deeper into values like this in future chapters, but for now it's enough to know that a `Cmd msg` is a **command which produces no message after it completes**.

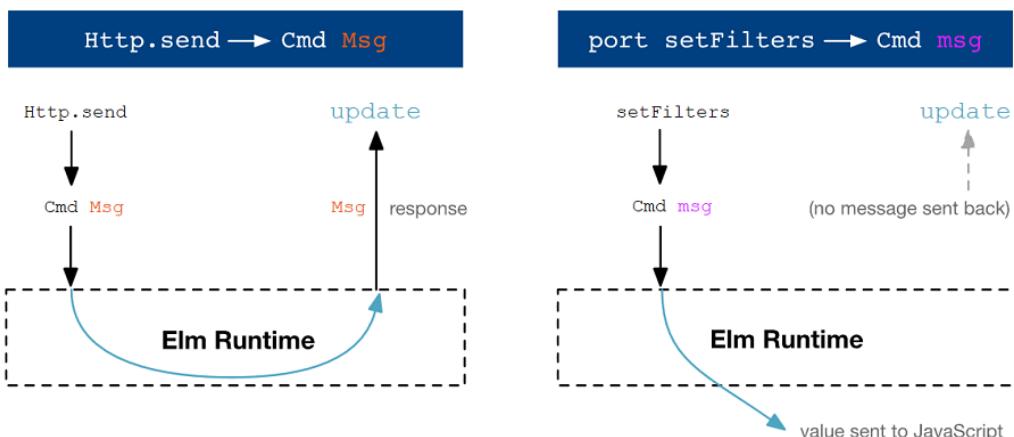
Both `Cmd.none` and `setFilters` produce no message after completing. The difference is that `Cmd.none` has no effect, whereas `setFilters` will perform the effect of sending data to JavaScript. (Specifically, it will send the `FilterOptions` value we pass it.) You can think of `setFilters` as a "fire and forget" command.

Table 5.7 Comparing Cmd Msg and Cmd msg.

Expression	Type	Effect	Message sent back to update
list photoDecoder > Http.get "http://..." > Http.send LoadPhotos	Cmd Msg	Send HTTP request	LoadPhotos
Cmd.none	Cmd msg	(none)	(no message sent back)
setFilters filterOptions	Cmd msg	Send filterOptions to JavaScript	(no message sent back)

NOTE While HTTP requests can fail, sending data to JavaScript cannot. We don't miss out on any error handling opportunities just because `setFilters` sends no message back to update.

Figure 5.5 shows the similarities and differences between a `Cmd` originating from an `Http.send` and a `Cmd` originating from a `port`.

**Figure 5.5 Comparing a Cmd from Http.send to a Cmd from a port**

PORT MODULES

Any module that uses the `port` keyword must be declared using the `port module` keyword, which means we'll need to change the first line of `PhotoGroove.elm` to this:

```
port module PhotoGroove exposing(..)
```

TIP If you're ever wondering whether a given module talks directly to arbitrary JavaScript code, checking to see if it's a `port module` is usually the quickest way.

CALLING SETFILTERS WHEN USERS SELECT PHOTOS

Now that we have our port set up, we need to call `setFilters` in order to send some data to JavaScript. We'll want to apply the filters every time the user selects a photo, so that's when we'll want `update` to return the command we get from `setFilter`.

Here's one way we could modify the `SelectByUrl` branch of our `update` function to do this.

Listing 5.4 SelectByUrl

```
SelectByUrl selectedUrl ->
  let
    filters =
      [ { name = "Hue", amount = model.hue }
      , { name = "Ripple", amount = model.ripple }
      , { name = "Noise", amount = model.noise }
      ]
    url =
      urlPrefix ++ "large/" ++ selectedUrl
    cmd =
      setFilters { url = url, filters = filters }
  in
    ( model, cmd )
```

However, the `SelectByUrl` branch is not the only way a user can select a photo! Users can also click the *Surprise Me!* button to select a photo at random, and we'll want to apply filters in that situation as well. This means we'll want to reuse the above code in two places: both the `SelectByUrl` branch of our *case-expression* as well as the `SurpriseMe` branch.

SHARING CODE BETWEEN UPDATE BRANCHES

Usually the simplest way to share code is to extract common logic into a helper function and call it from both places. This is just as true for `update` as it is for any function, so let's do that!

NOTE The structure of `update` permits clever alternatives to this venerable code sharing technique—like having `update` call itself, passing a different `Msg`. This saves us from writing another function, but it's more error-prone and less performant. If we do this, and later a teammate innocently changes how that other `Msg` responds to user input, our code breaks! Having `update` call itself also runs a bit slower because we create an unnecessary `Msg` and run an unnecessary *case-expression* on it. A helper function not only runs faster, it's less error-prone to maintain because it explicitly signals to future maintainers that code is being reused.

We'll name the helper function `applyFilters`, and add it right below `update`:

Listing 5.5 applyFilters

```
applyFilters: Model -> ( Model, Cmd Msg )
applyFilters model =
  case model.selectedUrl of
    Just selectedUrl ->
```

```

let
    filters =
        [ { name = "Hue", amount = model.hue }
        , { name = "Ripple", amount = model.ripple }
        , { name = "Noise", amount = model.noise }
        ]

    url =
        urlPrefix ++ "large/" ++ selectedUrl
in
    ( model, setFilters { url = url, filters = filters } )

Nothing ->
    ( model, Cmd.none )

```

Now we can have both `SelectByIndex` and `SelectByUrl` call `applyFilters` directly:

```

SelectByIndex index ->
    let
        newSelectedUrl : Maybe String
        newSelectedUrl =
            ...
    in
        applyFilters { model | selectedUrl = newSelectedUrl }

SelectByUrl selectedUrl ->
    applyFilters { model | selectedUrl = Just selectedUrl }

```

Lovely! Now whenever a user clicks either a photo or the *Surprise Me!* button, `setFilters` will return a `Cmd` that sends the appropriate `FilterOptions` value over to JavaScript.

Next we'll wire up the logic on the JavaScript side, which will receive that `FilterOptions` value and use it to apply some filters!

5.2.2 Receiving Data from Elm

Now we're going to write some JavaScript code. Whenever we access a JS library from Elm, it's best to write as little JS as possible. This is because if something crashes at runtime, it's a safe bet that the culprit is somewhere in our JS code—so the less of it we have, the less code we'll have to sift through to isolate the problem.

ADDING PASTA

The JavaScript code we need comes from an open-source image filtering library intuitively named `Pasta.js`. We can import Pasta by adding this `<script>` tag to `index.html`, right before the `<script>` that imports our compiled `elm.js` file:

```

<div id="elm-area"></div>

<script src="http://elm-in-action.com/pasta.js"></script>
<script src="elm.js"></script>

```

This `<script>` adds a global JavaScript function called `Pasta.apply` to the page. It takes two arguments:

1. A `<canvas>` element, which is where `apply` will draw the filtered photos.
2. An `options` object, which—in a remarkable coincidence—has the same structure as the `FilterOptions` record Elm will be sending to JavaScript via our `setFilter` port.

Let's introduce that `<canvas>`. We can do this with a quick change to `PhotoGroove.elm`: having `viewLarge` render a `canvas` instead of an `img`.

Let's replace the `Just url` branch of the `case-expression` in `viewLarge` with this:

```
Just url ->


```

`Pasta.apply` will take care of drawing the filtered photo onto this `canvas`. (The photo URL to draw will be sent to `Pasta.apply` via the `FilterOptions` we're sending through our port.)

RECEIVING DATA FROM THE SETFILTERS PORT

We can set up a callback function that receives data from our `setFilters` function like so:

```
<script src="http://elm-in-action.com/pasta.js"></script>
<script src="elm.js"></script>
<script>
  var app = Elm.PhotoGroove.embed(document.getElementById("elm-area"));

  app.ports.setFilters.subscribe(function(options) {
    Pasta.apply(document.getElementById("main-canvas"), options);
  });
</script>
```

We've never needed the return value of `Elm.PhotoGroove.embed` before, but now we do!

This object, typically called `app`, lets us subscribe to data Elm sends to JavaScript via ports like `setFilters`. When the Elm runtime executes the `Cmd` returned by `setFilters`, the callback function we've passed to `app.ports.setFilters.subscribe` will run. The `options` argument it accepts is the `FilterOptions` record Elm sent over, but converted from an Elm record to a JavaScript object.

WARNING Like Custom Elements, ports invoke JavaScript code—which may throw runtime exceptions. If we make a mistake in any of our JavaScript code here, Elm's compiler has no way to help us catch it.

Table 5.8 shows how a port translates immutable Elm values (like the `FilterOptions` record) into brand-new (potentially mutable) JavaScript values.

Table 5.8 Translating Elm values into JavaScript values via ports

Elm Value	Elm Type	JavaScript Value	JavaScript Type
"foo"	String	"foo"	string
4.2	Float	4.2	number
True	Bool	true	boolean
("foo", True, 4.2)	(String, Bool, Float)	["foo", true, 4.2]	Array
["drive", "line"]	List String	["drive", "line"]	Array
{ name = "Shadow" }	{ password : String }	{"name": "Shadow"}	object
Nothing	Maybe val	null	object

TIP Since all Elm values must be immutable, mutable values (like JavaScript objects and arrays) can't be sent through the port. Instead they automatically get copied into immutable data structures. This process has some overhead. Passing in values like strings and numbers, which are already immutable, has no such overhead.

Since we gave our `<canvas>` an `id` of "main-canvas", we can easily pass it to `Pasta.apply` by calling `document.getElementById("main-canvas")` to find it on the page by `id`.

TRYING IT OUT

Let's try it out! If we open the page, we see...well, an empty rectangle.

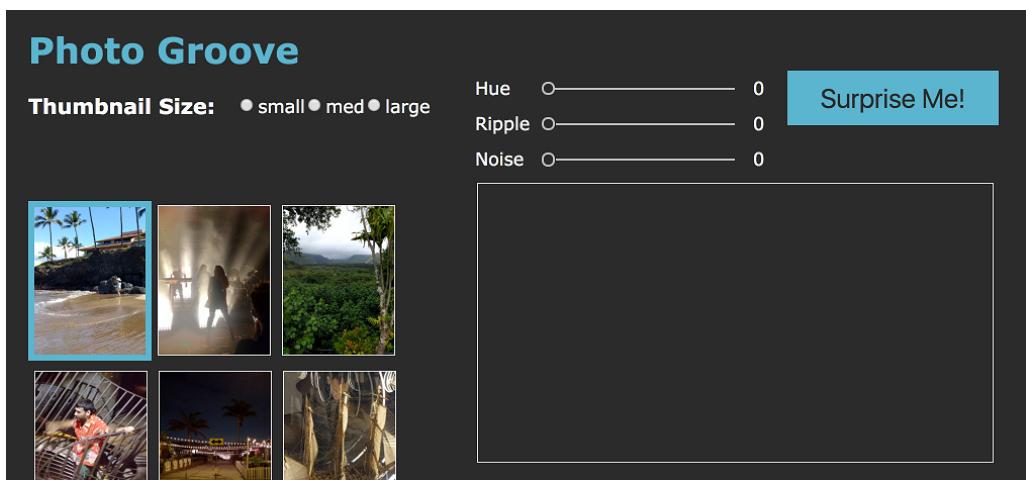


Figure 5.6 An empty rectangle. Yep.

That's not what we wanted! We'll figure that bug out later, but first let's see what these sliders do! Crank up that Ripple value, and then...hm, still nothing happens. What's the deal?

One last idea: maybe try selecting a new photo?

Oh hey, look at that!

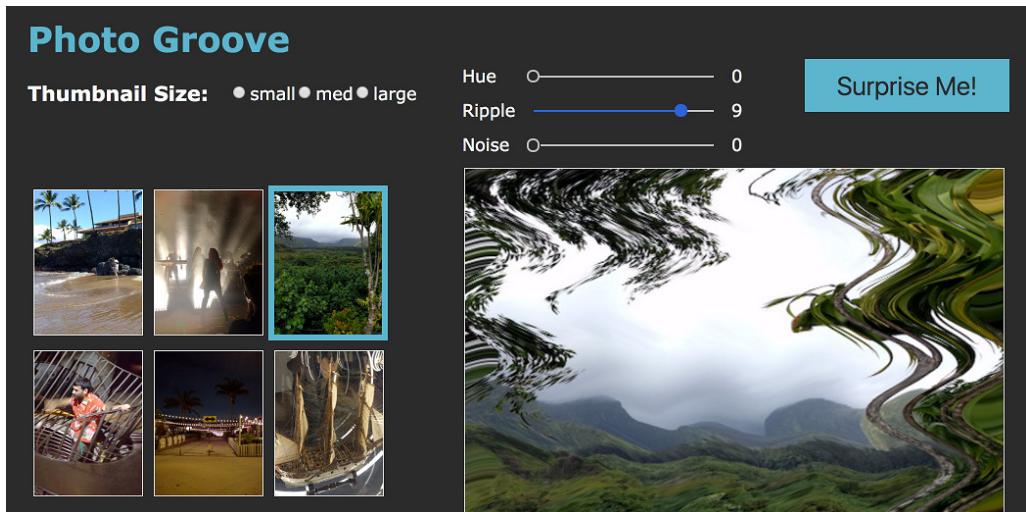


Figure 5.7 Increasing Ripple, then selecting another photo.

Sure enough, things changed...quite a lot! We are definitely talking to that JavaScript library.

However, there are some bugs to iron out. Playing around with it some more, we observe the following problems.

1. All of the Ripple and Noise values between 1 and 11 have the same effect.
2. Changing Hue doesn't seem to do anything.
3. When we initially load the page, we see an empty rectangle instead of the first photo.

Clearly we have some fixing to do!

CONVERTING THE AMOUNT VALUE TO A PERCENTAGE

The reason Hue isn't working is the same reason that the Ripple and Noise values between 1 and 11 do the same thing: we're not quite using the Pasta.js API correctly.

We're sending an `Int` between 1 and 11, but actually Pasta.js is expecting a percentage—a `Float` between 0 and 1. Since JavaScript does not draw a distinction between `Int` and `Float`, this mistake does not result in a type mismatch error—not even at runtime!—and instead, our code simply does not work as expected.

We can fix this by dividing our `Model`'s `hue`, `ripple`, and `noise` fields by 11 before sending them to JavaScript. Let's revise our logic in `applyFilters`:

©Manning Publications Co. We welcome reader comments about anything in the manuscript - other than typos and other simple mistakes. These will be cleaned up during production of the book by copyeditors and proofreaders.
<https://forums.manning.com/forums/elm-in-action>

```
filters =
  [ { name = "Hue", amount = toFloat model.hue / 11 }
  , { name = "Ripple", amount = toFloat model.ripple / 11 }
  , { name = "Noise", amount = toFloat model.noise / 11 }
]
```

We need that `toFloat` because Elm's division operator (`/`) only works if you give it two `Float` values. The `Basics.toFloat` function converts an `Int` to a `Float`, so `toFloat model.hue` converts `model.hue` from an `Int` to a `Float`—at which point we can divide it by 11 as normal.

NOTE This is another example of Elm's design emphasizing *being explicit*. If JavaScript required similar explicitness, we'd have caught this bug earlier on, when we tried to pass an `Int` to an API expecting a `Float`.

We'll also need to update `FilterOptions` to expect a `Float` for the `amount` field.

```
type alias FilterOptions =
  { url : String
  , filters : List { name : String, amount : Float }
}
```

Now if we recompile and try our sliders again, we see a much more interesting range of Hue, Ripple, and Noise values. Lovely! There are still some bugs left, but we're making progress.

CHANGING AS WE SLIDE

This will be way more fun if the photos update every time we slide, right? We know that `applyFilters` calls our `setFilters` port, so to apply the filters every time we slide, all we need to do is to run our `applyFilters` command every time we slide:

```
update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
  case msg of
    SetHue hue ->
      applyFilters { model | hue = hue }

    SetRipple ripple ->
      applyFilters { model | ripple = ripple }

    SetNoise noise ->
      applyFilters { model | noise = noise }
```

Recompile, and...presto! We're still greeted with an empty rectangle—we'll fix that next—but now whenever we slide, the photos update in realtime to reflect the new filter settings. Whee!

Let's fix that pesky bug where we see an empty rectangle instead of a photo on page load.

5.2.3 Timing DOM Updates

Showing the initial photo when the page loads takes two steps—a straightforward step and a tricky step. Let's start with the straightforward one: what we do after loading the photos.

APPLYING FILTERS AFTER LOADING PHOTOS

The problem here is that we currently load the photos and then re-render the view, but re-rendering the view is no longer enough to make a photo show up! We now need to call `applyFilters` so that `Pasta` can render something to the canvas.

To fix this, we need the `LoadPhotos` branch of our `update` function's *case-expression* to call `applyFilters` after updating `model`:

```
LoadPhotos (Ok photos) ->
  applyFilters
    { model
      | photos = photos
      , selectedUrl = Maybe.map .url (List.head photos)
    }
```

However, if you recompile and reload, you'll see the same behavior. Why wasn't that enough?

TIMING PORTS AND RENDERING

The remaining problem is a matter of timing. Let's break down what is happening:

1. We initialize `model.photos` to `[]` and `model.selectedUrl` to `Nothing`.
2. We request a list of photos from the server.
3. `view` runs, passing `Nothing` to `viewLarge` because `model.selectedUrl` is `Nothing`.
4. Because `viewLarge` received `Nothing`, it declines to render the `<canvas>`!
5. The server responds with our photos, meaning `update` gets a `LoadPhotos` message.
6. `update` runs its `LoadPhotos` branch and returns a new `model` (with a `selectedUrl` that is no longer `Nothing`), as well as a `Cmd` that will instruct our JavaScript code to have `Pasta` render to the `<canvas>`.

See the problem? Step 6 tells `Pasta` to render to a `<canvas>`, but as we noted in Step 4, no such `<canvas>` has been created yet!

This means our JavaScript call to `Pasta.apply(document.getElementById("main-canvas"), options)` will silently fail. Shortly after this happens, `view` will run again with the new `model`. This time, `model.selectedUrl` will not be `Nothing`—meaning `viewLarge` will happily ask for a fresh, blank `<canvas>`. Great.

We were so close, though! If Elm would have rendered the `view` before running the `Cmd` that invoked `Pasta`, the `<canvas>` would have been on the DOM before the JavaScript executed, and everything would have been fine! So why doesn't Elm do this?

OPTIMIZED DOM UPDATES

One reason the Elm Runtime has good performance is that it skips unnecessary renders.

See, browsers only repaint the DOM as pixels on users' screens every so often. If the Elm Runtime changes part of the DOM, and then changes it again before the next repaint, the first change will have been wasted time; only the second change will be painted for users to see.

You might think that if your `update` function gets called a million times in a single second, that your `view` would also be called a million times. Not so! Although those million updates will result in a million potential `Model` changes, Elm waits until the browser's next repaint to call `view` even once—with whatever `Model` is at that moment. Invoking `view` more frequently than that would result in DOM updates that the browser wouldn't bother to paint anyway.

SYNCHRONIZING WITH THE ELM RUNTIME

The JavaScript function `requestAnimationFrame` allows code to run just before the browser's next repaint. Since this is when the Elm runtime will schedule its next DOM update, we can use `requestAnimationFrame` to delay our call to `document.getElementById("main-canvas")` until after our next `view` has added the `<canvas>` we need to the DOM!

Since `requestAnimationFrame` accepts a single callback function, this means we can finally fix this bug by wrapping our call to `Pasta.apply` in a `requestAnimationFrame` callback function, like so:

```
app.ports.applyFilters.subscribe(function(options) {
  requestAnimationFrame(function() {
    Pasta.apply(document.getElementById("main-canvas"), options);
  });
});
```

TIP Anytime you need to trigger some JavaScript port code to run **after the next time** `view` results in a DOM update, you can synchronize things by wrapping your port code in `requestAnimationFrame` like this.

Let's recompile and bring up the page. Now the initial photo loads right after the page does!

5.3 Receiving Data from JavaScript

Now we're ready for the final piece of the puzzle: the status bar, which displays information reported by Pasta. First we'll start subscribing to Pasta's realtime status updates, which Pasta broadcasts as they occur. Then we'll set an initial status message which displays the version of Pasta we're using. Each change involves a different Elm concept we haven't used before: first *subscriptions*, and then *flags*. Let's put them to work!

5.3.1 Receiving realtime data from JavaScript via Ports

Here we'll be receiving realtime data from JavaScript in the same way we receive user input: through a `Msg`. Let's start by introducing a status string to our `Model`, `initialModel`, and `Msg`.

<code>Model</code>	<code>initialModel</code>	<code>Msg</code>
<pre>type alias Model = { photos : List Photo , status : String }</pre>	<pre>initialModel : Model initialModel = { photos = [] , status = "" }</pre>	<pre>type Msg = SelectByUrl String SelectByIndex Int SetStatus String</pre>

With these in place, our change to `update` is straightforward: set the `status` field on `model` when we receive a `SetStatus` message.

```
update msg model =
  case msg of
    SetStatus status ->
      ( { model | status = status }, Cmd.none )
```

We'll also need `view` to display the status:

```
, button
  [ onClick SurpriseMe ]
  [ text "Surprise Me!" ]
, div [ class "status" ] [ text model.status ]
, div [ class "filters" ]
```

Nice! We're getting faster at this.

Now all that's missing is a source of `SetStatus` messages. We'll get those from JavaScript!

SENDING DATA FROM JAVASCRIPT TO ELM

Let's modify `index.html` to add a bit more JavaScript code right below our `setFilters` code.

```
var app = Elm.PhotoGroove.embed(document.getElementById("elm-area"));

app.ports.setFilters.subscribe(function(options) {
  requestAnimationFrame(function() {
    Pasta.apply(document.getElementById("main-canvas"), options);
  });
});

Pasta.addStatusListener(function(status) {  
  ①  
  console.log("Got a new status to send to Elm: ", status); ②  
});
```

- ① Pasta calls this function whenever its status changes
- ② Log the new status to the developer console

This uses Pasta's `addStatusListener` function to log status changes to the browser's JavaScript console. If we reload the page and open the browser's console, we'll see various status messages flowing through as we play around with the sliders and select photos.

Next we need to send those `status` strings to Elm instead of to the console!

SUBSCRIPTIONS

Earlier in the chapter we sent data from Elm to JavaScript using a command. Now we'll do the reverse: send some data from JavaScript to Elm—using not a command, but a *subscription*.

DEFINITION A *subscription* represents a way to translate certain events outside our program into messages that get sent to our `update` function.

One use for subscriptions is handling user inputs that aren't tied to a particular DOM element. For example, we've used `onClick` to translate a click event on a specific element into a message that gets sent to `update`. What if we instead want to detect when the user resizes the entire browser window?

We do this by adding a subscription to our program, which translates global window resize events into messages. Those messages get sent to `update` just like the ones for `onClick` do.

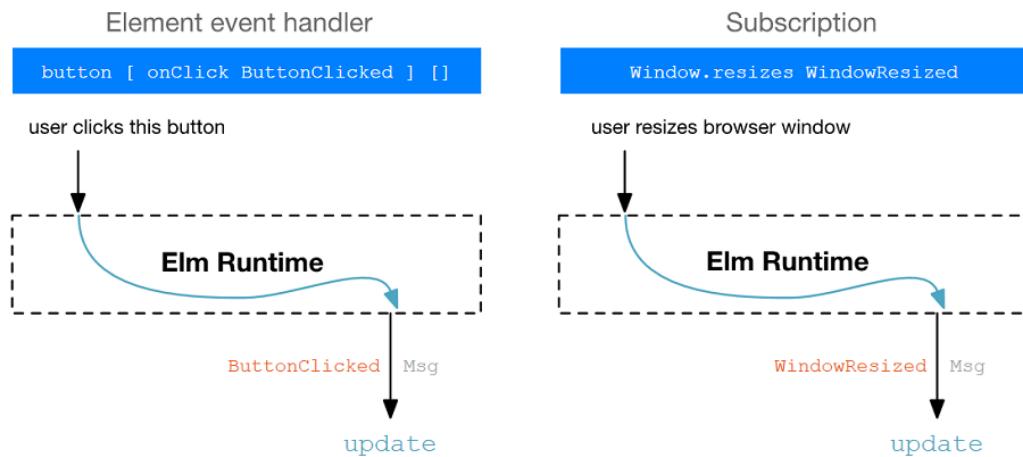


Figure 5.8 A subscription that translates browser window resize events into messages

NOTE If you'd like to try out `Window.resizes` you'll need to run `elm-package install elm-lang/window` first. You can also find subscriptions for global mouse events in the `elm-lang/mouse` package, and subscriptions for global keyboard events in the `elm-lang/keyboard` package.

DEFINING A SUBSCRIPTION PORT

There are also subscriptions which translate data from JavaScript into messages that are sent to `update`. We can get one of these subscriptions by using a slightly different `port` declaration.

Let's add another `port` right below the code that defines `setFilters`:

```
port setFilters : FilterOptions -> Cmd msg
port statusChanges : (String -> msg) -> Sub msg
```

We can already see some similarities and differences between our command port and our subscription port. Both define a function that takes one argument, but whereas the first function returned a command (`Cmd msg`), the second returns a subscription (`Sub msg`).

`Cmd` and `Sub` are both parameterized on the type of message they produce. We noted earlier how `setFilters` returns a `Cmd msg` (as opposed to `Cmd Msg`) because it is a command

which produces no message after it completes. In contrast, `statusChanges` returns a `Sub msg`, but here `msg` refers to the type of message returned by the `(String -> msg)` function we pass to `statusChanges`.

Table 5.9 shows how various calls to `statusChange` can yield different return values.

Table 5.9 Calling the `statusChanges` function

Function to pass to <code>statusChanges</code>	Expression	Return Type
<code>String.length : String -> Int</code>	<code>statusChanges String.length</code>	<code>Sub Int</code>
<code>String.reverse : String -> String</code>	<code>statusChanges String.reverse</code>	<code>Sub String</code>
<code>SetStatus : String -> Msg</code>	<code>statusChanges SetStatus</code>	<code>Sub Msg</code>

NOTE Whereas it's normal for `setFilters` to return `Cmd msg`, it would be bizarre for `statusChanges` to return `Sub msg`. After all, a `Cmd msg` is a command that has an effect but never sends a message to update...but subscriptions do not run effects. Their whole purpose is to send messages to update! Subscribing to a `Sub msg` would be like listening to a disconnected phone line: not terribly practical.

If we call `statusChanges SetStatus`, we'll get back a `Sub Msg` subscription. That's all well and good, but what do we do with a `Sub Msg`?

PASSING SUBSCRIPTIONS TO HTML.PROGRAM

We use it with that `subscriptions` field we set up with `Html.program!`

```
main : Program Never Model Msg
main =
    Html.program
        { init = ( initialModel, initialCmd )
        , view = viewOrError
        , update = update
        , subscriptions = \_ -> Sub.none
        , subscriptions = \_ -> statusChanges SetStatus
        }
```

Ever since Chapter 3 we've been setting this `subscriptions` field to an anonymous function that always returned `Sub.none`. Now we've made that anonymous function return `statusChanges SetStatus` instead, which means that whenever JavaScript sends a string to the `statusChanges` port, it will result in a `SetStatus` message being sent to `update`.

NOTE The argument this anonymous `subscriptions` function accepts is a `Model`. Whenever our model changes, the new model is passed to this function, giving us a chance to return a different `Sub` depending on what's in the new model. This lets us dynamically control which subscriptions our program pays attention to.

Figure 5.9 illustrates how subscriptions fit into Elm programs.

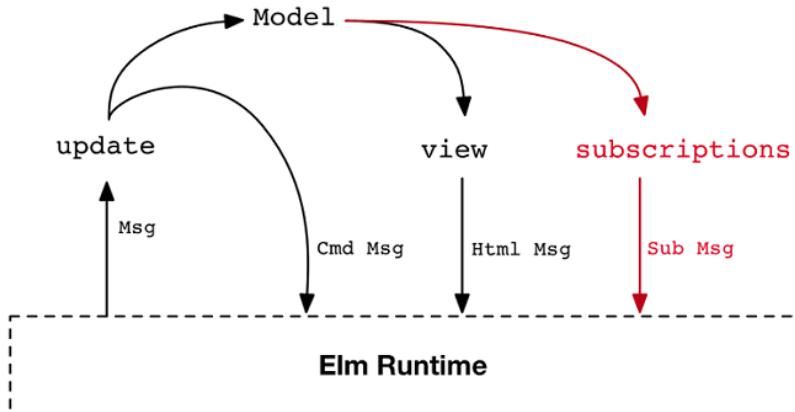


Figure 5.9 How Subscriptions fit into Elm programs

Now our `statusChanges` port is fully connected to the rest of our application, meaning we're ready to have our JavaScript code start sending it data!

CALLING APP.PORTS.STATUSCHANGES.SEND

Much like how we used `app.ports.setFilters.subscribe` to receive data from Elm, we can use `app.ports.statusChanges.send` to send data to Elm. Let's replace our `console.log` in `index.html` with a call to `app.ports.statusChanges.send`:

```
Pasta.addEventListener("statusChange", function(status) {
    console.log("Got a new status to send to Elm: ", status);
    app.ports.statusChanges.send(status);
});
```

Remember, `subscriptions = _ -> statusChanges SetStatus` specifies that we'll wrap whatever `status` string we receive from the `statusChanges` port in a `SetStatus` message.

This means that calling `app.ports.statusChanges.send("Reticulating splines")` from JavaScript will ultimately result in a `SetStatus "Reticulating splines"` message being sent to `update` on the Elm side. At that point, our existing status-rendering logic should kick in, which we can confirm by recompiling and seeing how things look.

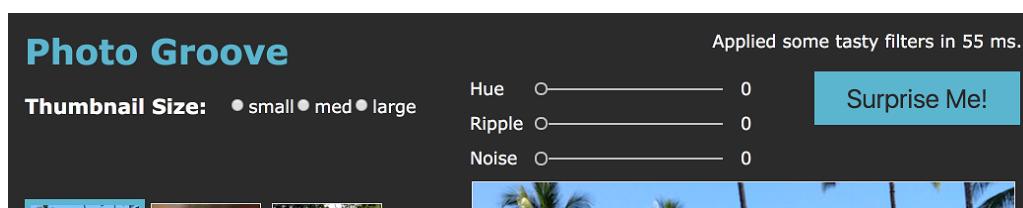


Figure 5.10 Showing the status when filters change

Fantastic! There's just one more problem to address. Right when we refresh the page, at first there is no status to show, leading to a brief flicker of nothingness before the first status update arrives. We'll fix this by showing an initial status, using data obtained from Pasta prior to the first status change it sends us.

5.3.2 Receiving initialization arguments via Flags

We'd like to display an initial status which reports the version of Pasta we're using. It will look something like this:

Initializing with Pasta v4.2...

We can easily obtain the current version from Pasta in JavaScript; all we have to do is reference `Pasta.version`. But how do we get that value into Elm in time for the initial render, to avoid that flicker of missing content?

SENDING INITIALIZATION DATA TO PROGRAMWITHFLAGS

Calling `app.ports.statusChanges.send(Pasta.version)` won't cut it here. Elm would still render `initialModel` (with `model.status` being "", resulting in the flicker of missing content) before receiving a message from `statusChanges` and re-rendering with the new `status`.

What we want is a way to give Elm a value from JavaScript that's available right away, early enough to use in our initial. This means we don't want ports...we want *flags*.

DEFINITION *Flags* are values passed to an Elm program's initialization function.

To use flags we'll need to switch from `Html.program` to `Html.programWithFlags`. The primary difference between the two is the type of `init` that each accepts. `program` wants `init` to be a `(Model, Cmd Msg)` tuple, whereas `programWithFlags` wants it to be a function which accepts initialization flags as an argument and returns a `(Model, Cmd Msg)` tuple.

This change lets us receive a `Float` from JavaScript—namely `Pasta.version`—as an argument to `init`, giving us the chance to incorporate it into our initial model like so:

```
main =
    Html.programWithFlags
        { init = init
        , view = viewOrError
        , update = update
        , subscriptions = (\_ -> Sub.none)
        }

init : Float -> ( Model, Cmd Msg )
init flags =
    let
        status =
            "Initializing Pasta v" ++ toString flags
    in
```

```
( { initialModel | status = status }, initialCmd )
```

Switching to `programWithFlags` affects our type annotation for `main` a bit:

```
main : Program Never Model Msg
main : Program Float Model Msg
```

`Program Never Model Msg` means “a `Program` with `Model` as its model type, `Msg` as its message type, and which has no flags.” `Program Float Model Msg` means the same thing, except that its flags are a `Float`.

The Never Type

The type of a `Html.program` always begins with `Program Never`. What is this `Never` type, exactly?

The definition of `Never` lives in the `Basics` module, and its implementation looks something like this:

```
type Never
  = PassMeNever Never
```

How might we obtain a value of type `Never`? That’s easy, just call `PassMeNever!`! What does `PassMeNever` take as an argument? Oh, a `Never`. Hm...how can we obtain a `Never` to pass to `PassMeNever`? Ah yes—call `PassMeNever!`! Now what type of argument does it take again?

To make an endlessly long story short, it’s not possible to instantiate a `Never`. Any function which accepts a `Never` for an argument is a function that can never be called! The reason `Html.program` returns a `Program Never Model Msg` is that it can never be passed any flags, `Float` or otherwise. As you may recall, `Html.program`’s `init` is not even a function, so we can never call it!

SENDING FLAGS TO ELM

On the JavaScript side, we can send flags to Elm by passing an extra argument to the `Elm.PhotoGroove.embed` method we’ve been using to start the program running.

```
var app = Elm.PhotoGroove.embed(document.getElementById("elm-area"), Pasta.version);
```

Decoding JavaScript Values

We can use the same Decoders we used in Chapter 4 to translate values from JavaScript into Elm values while gracefully handling errors.

In Chapter 4 we used the `Json.Decode.decodeString` function to decode a String of JSON into various Elm values. There is a similar function, `Json.Decode.decodeValue`, which takes a `Json.Decode.Value` instead of a String. We can write our `statusChanges` port to expect a `Value` from JavaScript, like so:

```
import Json.Decode exposing (Value)

port statusChanges : (String --> msg) --> Sub msg
port statusChanges : (Value -> msg) -> Sub msg
```

From there, writing a Decoder Msg and passing it to Json.Decode.decodeValue will give us a Result String Msg. We can translate that into a Msg by introducing an error-handling Msg which ideally displays a helpful message to end users, explaining that something went wrong.

We can use the same technique for Flags:

```
main : Program Value Model Msg
```

Using this decoder style is generally a better choice for production applications. For one, it lets us give users a better experience by gracefully handling errors. For another, any incoming type other than Value is liable to result in a runtime exception if JavaScript passes in something invalid!

Elm will check this immediately upon receiving the value, and will throw the exception early rather than letting it snake its way through our program, but it's better still to do the error handling ourselves. Besides, since Elm performs the automatic check using decoders anyway, it's not as if there's a performance penalty for doing it explicitly!

That's it! Now when we start up Photo Groove, we'll see the version number proudly displaying for the briefest of moments before the first photo loads. Figure 5.11 shows the initial status, captured using an incredible high-speed screenshotting camera.



Figure 5.11 Showing the status on page load

Nicely done!

5.4 Summary

Photo Groove is now substantially groovier than before. Specifically:

- We added some cool sliders using Custom Elements written in JavaScript
- As users slide their values around, it changes the filters on our large photo
- The filters come from a JavaScript library, which Photo Groove talks to like a server
- When the JS filtering library sends status reports, Photo Groove displays them
- The initial status report includes the JS filtering library's version number

In the process we learned some new concepts and techniques for building Elm applications:

- We can resolve cyclical dependency errors by making functions more flexible
- Asking “Which approach rules out more bugs?” is a good way to decide between different ways to model data
- Once Custom Elements have been registered on the page, we can access them using the `Html.node` function
- The `as` keyword lets us alias imported modules, for example `import Html.Attributes`

- as Attr
- The `Html.Attributes.on` function lets us create a custom event handler using a Decoder
- `requestAnimationFrame` can delay our JavaScript code's execution until the next time Elm calls `view` and renders the result to the DOM
- A *subscription* represents a way to translate certain events outside our program into messages that get sent to our `update` function
- A `port` module can use the `port` keyword to define functions that return commands and subscriptions which talk to JavaScript
- `Html.programWithFlags` lets us send initialization flags from JavaScript to Elm

In Chapter 6 we'll crank our application's maintainability into overdrive by building it a hearty suite of automated tests!

Listing 5.6 index.html

```
<!doctype html>
<html>
  <head>
    <link rel="stylesheet" href="http://elm-in-action.com/styles.css">
    <link rel="import" href="http://elm-in-action.com/wc/slider/paper-slider.html">
  </head>

  <body>
    <div id="elm-area"></div>

    <script src="http://elm-in-action.com/pasta.js"></script>
    <script src="elm.js"></script>
    <script>
      var app = Elm.PhotoGroove.embed(document.getElementById("elm-area"),
Pasta.version);

      app.ports.setFilters.subscribe(function(options) {
        requestAnimationFrame(function() {
          Pasta.apply(document.getElementById("main-canvas"), options);
        });

        Pasta.addStatusListener(function(status) {
          app.ports.statusChanges.send(status);
        });
      });
    </script>
  </body>
</html>
```

Listing 5.7 PhotoGroove.elm

```
port module PhotoGroove exposing (..)

import Html exposing (..)
import Html.Events exposing (onClick)
import Array exposing (Array)
```

```

import Random
import Http
import Html.Attributes exposing (id, class, classList, src, name, max, type_, title)
import Json.Decode exposing (string, int, list, Decoder)
import Json.Decode.Pipeline exposing (decode, required, optional)
import Slider exposing (paperSlider)

photoDecoder : Decoder Photo
photoDecoder =
    decode buildPhoto
        |> required "url" string
        |> required "size" int
        |> optional "title" string "(untitled)"

buildPhoto : String -> Int -> String -> Photo
buildPhoto url size title =
    { url = url, size = size, title = title }

urlPrefix : String
urlPrefix =
    "http://elm-in-action.com/"

type ThumbnailSize
    = Small
    | Medium
    | Large

view : Model -> Html Msg
view model =
    div [ class "content" ]
        [ h1 [] [ text "Photo Groove" ]
        , button
            [ onClick SurpriseMe ]
            [ text "Surprise Me!" ]
        , div [ class "status" ] [ text model.status ]
        , div [ class "filters" ]
            [ viewFilter SetHue "Hue" model.hue
            , viewFilter SetRipple "Ripple" model.ripple
            , viewFilter SetNoise "Noise" model.noise
            ]
        , h3 [] [ text "Thumbnail Size:" ]
        , div [ id "choose-size" ]
            (List.map viewSizeChooser [ Small, Medium, Large ])
        , div [ id "thumbnails", class (sizeToString model.chosenSize) ]
            (List.map (viewThumbnail model.selectedUrl) model.photos)
        , viewLarge model.selectedUrl
        ]

viewFilter : (Int -> Msg) -> String -> Int -> Html Msg
viewFilter toMsg name magnitude =
    div [ class "filter-slider" ]
        [ label [] [ text name ] ]

```

```

    , paperSlider [ Attr.max "11", Slider.onImmediateValueChange toMsg ] []
    , label [] [ text (toString magnitude) ]
]

viewLarge : Maybe String -> Html Msg
viewLarge maybeUrl =
  case maybeUrl of
    Nothing ->
      text ""
    Just url ->
      canvas [ id "main-canvas", class "large" ] []

viewThumbnail : Maybe String -> Photo -> Html Msg
viewThumbnail selectedUrl thumbnail =
  img
    [ src (urlPrefix ++ thumbnail.url)
    , title (thumbnail.title ++ " [" ++ toString thumbnail.size ++ " KB]")
    , classList [ ( "selected", selectedUrl == Just thumbnail.url ) ]
    , onClick (SelectByUrl thumbnail.url)
    ]
  []

viewSizeChooser : ThumbnailSize -> Html Msg
viewSizeChooser size =
  label []
    [ input [ type_ "radio", name "size", onClick (SetSize size) ] []
    , text (sizeToString size)
    ]

sizeToString : ThumbnailSize -> String
sizeToString size =
  case size of
    Small ->
      "small"
    Medium ->
      "med"
    Large ->
      "large"

port setFilters : FilterOptions -> Cmd msg

port statusChanges : (String -> msg) -> Sub msg

type alias FilterOptions =
  { url : String
  , filters : List { name : String, amount : Float }
  }

```

```

type alias Photo =
    { url : String
    , size : Int
    , title : String
    }

type alias Model =
    { photos : List Photo
    , status : String
    , selectedUrl : Maybe String
    , loadingError : Maybe String
    , chosenSize : ThumbnailSize
    , hue : Int
    , ripple : Int
    , noise : Int
    }

initialModel : Model
initialModel =
    { photos = []
    , status = ""
    , selectedUrl = Nothing
    , loadingError = Nothing
    , chosenSize = Medium
    , hue = 0
    , ripple = 0
    , noise = 0
    }

photoArray : Array Photo
photoArray =
    Array.fromList initialModel.photos

getPhotoUrl : Int -> Maybe String
getPhotoUrl index =
    case Array.get index photoArray of
        Just photo ->
            Just photo.url
        Nothing ->
            Nothing

type Msg
    = SelectByUrl String
    | SelectByIndex Int
    | SetStatus String
    | SurpriseMe
    | SetSize ThumbnailSize
    | SetHue Int
    | SetRipple Int
    | SetNoise Int
    | LoadPhotos (Result Http.Error (List Photo))

```

```

randomPhotoPicker : Random.Generator Int
randomPhotoPicker =
    Random.int 0 (Array.length photoArray - 1)

update : Msg -> Model -> ( Model, Cmd Msg )
update msg model =
    case msg of
        SetStatus status ->
            ( { model | status = status }, Cmd.none )

        SetHue hue ->
            applyFilters { model | hue = hue }

        SetRipple ripple ->
            applyFilters { model | ripple = ripple }

        SetNoise noise ->
            applyFilters { model | noise = noise }

        SelectByIndex index ->
            let
                newSelectedUrl : Maybe String
                newSelectedUrl =
                    model.photos
                        |> Array.fromList
                        |> Array.get index
                        |> Maybe.map .url
            in
                applyFilters { model | selectedUrl = newSelectedUrl }

        SelectByUrl selectedUrl ->
            applyFilters { model | selectedUrl = Just selectedUrl }

        SurpriseMe ->
            let
                randomPhotoPicker =
                    Random.int 0 (List.length model.photos - 1)
            in
                ( model, Random.generate SelectByIndex randomPhotoPicker )

        setSize size ->
            ( { model | chosenSize = size }, Cmd.none )

        LoadPhotos (Ok photos) ->
            applyFilters
                { model
                    | photos = photos
                    , selectedUrl = Maybe.map .url (List.head photos)
                }

        LoadPhotos (Err _) ->
            ( { model
                | loadingError = Just "Error! (Try turning it off and on again?)"
            }
            , Cmd.none
    )
}

```

```

        )

applyFilters : Model -> ( Model, Cmd Msg )
applyFilters model =
    case model.selectedUrl of
        Just selectedUrl ->
            let
                filters =
                    [ { name = "Hue", amount = (toFloat model.hue) / 11 }
                    , { name = "Ripple", amount = (toFloat model.ripple) / 11 }
                    , { name = "Noise", amount = (toFloat model.noise) / 11 }
                    ]
            in
                ( model, setFilters { url = url, filters = filters } )

        Nothing ->
            ( model, Cmd.none )

initialCmd : Cmd Msg
initialCmd =
    list photoDecoder
        |> Http.get "http://elm-in-action.com/photos/list.json"
        |> Http.send LoadPhotos

viewOrError : Model -> Html Msg
viewOrError model =
    case model.loadingError of
        Nothing ->
            view model

        Just errorMessage ->
            div [ class "error-message" ]
                [ h1 [] [ text "Photo Groove" ]
                , p [] [ text errorMessage ]
                ]

main : Program Float Model Msg
main =
    Html.programWithFlags
        { init = init
        , view = viewOrError
        , update = update
        , subscriptions = \_ -> statusChanges SetStatus
        }

init : Float -> ( Model, Cmd Msg )
init flags =
    let
        status =
            "Initializing Pasta v" ++ toString flags

```

```
in
  ( { initialModel | status = status }, initialCmd )

paperSlider : List (Attribute msg) -> List (Html msg) -> Html msg
paperSlider =
  node "paper-slider"

onImmediateValueChange : (Int -> msg) -> Attribute msg
onImmediateValueChange toMsg =
  at [ "target", "immediateValue" ] int
    |> Json.Decode.map toMsg
    |> on "immediate-value-changed"
```

A

Getting Set Up

This covers installing

- Node.js (4.0 or higher) and NPM
- The Elm Platform
- The Elm in Action Repository
- Recommended optional tools

INSTALLING NODE.JS AND NPM

In addition to having the Elm Platform installed, the examples in this book require having Node.js 4.0 or higher and NPM 2.0 or higher installed as well. If you haven't already, visit <https://nodejs.org> and follow the instructions to download and install them.

NOTE The Node.js installer on <https://nodejs.org> also installs NPM for you, so if you install Node that way, you will not need to install NPM separately.

To confirm that you have Node.js 4.0 or higher and NPM 2.0 or higher installed, you should be able to run these commands in your terminal and see similar output:

```
$ node --version
v4.0.0

$ npm --version
2.0.0
```

If you have Node installed but not NPM, then your installation of Node is probably not the one from <http://nodejs.org>. Please make sure you have NPM installed before continuing!

INSTALLING THE ELM PLATFORM

Now that you have NPM installed, you can use it to get the Elm Platform:

```
npm install -g elm
```

TIP If `npm` gives you a lengthy error involving the word `EACCES`, visit <https://docs.npmjs.com/getting-started/fixing-npm-permissions> for how to resolve it. If you are unable to resolve it, you can fall back on installing the Elm Platform directly from <http://elm-lang.org/install> - but this will only get you through Chapters 1 through 5. Starting in Chapter 6, being able to run `npm install -g` will be required to run the examples!

Let's verify that the Elm platform was installed properly:

```
elm-make --version
elm-make 0.18.0 (Elm Platform 0.18.0)
```

If you see a version higher than `0.18.0`, then as of this writing, you are living in the future! Hello from the past, where `0.18.0` is the latest release.

NOTE Any version that starts with `0.18` should work fine with this book, but according to semantic versioning, a version number beginning with `0.19` or higher indicates breaking changes. In such a case, there's a good chance your examples will not compile! Try `npm install -g elm@0.18.0` to get a compatible version.

OBTAINING THE ELM IN ACTION REPOSITORY

This book does not require using Git or any Git knowledge. If you are not a Git user, you can download and extract a zip archive of the repository at <https://github.com/rtfeldman/elm-in-action/archive/master.zip> and proceed directly to the next section!

Git users can download the repository by running this:

```
git clone https://github.com/rtfeldman/elm-in-action.git
```

The repository has been tagged with various checkpoints. Suppose you are about to read Chapter 2, Section 2.2.1. You can visit this URL to obtain all the code necessary to work through the examples in Section 2.2.1:

```
https://github.com/rtfeldman/elm-in-action/tree/2.2.1
```

You can replace the `2.2.1` in that URL with a different section number to view the code at the start of that section.

For example, if you want to peek ahead and see where things stand at the end of Section 2.2.1, you can bring up the `2.2.2` tag by visiting <https://github.com/rtfeldman/elm-in-action/tree/2.2.2> in a browser. Alternatively, you can run `git checkout 2.2.2` from a terminal if you ran `git clone` on the repository earlier.

INSTALLING RECOMMENDED OPTIONAL TOOLS

To get syntax highlighting and other niceties, visit <http://elm-lang.org/install#syntax-highlighting> and select your editor of choice to find an appropriate Elm plugin. Make sure your editor is configured to convert tabs to spaces, as tab characters are syntax errors in Elm!

I also strongly recommend installing *elm-format*: <https://github.com/avh4/elm-format>

I absolutely love *elm-format*. It automatically formats Elm code according to a nice, consistent standard. I use it for personal projects, my whole team at work uses it, and I have my editor configured to run it whenever I save. (Several editors enable this by default, but if not, there are easy instructions in the above link for how to enable format-on-save.)

All of the examples in this book are formatted with *elm-format*, so using it will make life easier as you work through them!